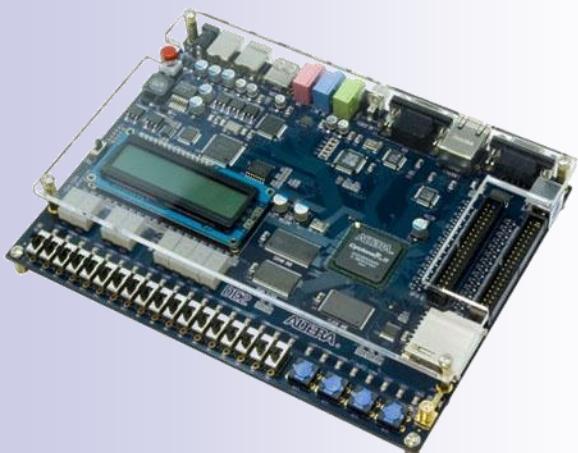


Logic Systems and Processors

cz:Logické systémy a procesory



Lecturer: Richard Šusta

richard@susta.cz, susta@fel.cvut.cz,

+420 2 2435 7359

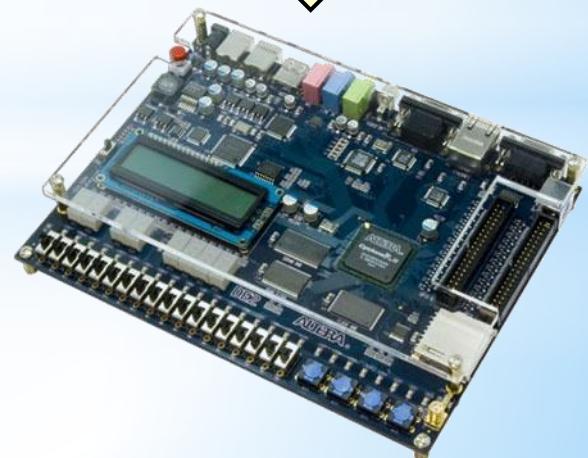
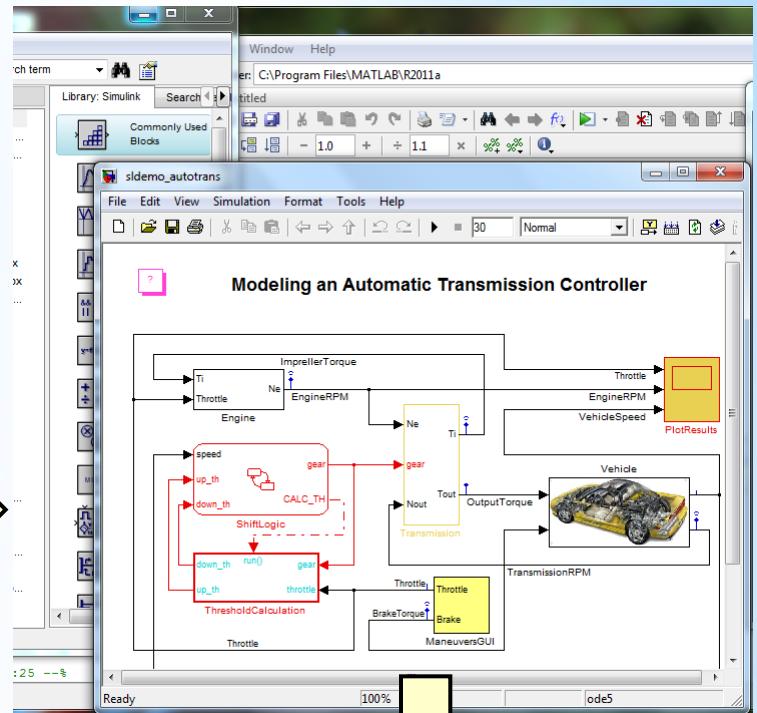
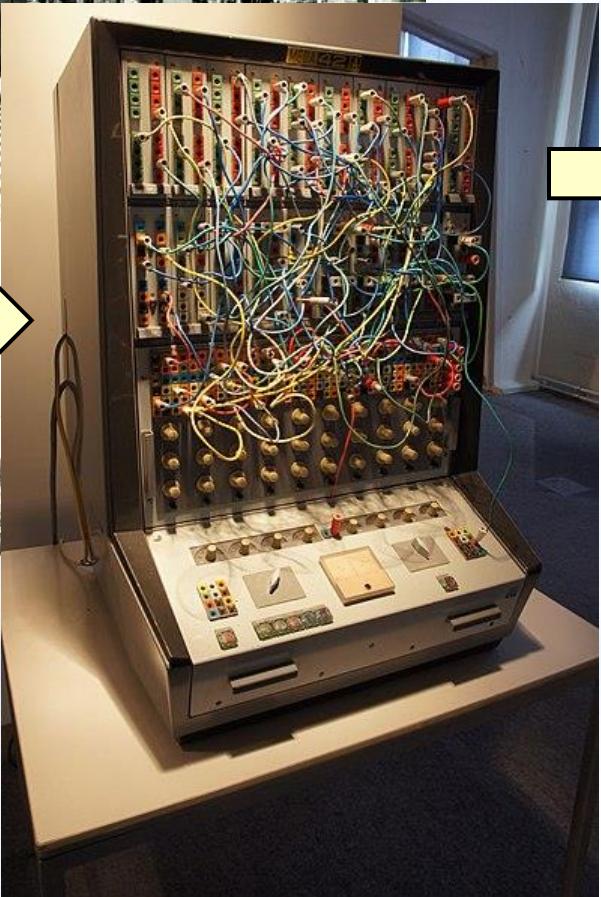
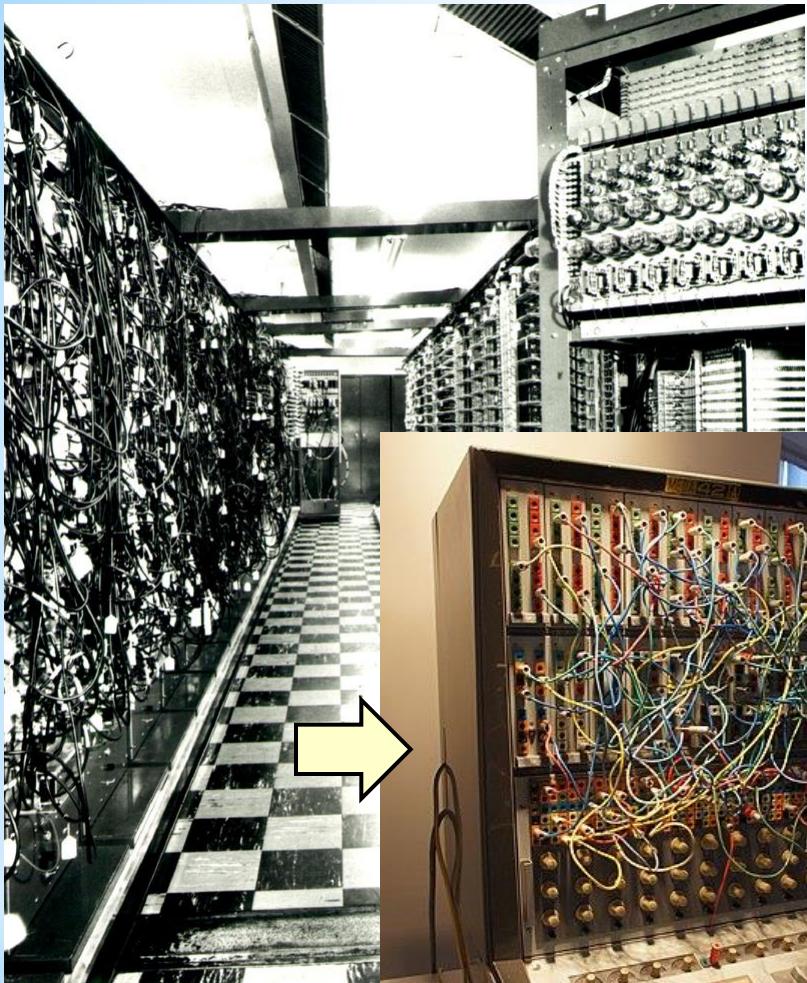
V1.1: Jan 7, 2025
extended slides 45 and 71

Pipeline ≈ Assembly Line



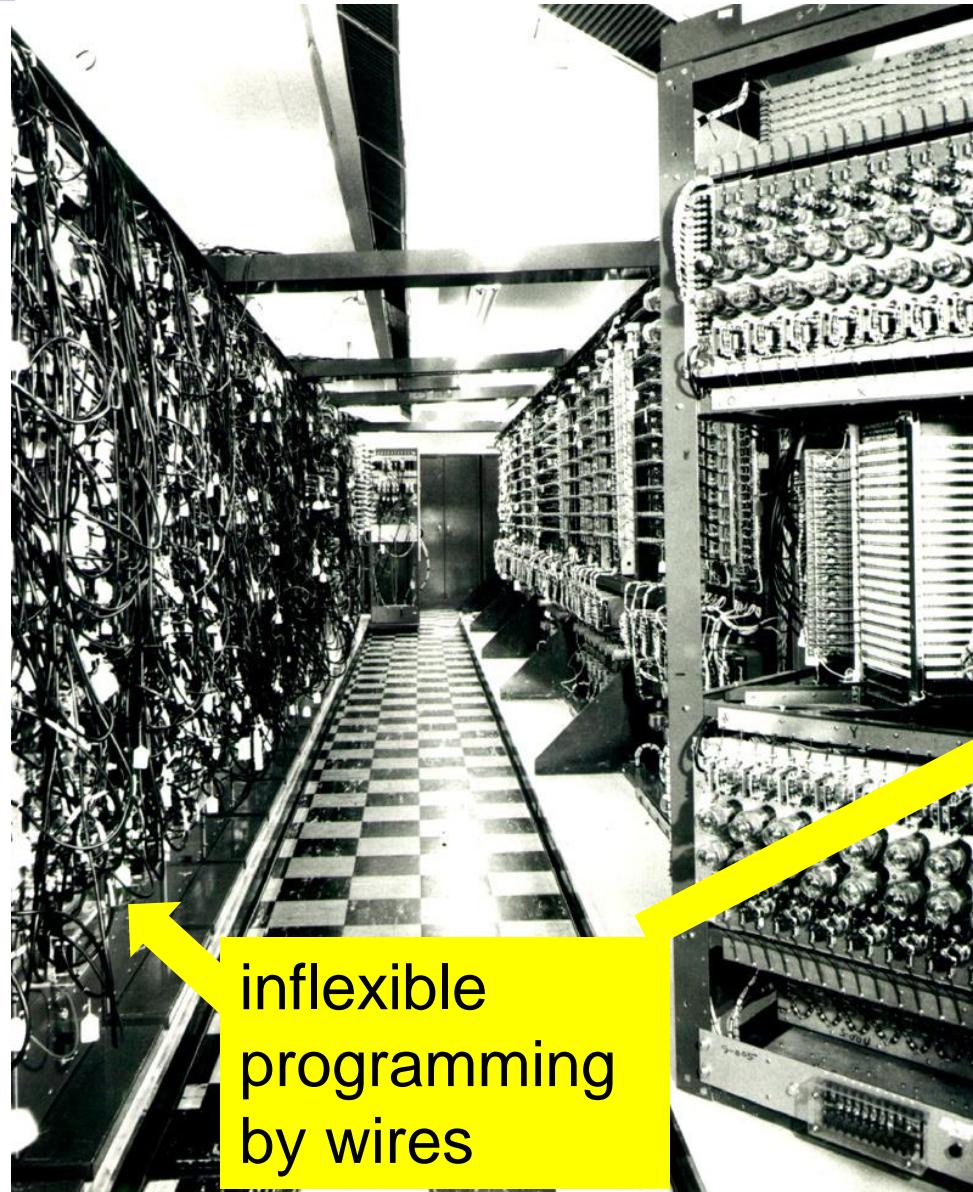
Charlie Chaplin - Factory Scene - Modern Times © Roy Export S.A.S. Scan, courtesy of Cineteca di Bologna

* Parallel Analog Operations



All images courtesy of Wikipedia

1943 MIT's Whirlwind Analog Computer



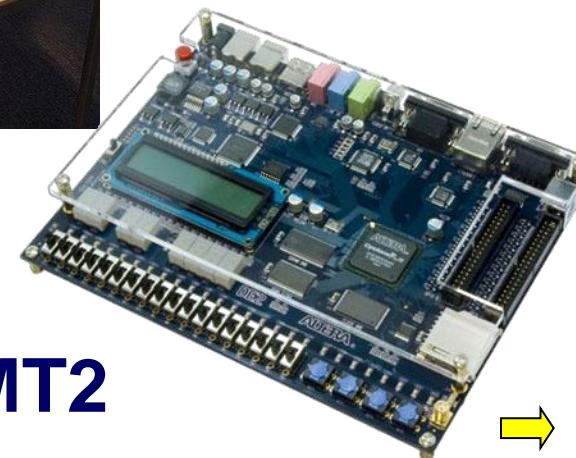
inflexible
programming
by wires

Designed as a flight simulator for training bomber crews, it simulated a realistic aerodynamic system operating in real time. But it was inaccurate and inflexible.



1980 MEDA
Aritma
Vysočany
1956-1987

2024
VEEK-MT2



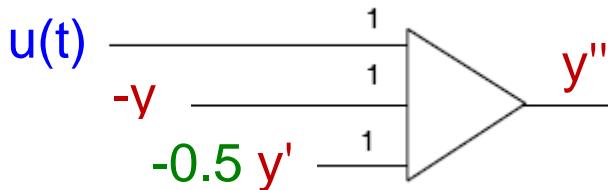
1. Laplace transformation

2. Its differential equations

3. Simulink

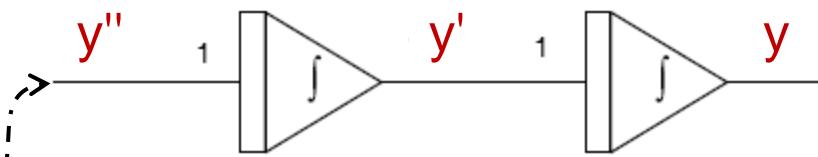
$$y'' = u(t) - 0.5 y' - y$$

3b. Summation member



$$F(p) = 1/(p^2 + 0.5 p + 1)$$

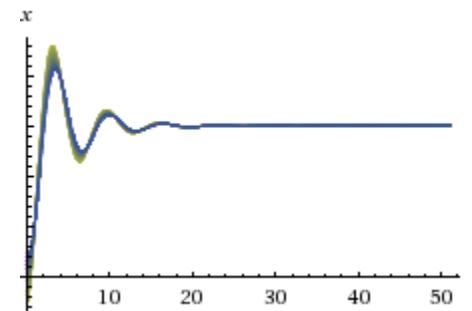
$$y'' + 0.5 y' + y = u(t)$$



3a. Two integrators

3c. Response to unit jump

$$u(t)=0.3; y(0)=0; y'(0)=1$$



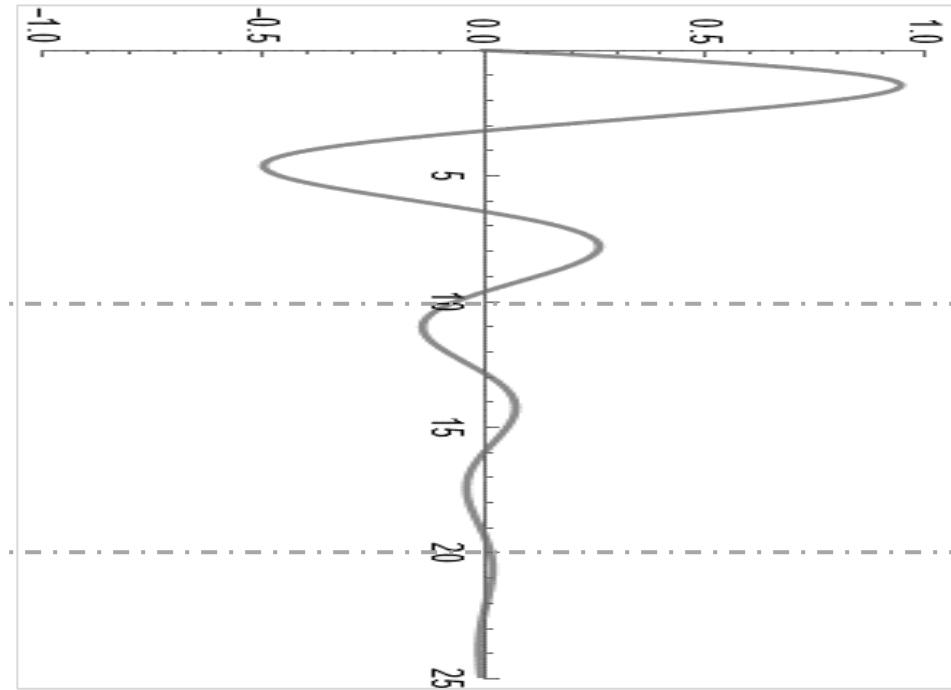
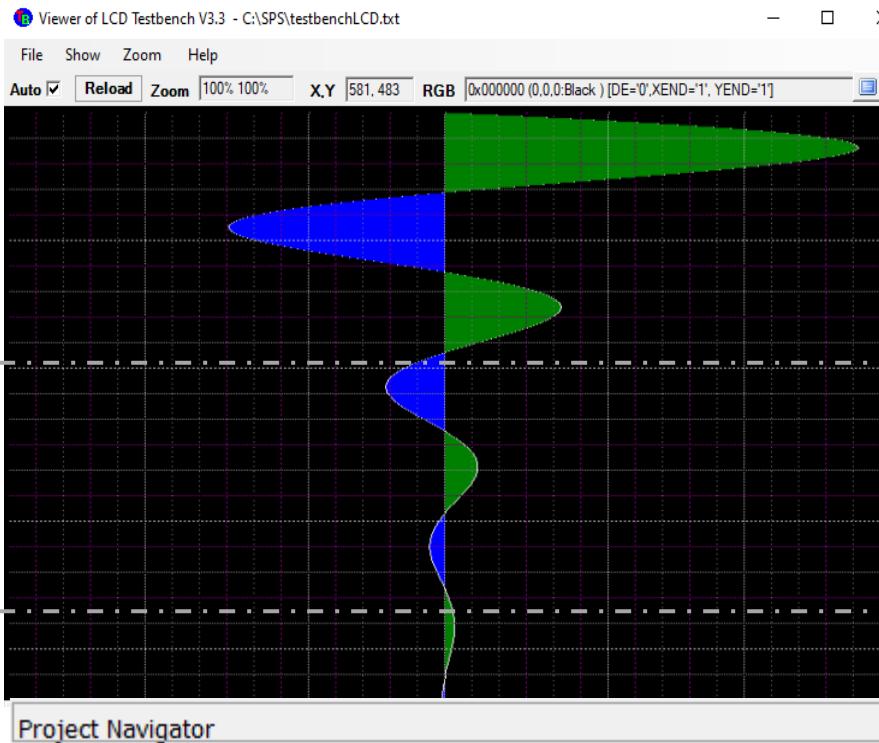
4. Numerical solution for unit jump $u(t)=0.3$ for $t \geq 0$

$$y(t) = 0.955336 e^{-0.25t} \sin(0.968246 t) - 0.3 e^{-0.25t} \cos(0.968246 t) + 0.3$$

[calculated using WolframAlpha]

Example: the VEEK-MT2

$$y''[t] + 0.4 y'[t] + y[t] == 0, \quad y'[0]==1.25, \quad y[0]==0$$



$$y[t] \rightarrow 1.27578 e^{-0.2 t} \sin[0.979796 t]$$

Wolfram Mathematica 14

Entity	Logic Cells	DSP 18x18
Cyclone IV E: EP4CE115F29C7		
Integrator	639 (1)	3
VeekMT2_LCDgenerator:inst	48 (47)	0
VeekMT2_LCDreg:inst1	44 (44)	0
LCDGraph:inst2	562 (168)	3
DifferentialEquation:integrator	400 (400)	2
lpm_mult:Mult0	0 (0)	1
lpm_mult:Mult1	0 (0)	1
lpm_mult:Mult0	0 (0)	1

In an FPGA, all elements of the iteration are emulated in parallel (concurrently), and so the response speed is independent of the number of elements, i.e., the sizes of modeled systems.

The famous mechanical integration solution



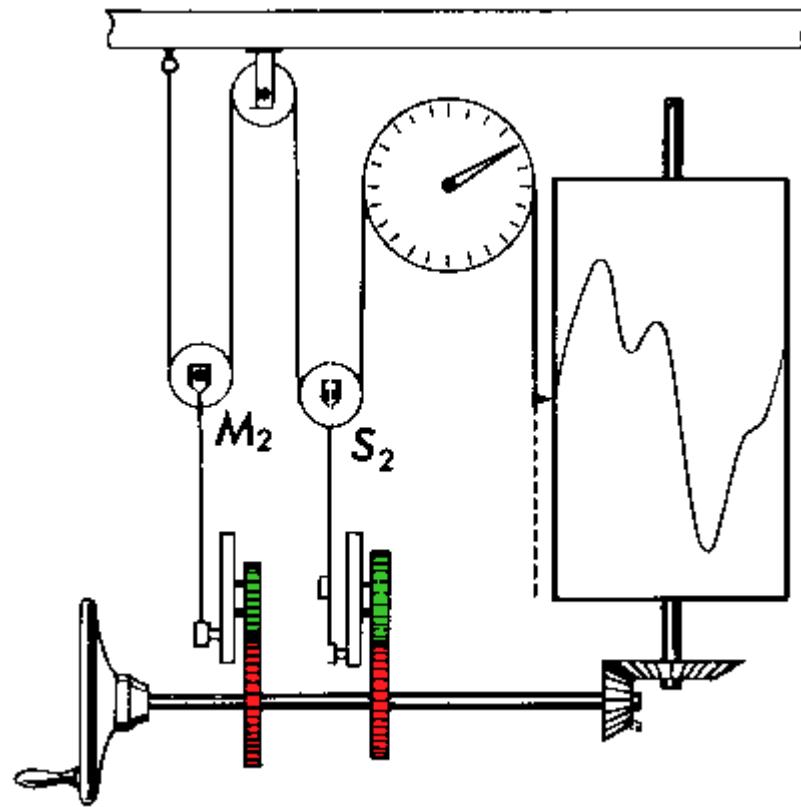
[Source: <http://www.mit.edu/~klund/analyzer/>]

Vannevar Bush
Differential Analyzer



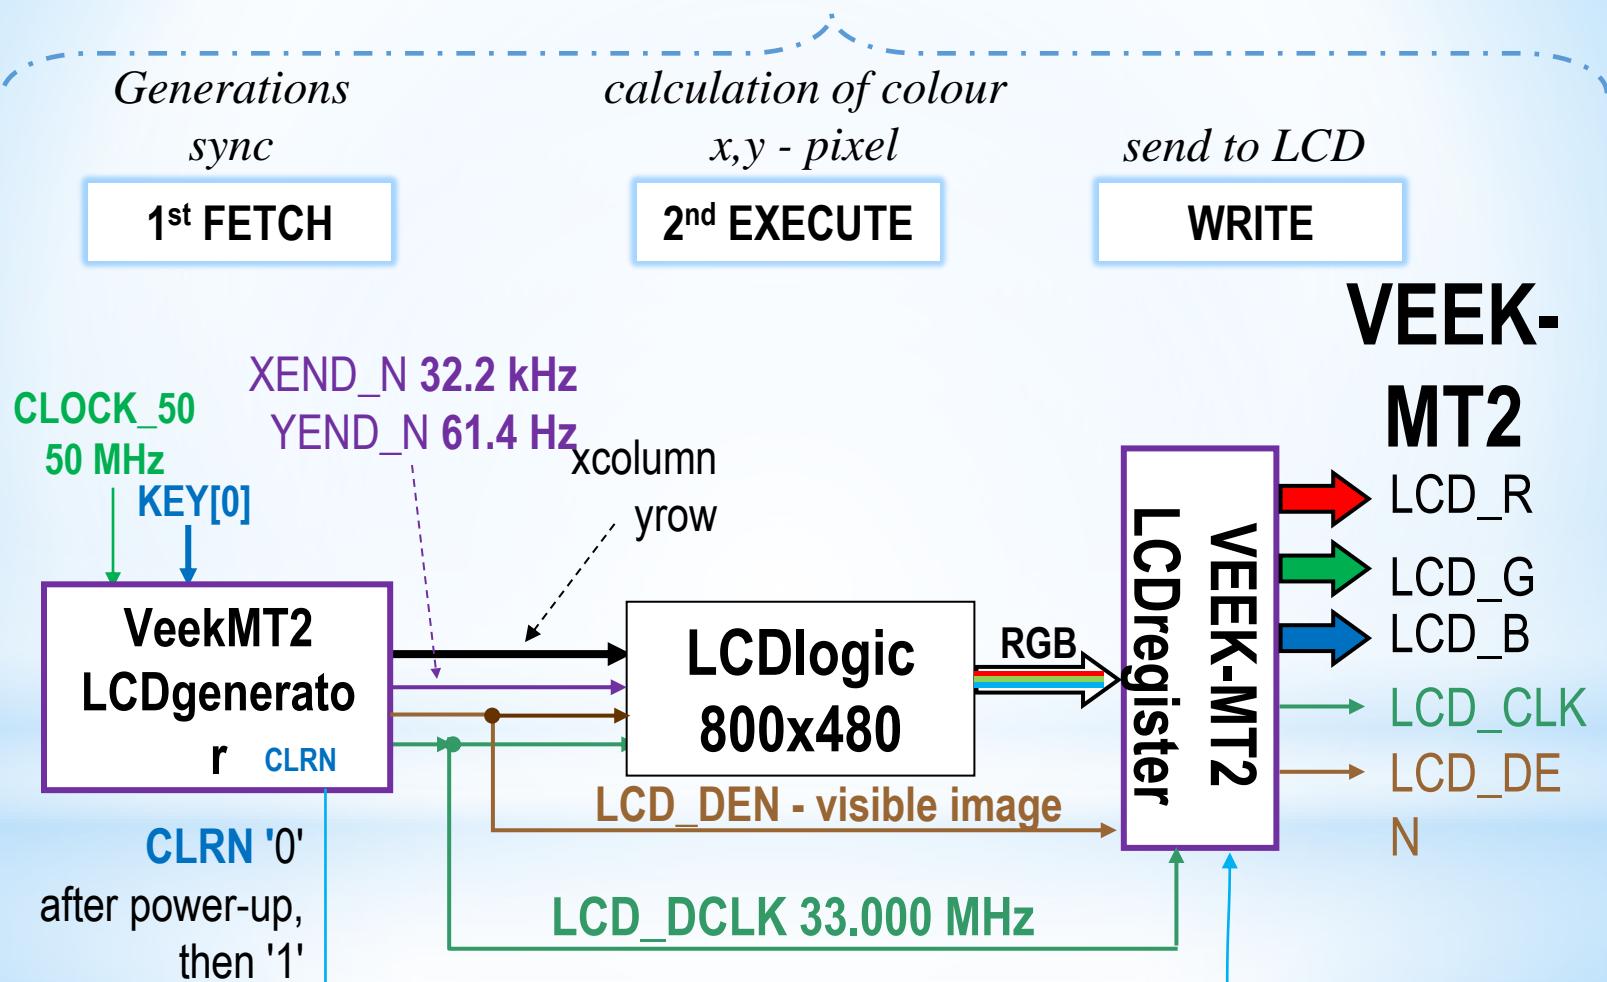
Video:

<http://www.mit.edu/~klund/analyzer/video3.mov>



Variable-size friction wheels
for simulating the behaviour
of differential equations

You know **Pipeline** from LCD it approximates analog processes



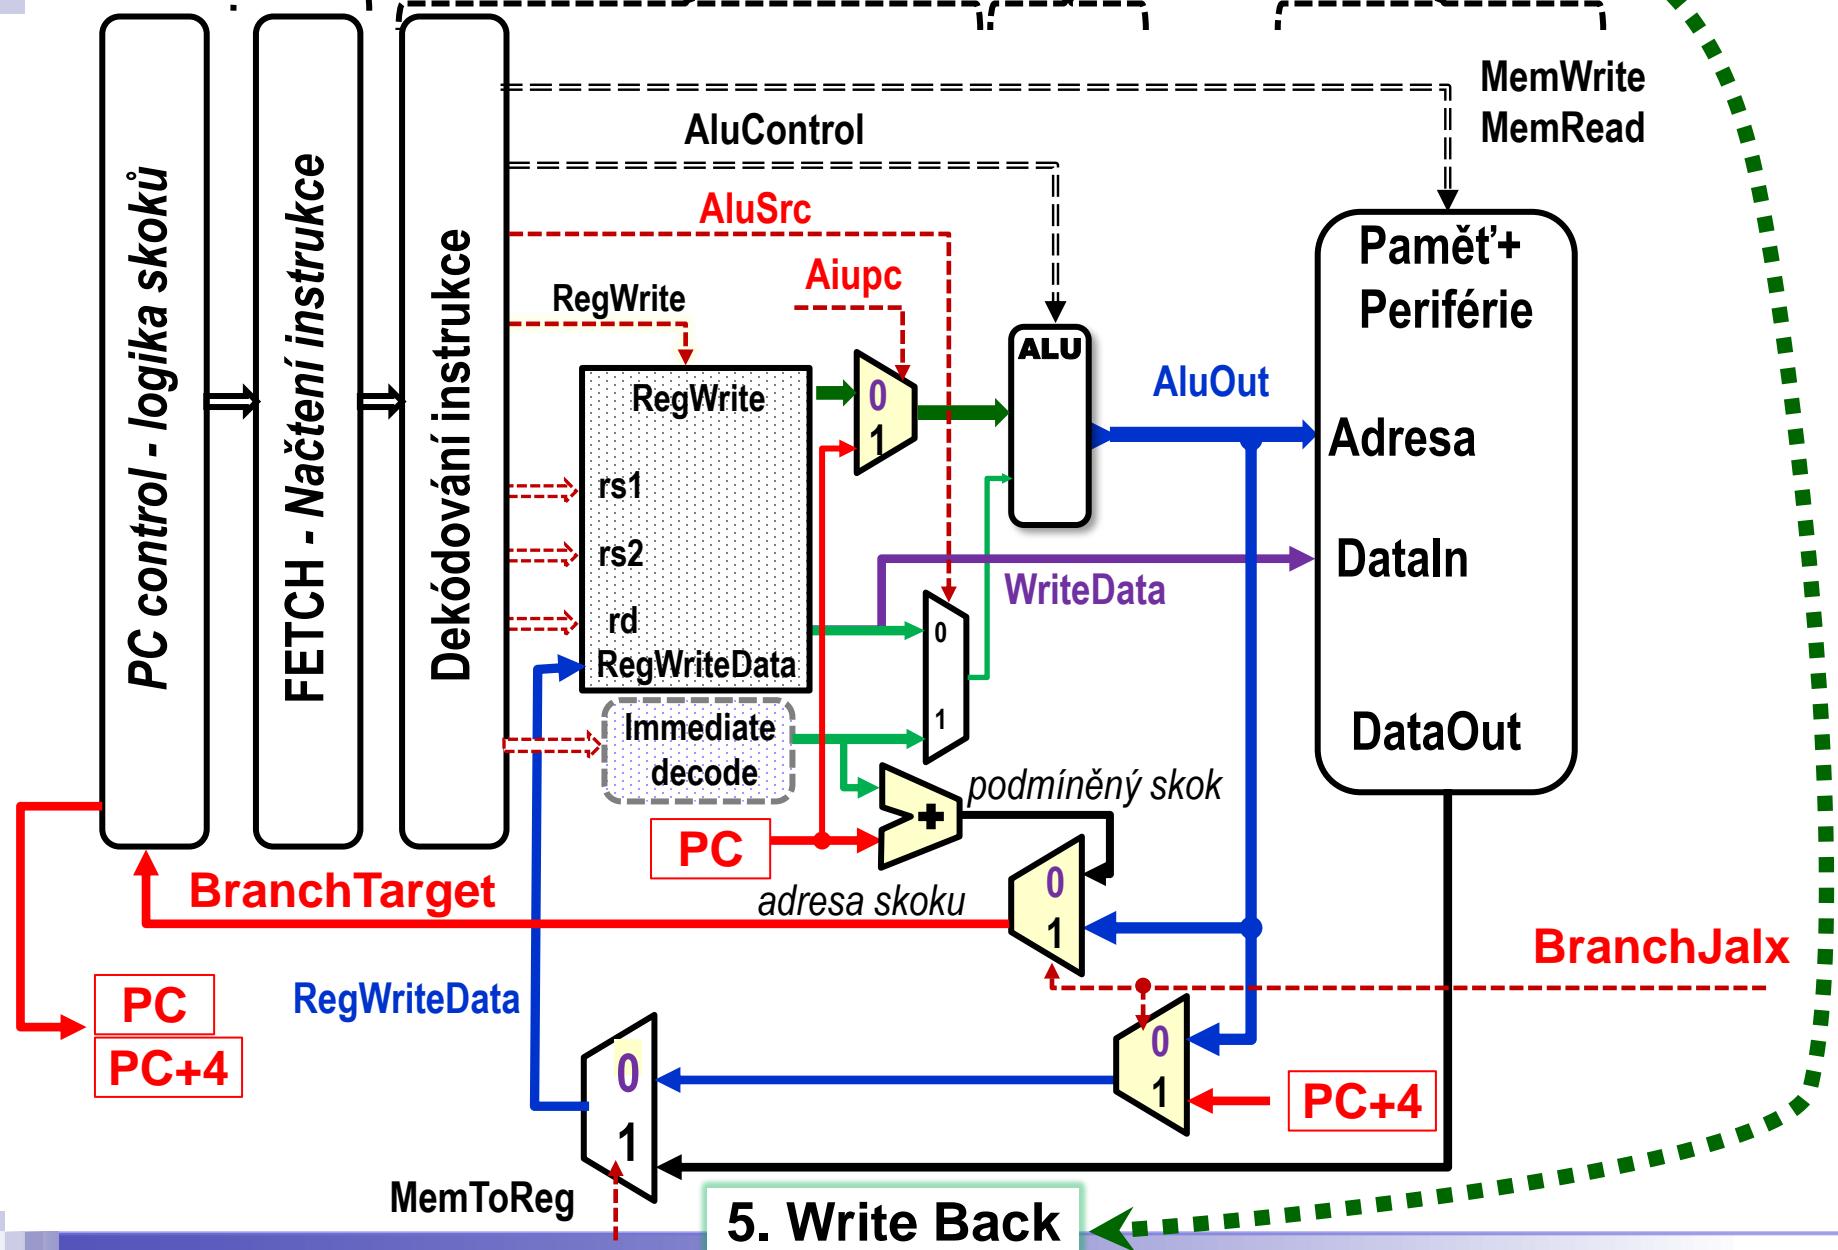
1. FETCH

2. DECODE

3. EXECUTE

4. MEMORY

Pipeline Stages



What the processor does in degrees

- **FETCH** - fetches an instruction from memory.
- **DECODE** - the instruction is decoded and the operand values are added to it.
- **EXECUTE** - ALU (Arithmetic Logic Unit) performs an operation with the values prepared in the DECODE phase.
- **MEMORY** - possible work with memory, writing or reading the address calculated by ALU, if there is a branch then writing the new address to PC (Program Counter).
- **WRITE-BACK** - values are saved in registers

Complete rabbit assembler model AD 1202

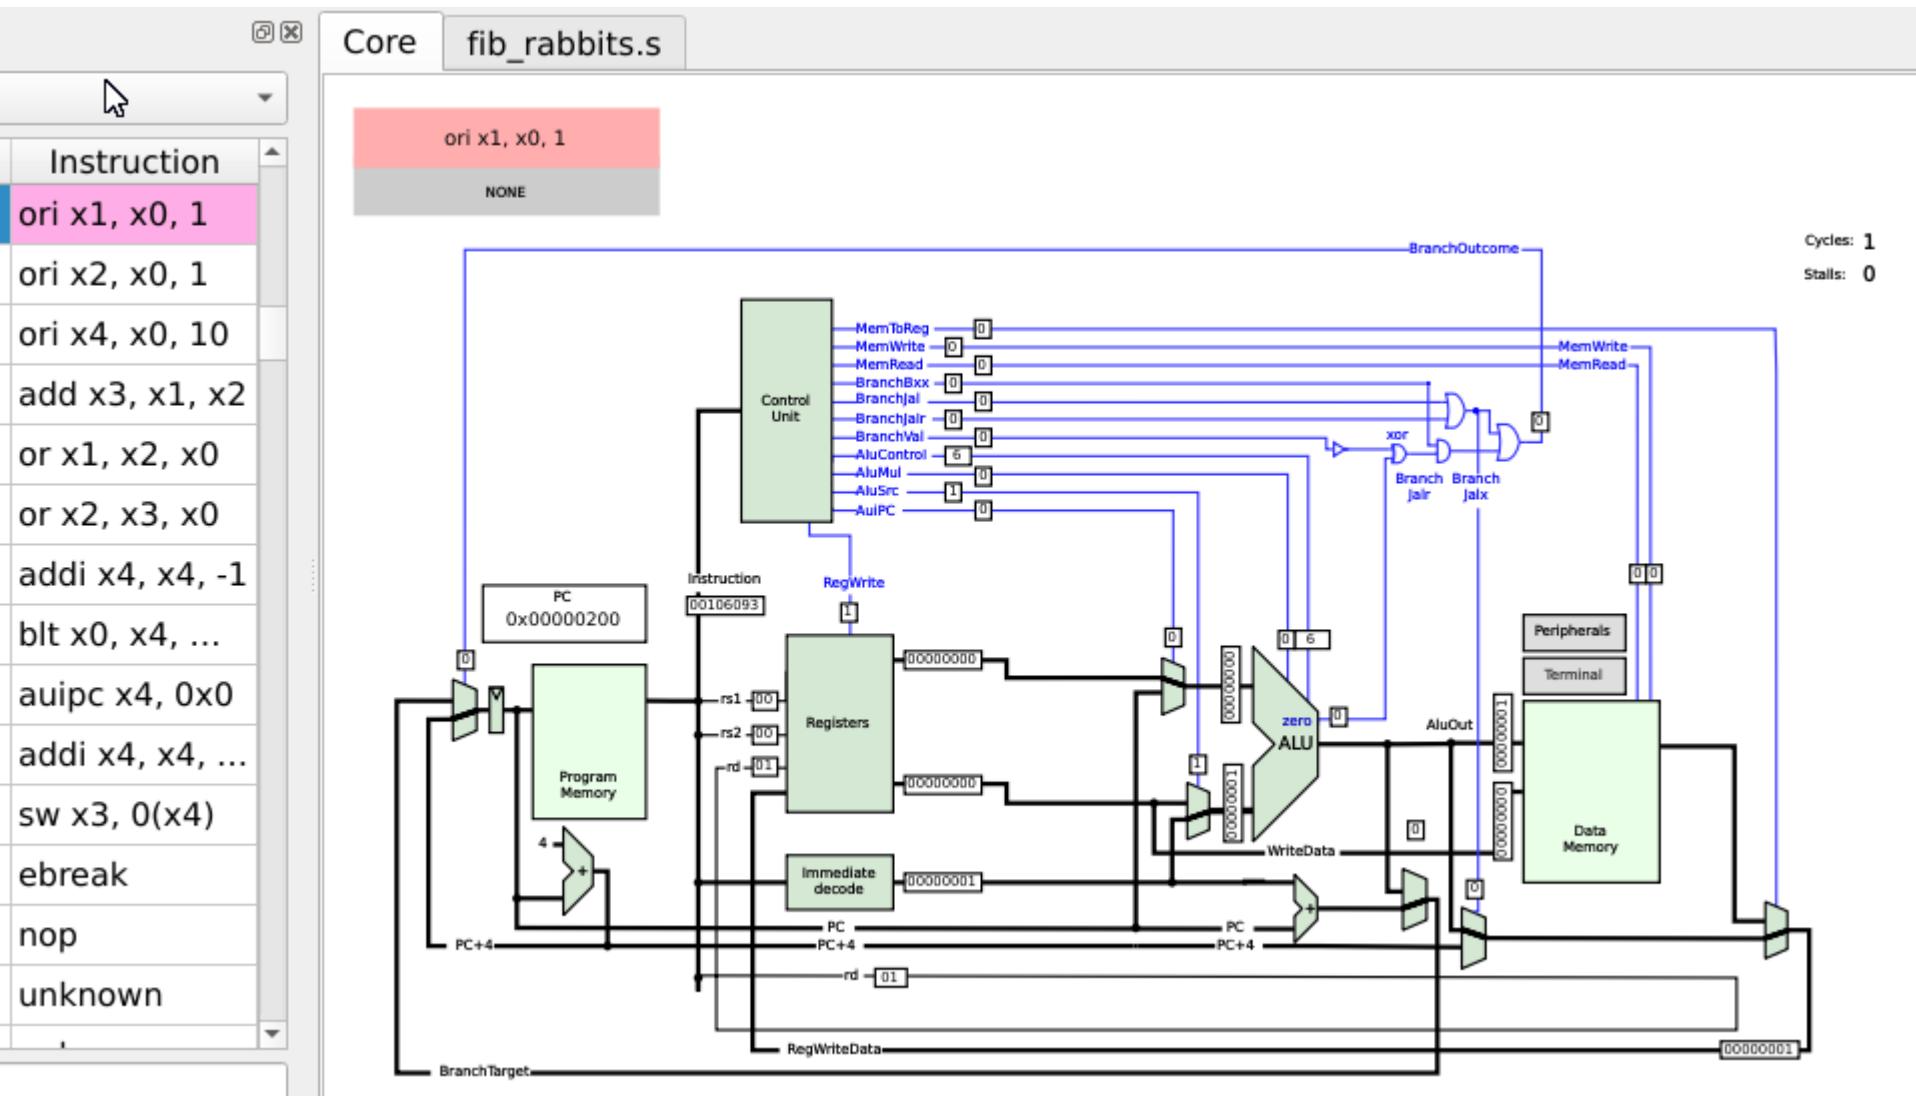
```
// Simulator directives to make interesting windows visible
#pragma qtmips show registers
.globl _start // entry point
.option norelax // no optimization of codes or instruction orders
.text // begin block of instructions
_start: ori x1, zero, 1 // x1←0 or 1
        ori x2, zero, 1 // x2←0 or 1
        ori x4, zero, 10 // x4←0 or 10
loop:
        add x3, x1, x2 // x3←x1 + x2
        or x1, x2, zero // x1←x2 or 0 ≡ x1 ← x2
        or x2, x3, zero // x2←x3 or 0 ≡ x2 ← x3
        addi x4, x4, -1 // x4←x3 + -1 ≡ x4--
        blt zero, x4, loop // if(0<x4) goto loop;
        la x4, fib13 // x4=&fib10
        sw x3, 0(x4) // *fib13=x3
stop: ebreak // break point
      nop
.data // begin block of data
.org 0x400 // start address of data segment
fib13: .word 0 // 32bit word initialized to 0
```

File
fib_rabbits.s

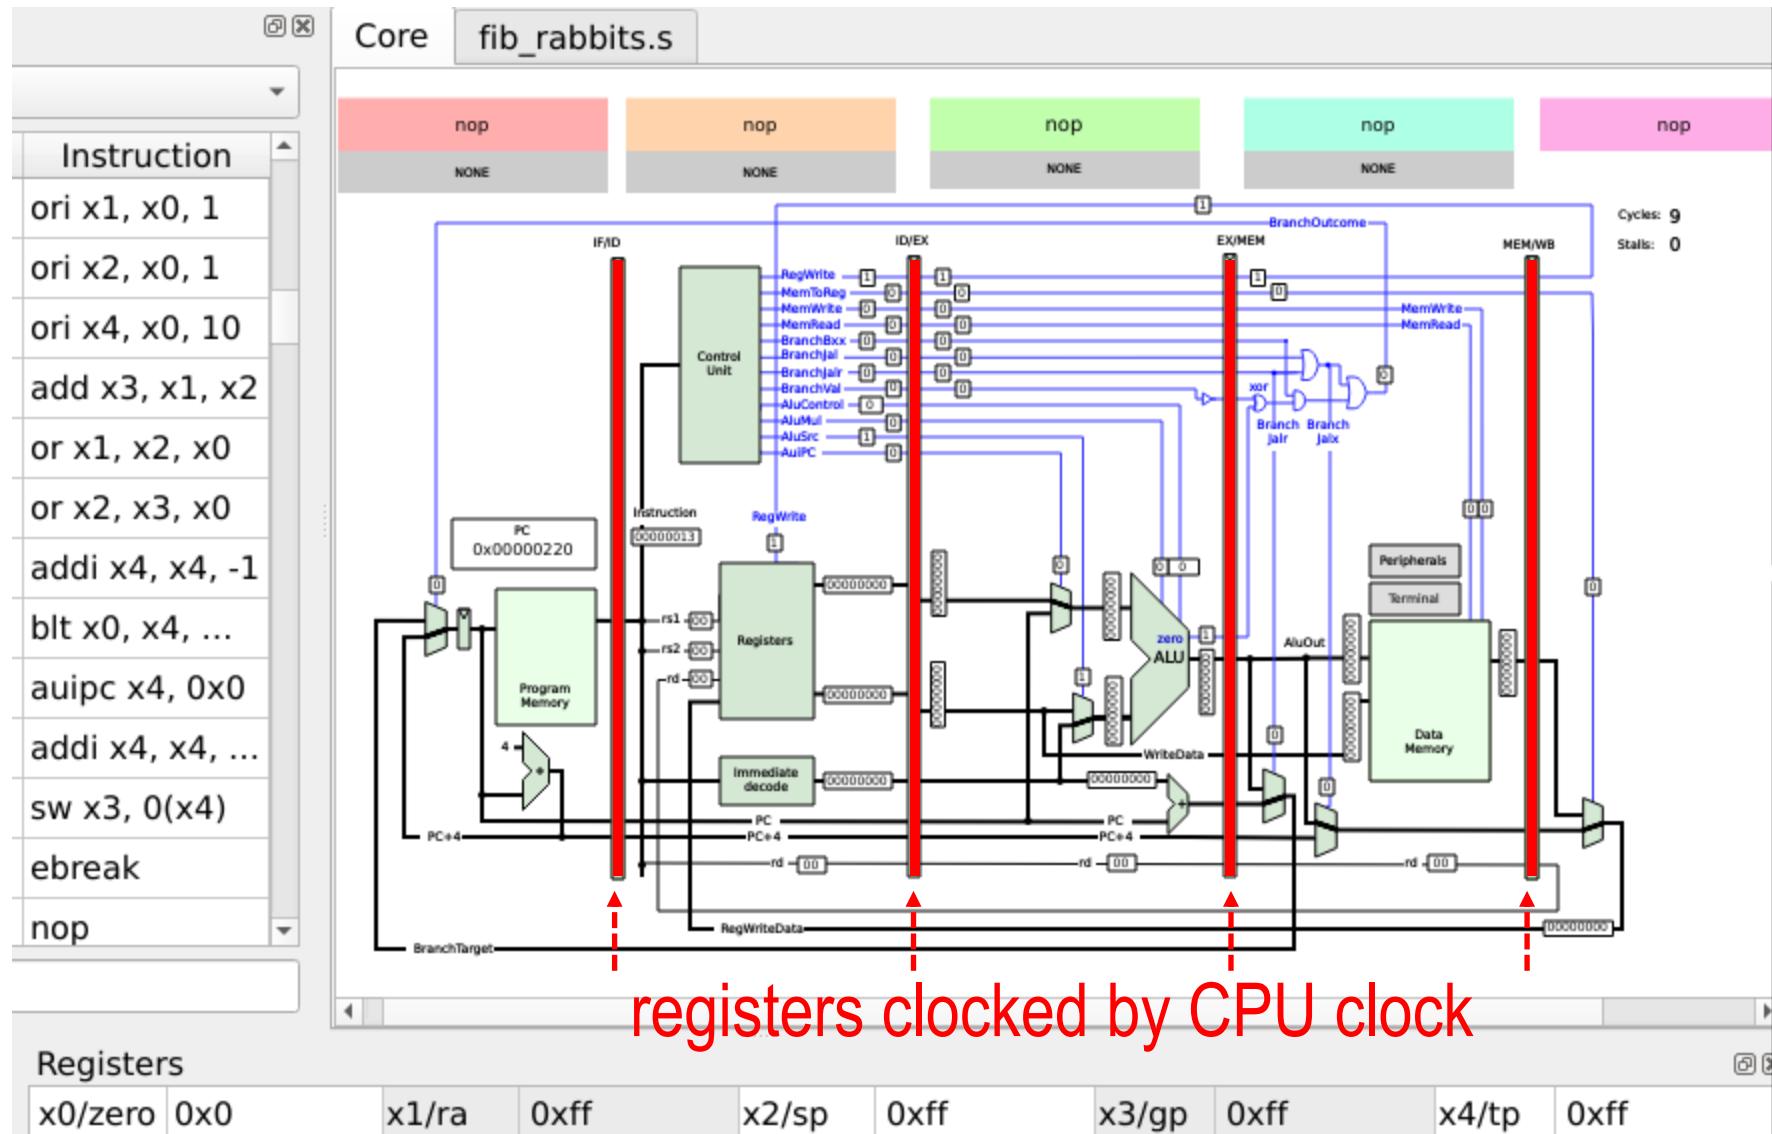
Registry	
zero	always 0
x1	step t-1
x2	step t
x3	step t+1
x4	counter

```
int x1, x2, x3, x4;
x1 = 1, x2 = 1; x4 = 10;
Loop: x3 = x1 + x2;
       x1 = x2; x2 = x3;
       x4--;
       if(0 < x4) goto Loop;
Fib13=x3; //&Fib13=0x400
```

Single cycle RISC V simulator



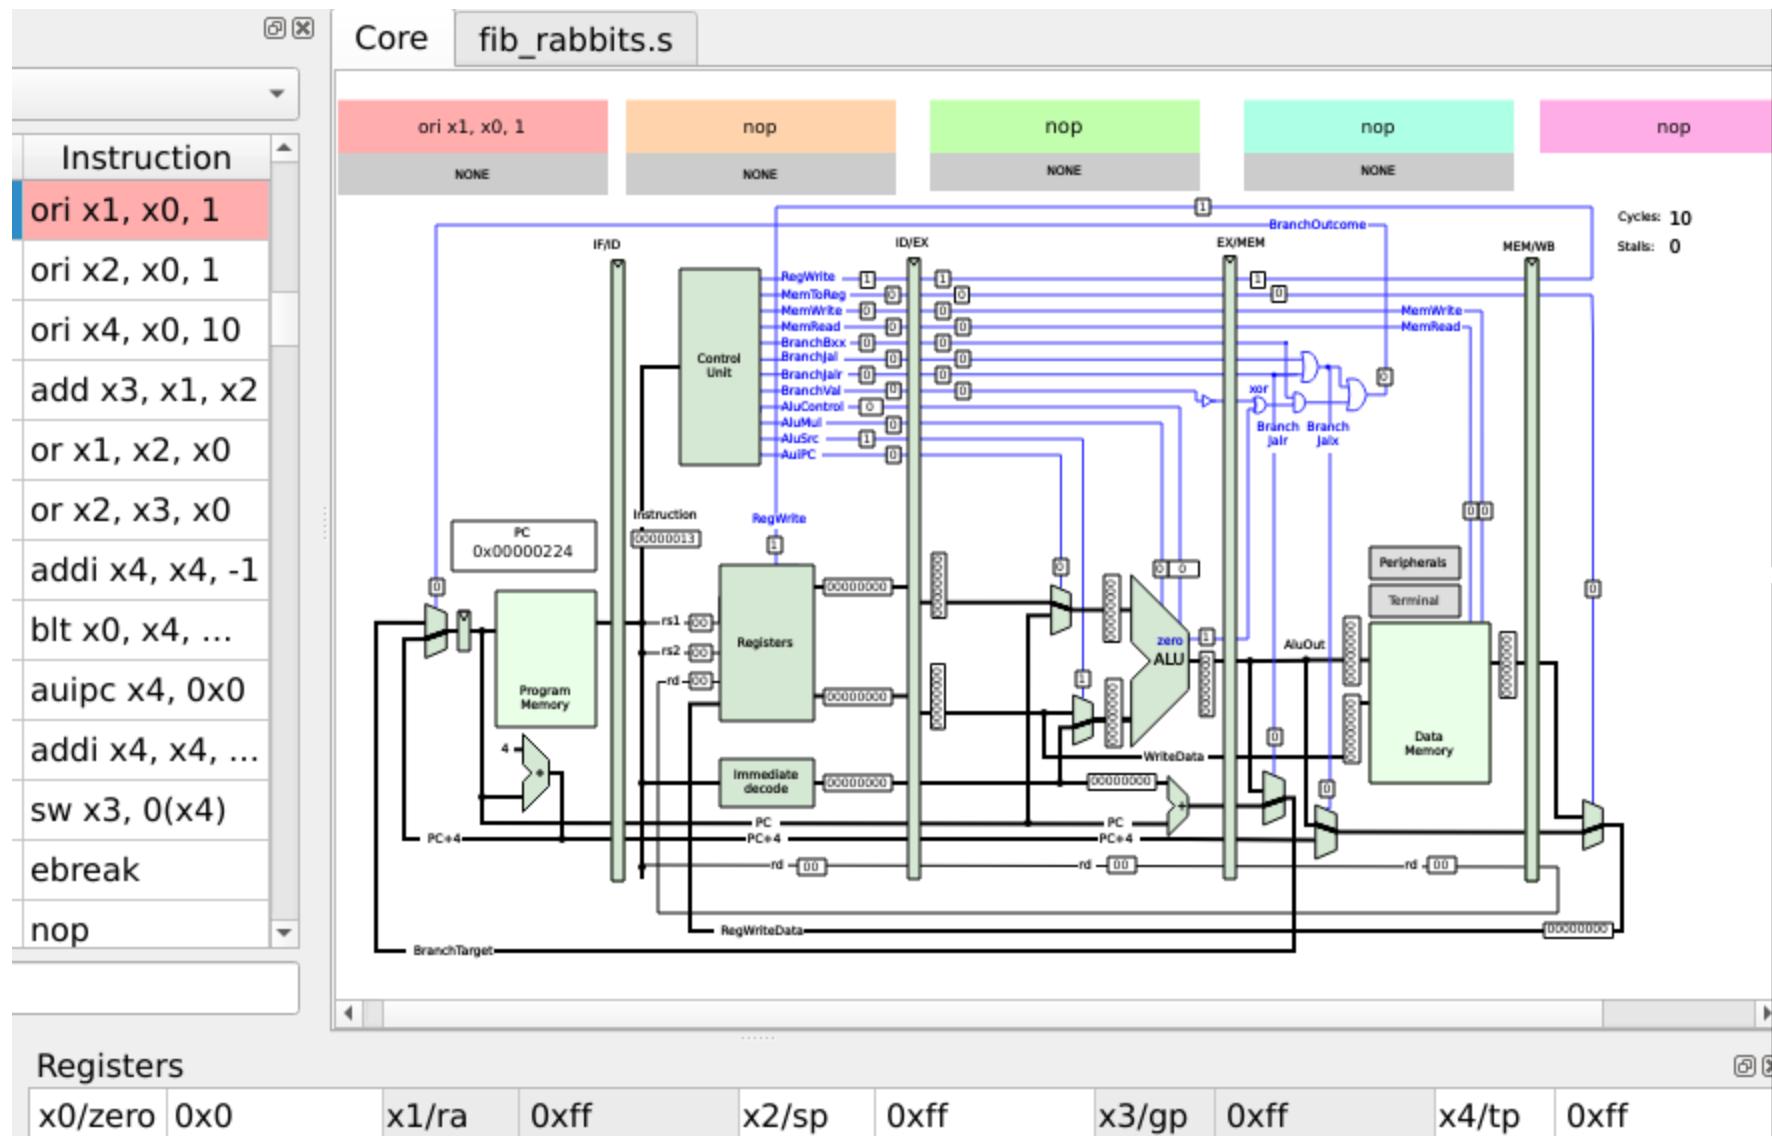
Pipeline Simulator RISC V



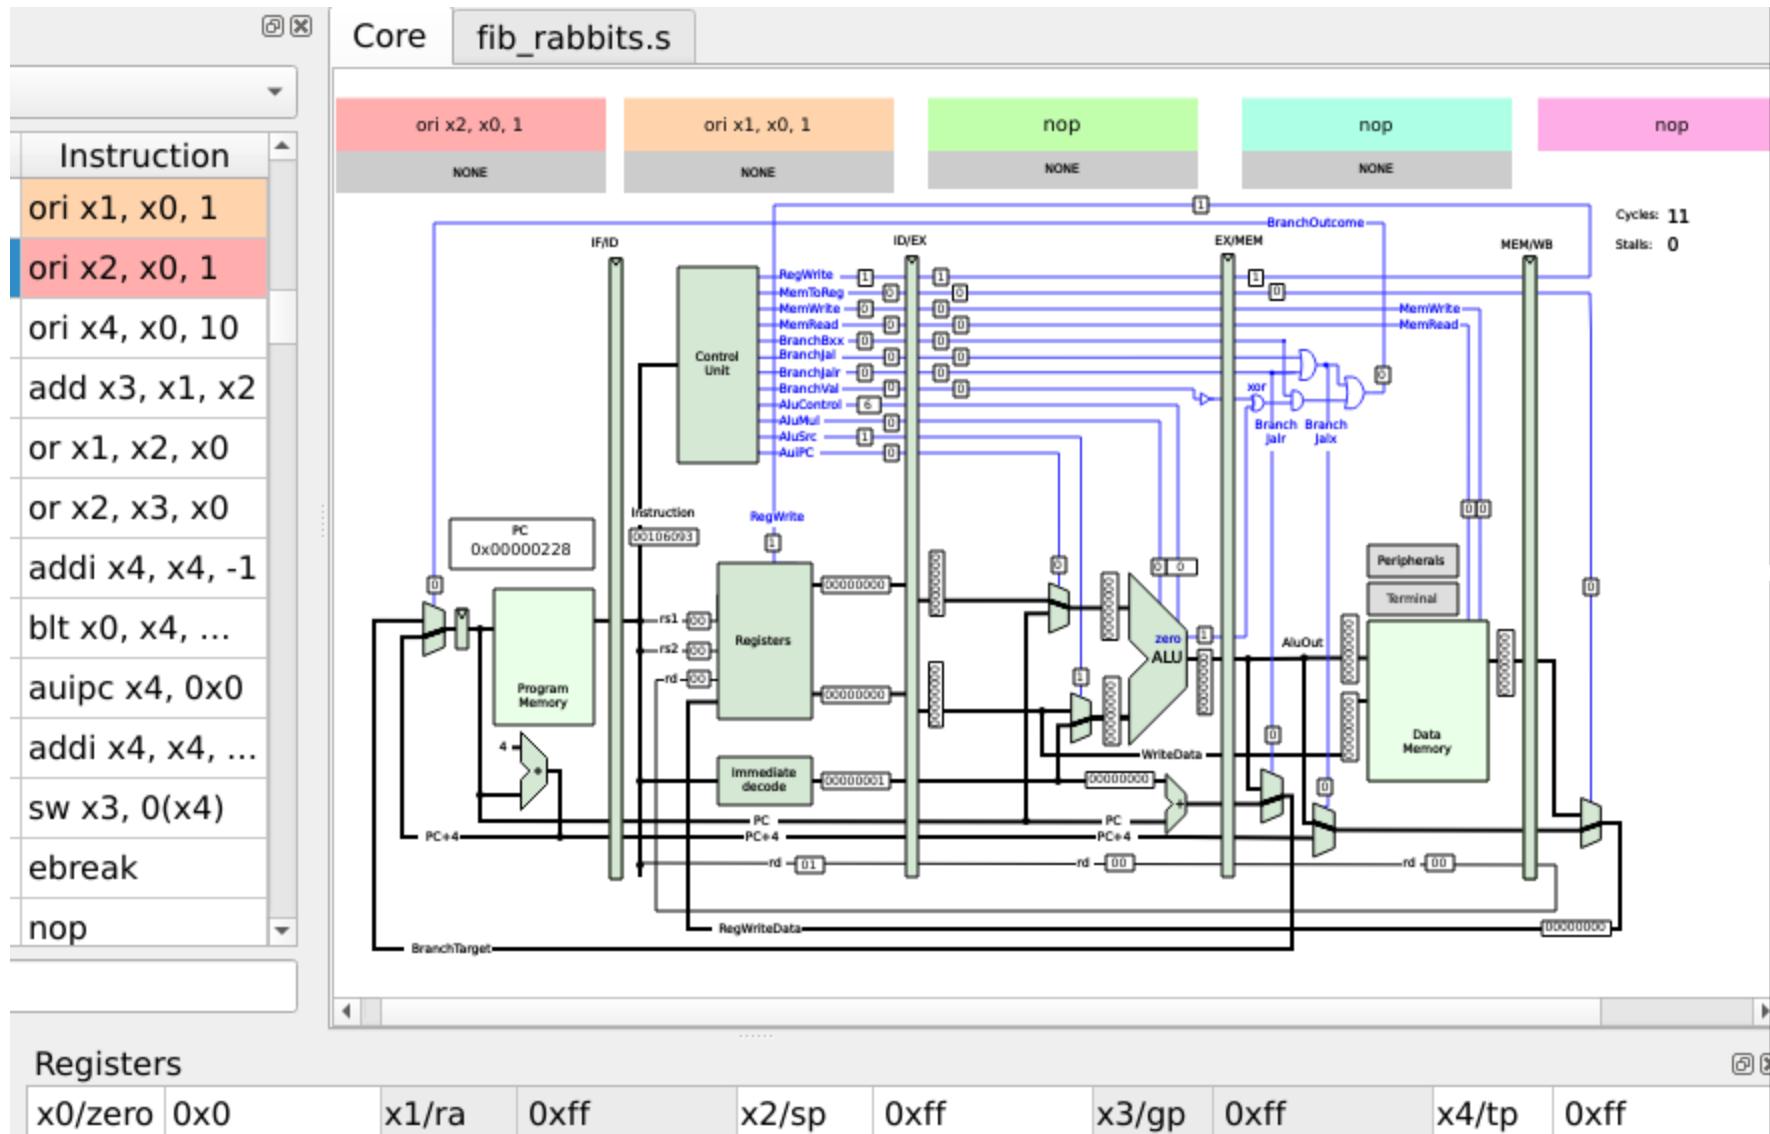
Data Hazards

x1	x2	x3	x4	Fetch	Decode	Execute	Memory	Write-Back
FF	FF	FF	FF	$x_1 \leftarrow 1$				
FF	FF	FF	FF	$x_2 \leftarrow 1$	$x_1 \leftarrow 1$			
FF	FF	FF	FF	$x_4 \leftarrow 10$	$x_2 \leftarrow 1$	$x_1 \leftarrow 1$		
FF	FF	FF	FF	$x_3 \leftarrow x_1 + x_2$	$x_4 \leftarrow 10$	$x_2 \leftarrow 1$	$x_1 \leftarrow 1$	
1	FF	FF	FF	$x_1 \leftarrow x_2$	$x_3 \leftarrow x_1 + x_2$	$x_4 \leftarrow 10$	$x_2 \leftarrow 1$	$x_1 \leftarrow 1$
1	1	FF	FF	$x_2 \leftarrow x_3$	$x_1 \leftarrow x_2$	$x_3 \leftarrow x_1 + x_2$	$x_4 \leftarrow 10$	$x_2 \leftarrow 1$
1	1	FF	10	$x_4 \leftarrow x_4 - 1$	$x_2 \leftarrow x_3$	$x_1 \leftarrow x_2$	$x_3 \leftarrow x_1 + x_2$	$x_4 \leftarrow 10$
1	1	100	10	loop $x_4 < x_0$	$x_4 \leftarrow x_4 - 1$	$x_2 \leftarrow x_3$	$x_1 \leftarrow x_2$	$x_3 \leftarrow x_1 + x_2$

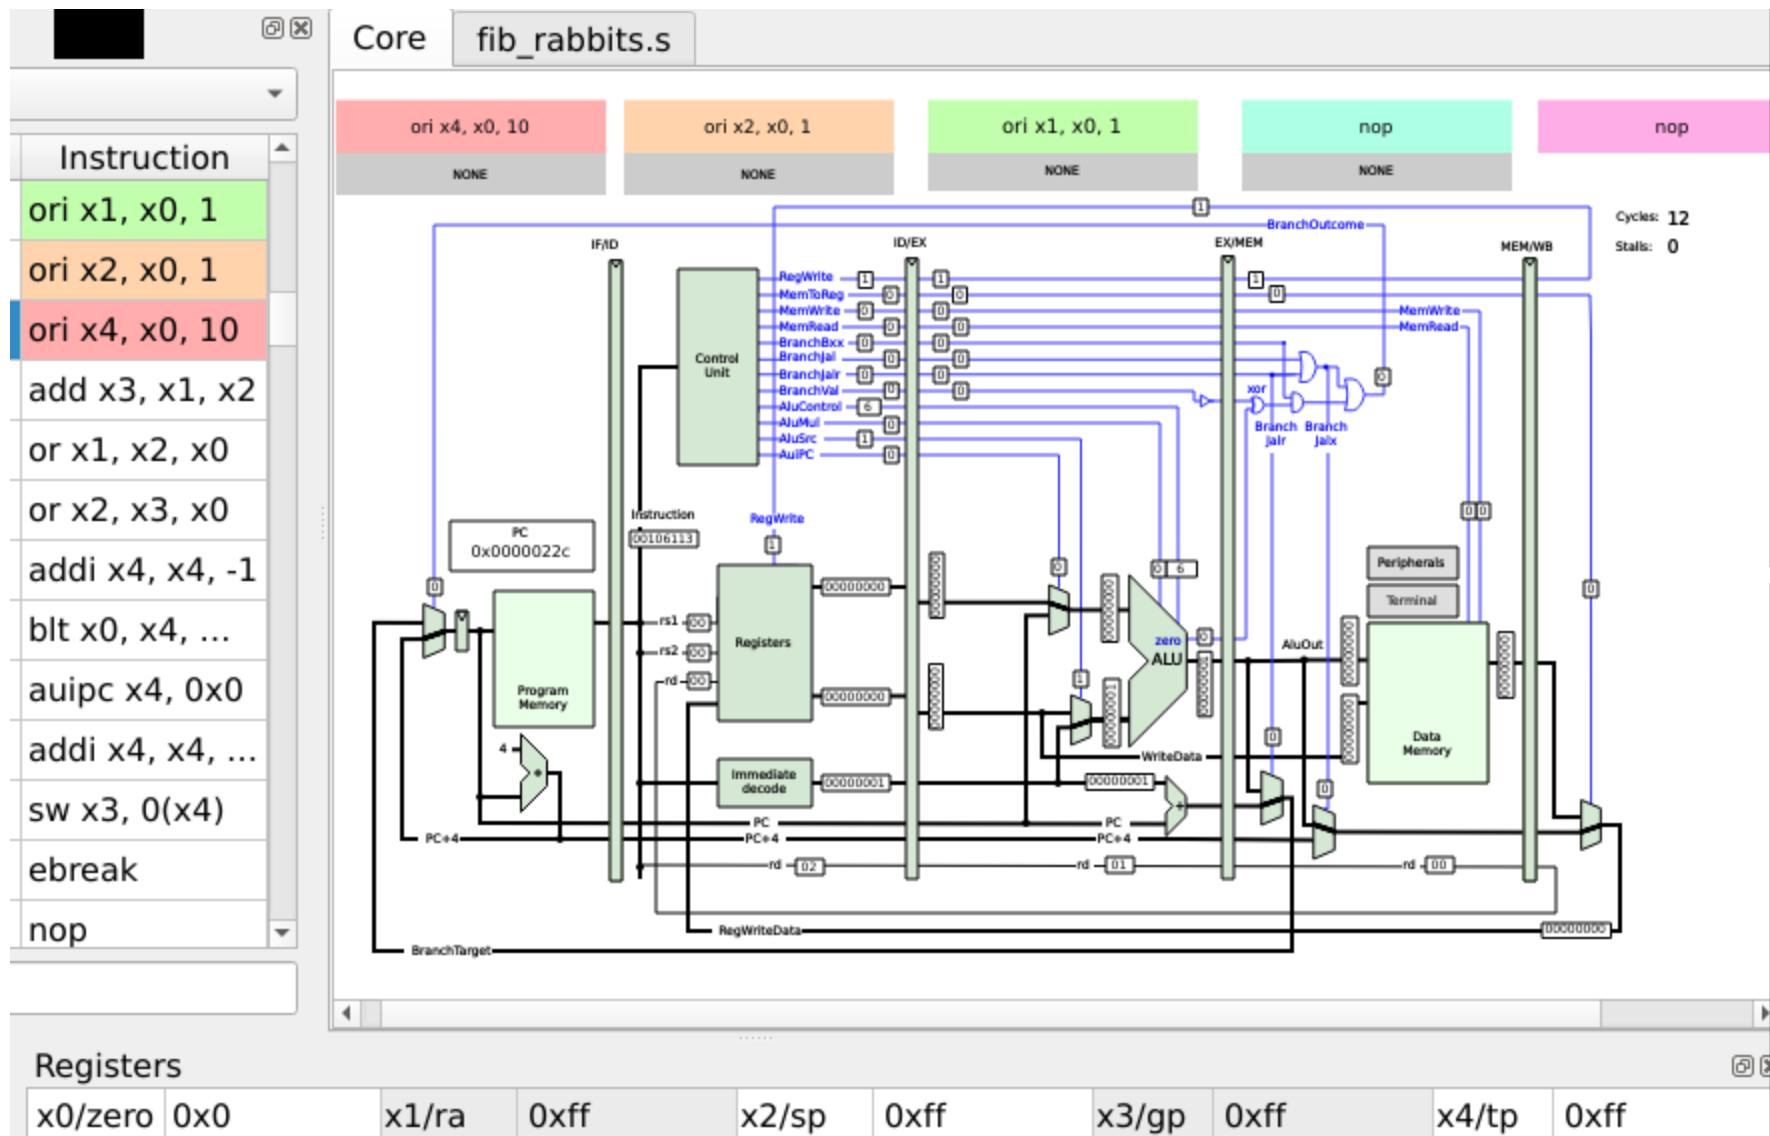
1st instructions in the pipeline



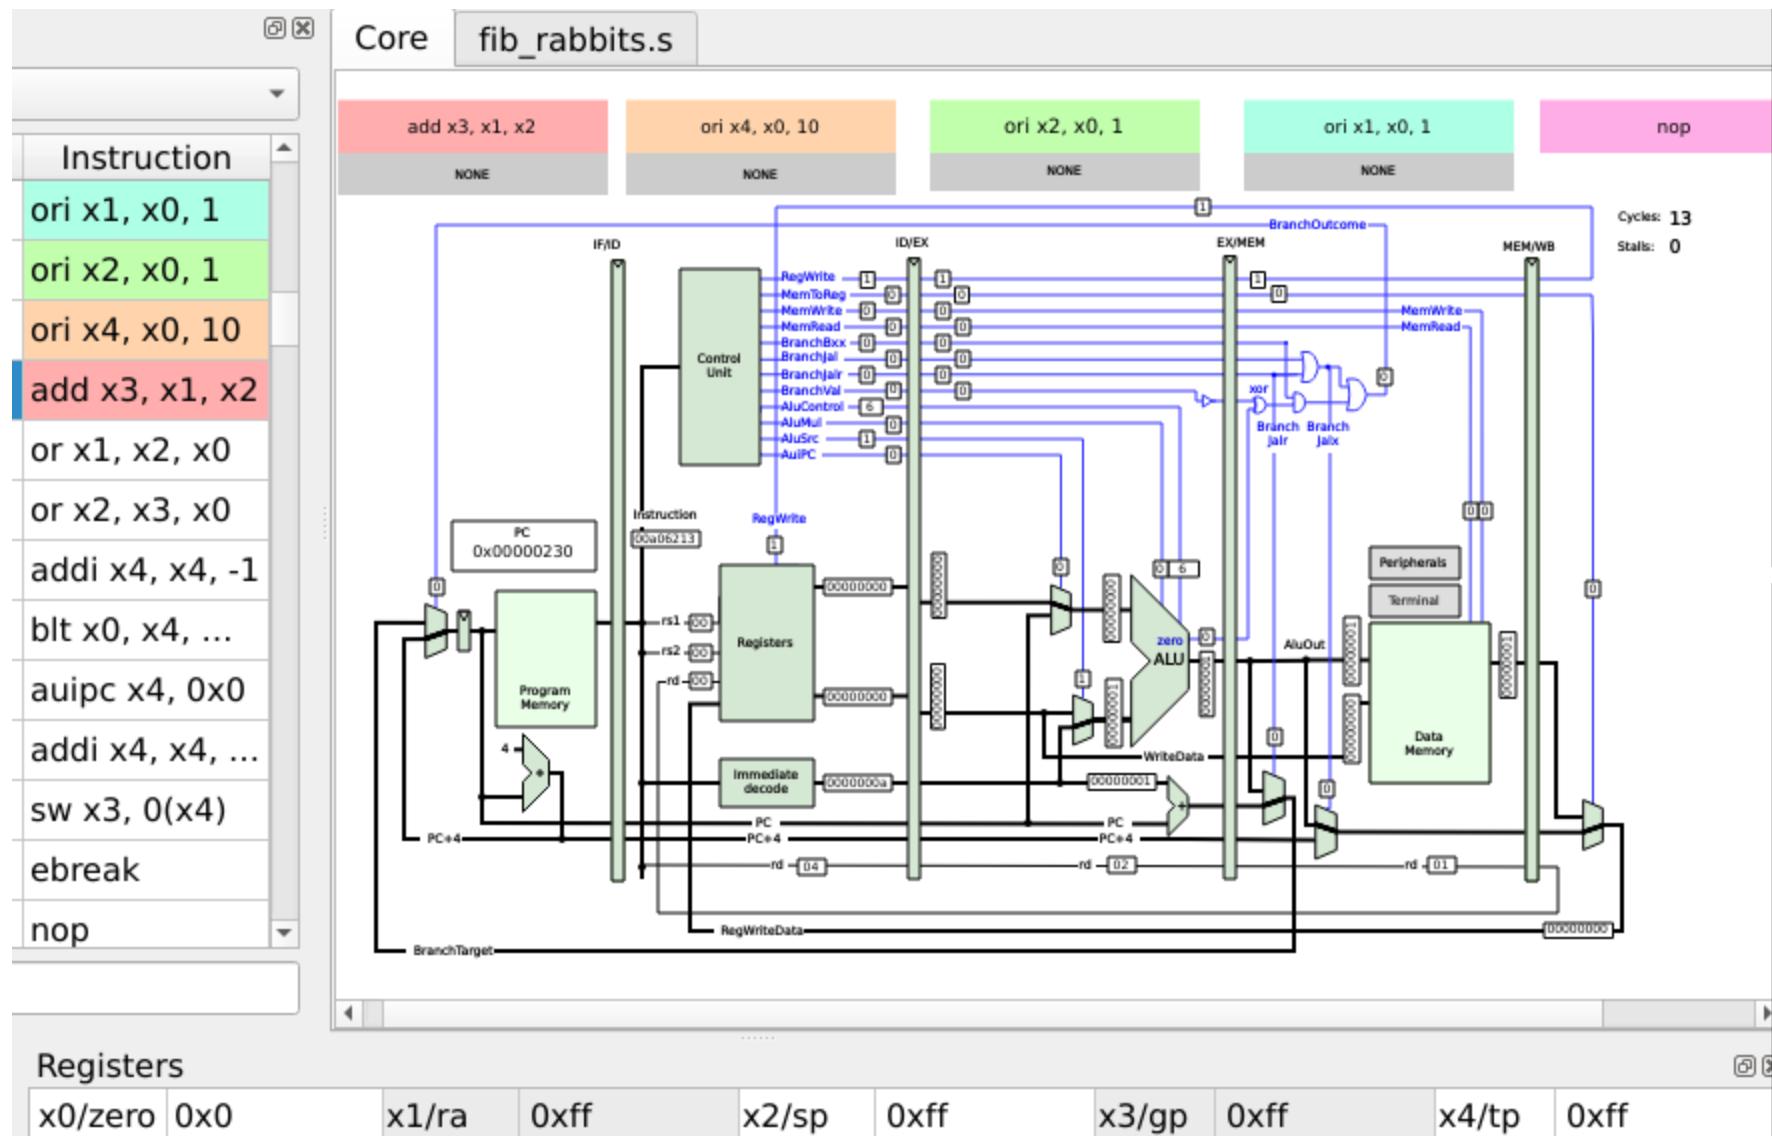
2nd instructions in the pipeline



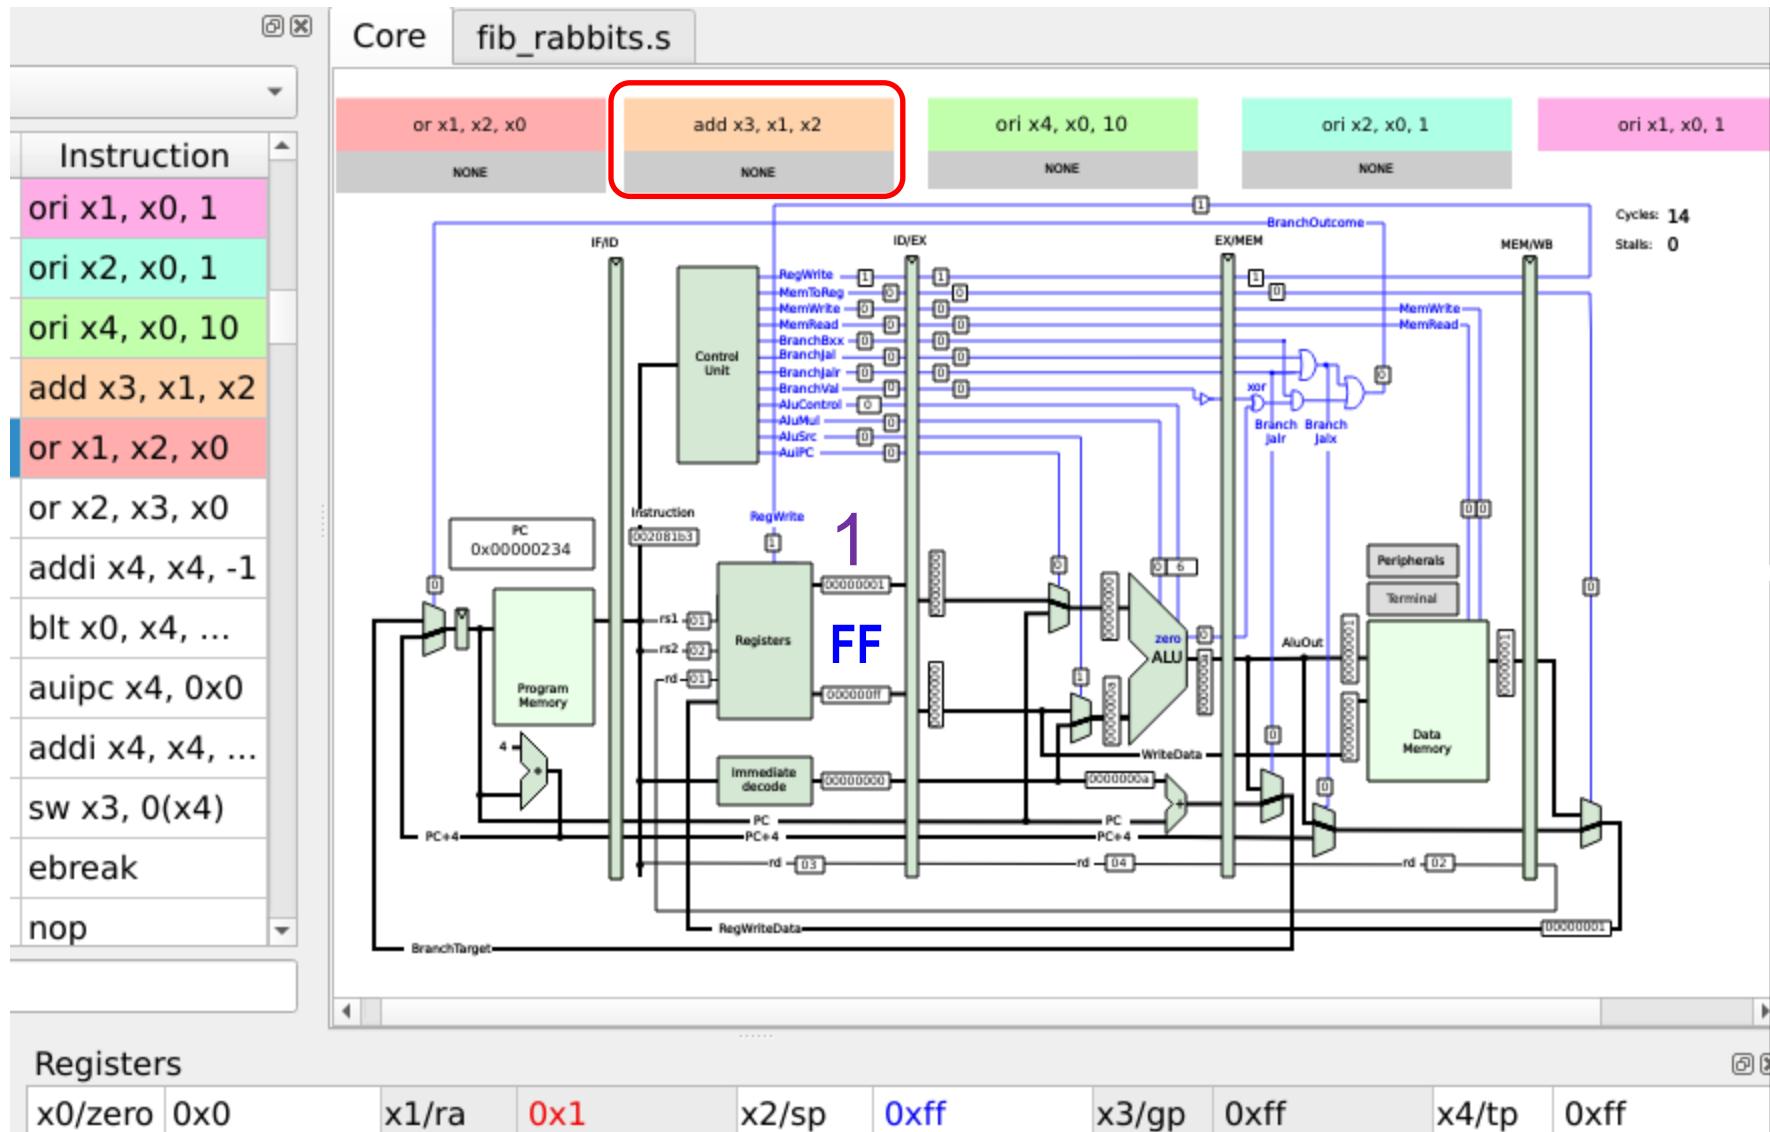
3rd instructions in the pipeline



4th instructions in the pipeline

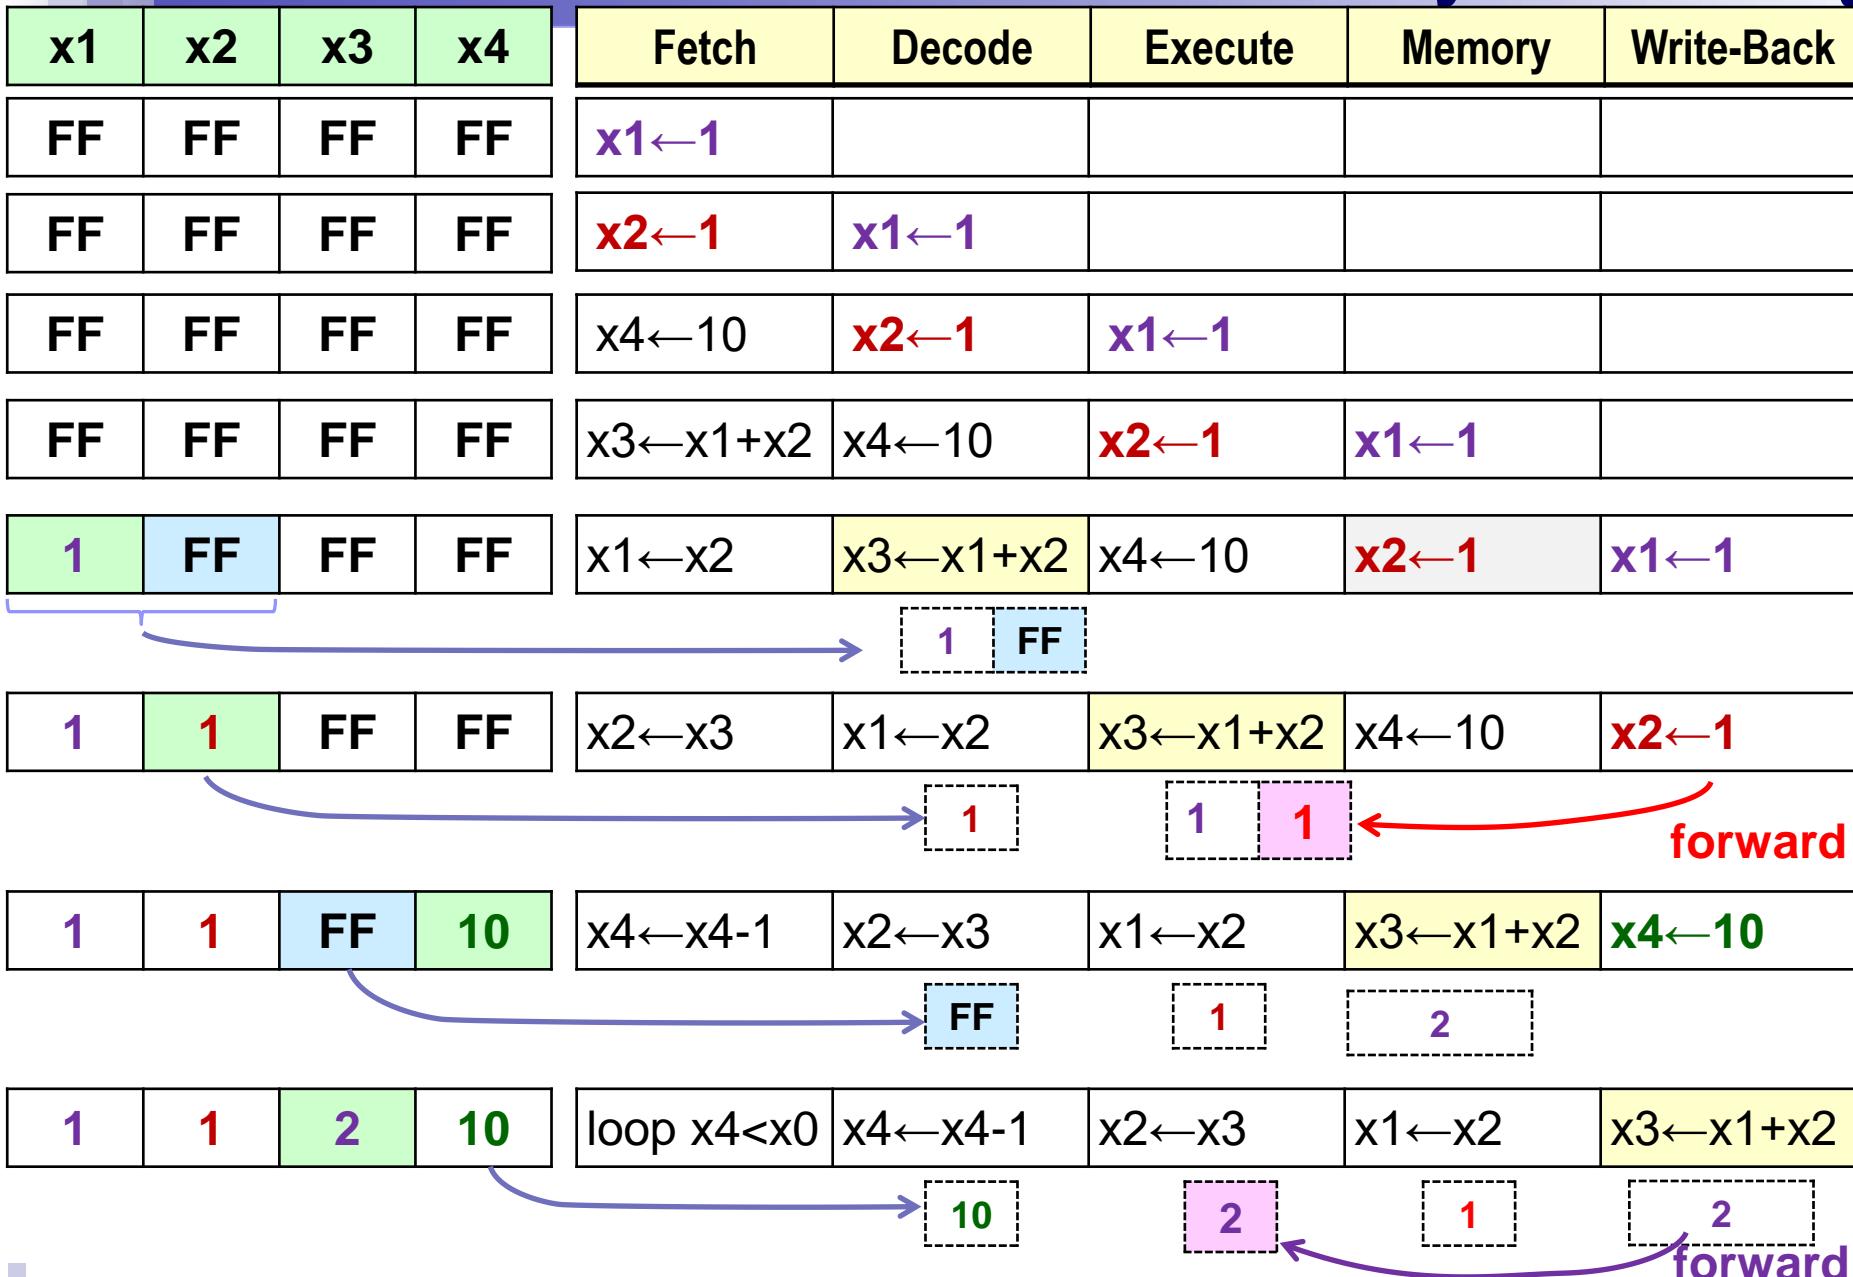


5th instructions in the pipeline

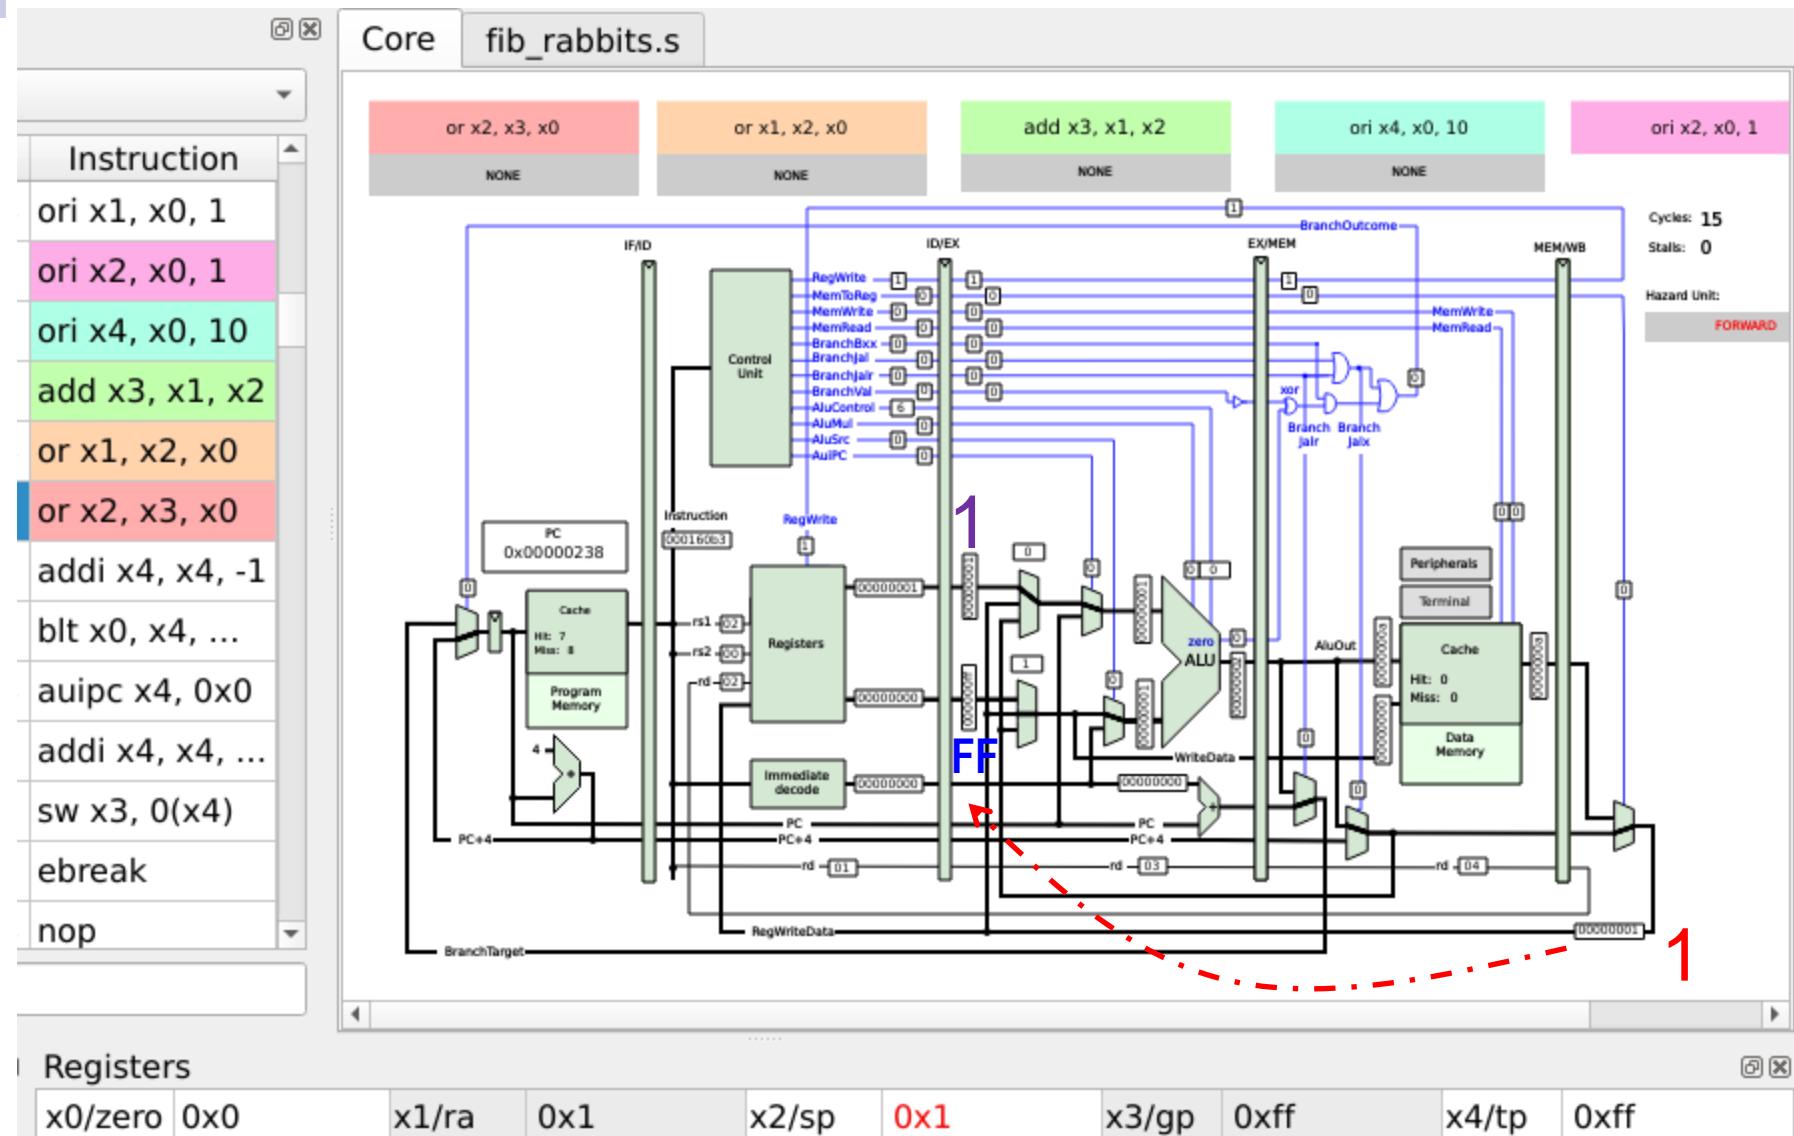


The add instruction got the operands in DECODE according to the immediate values in the registers

Treatment of data hazards by forwarding

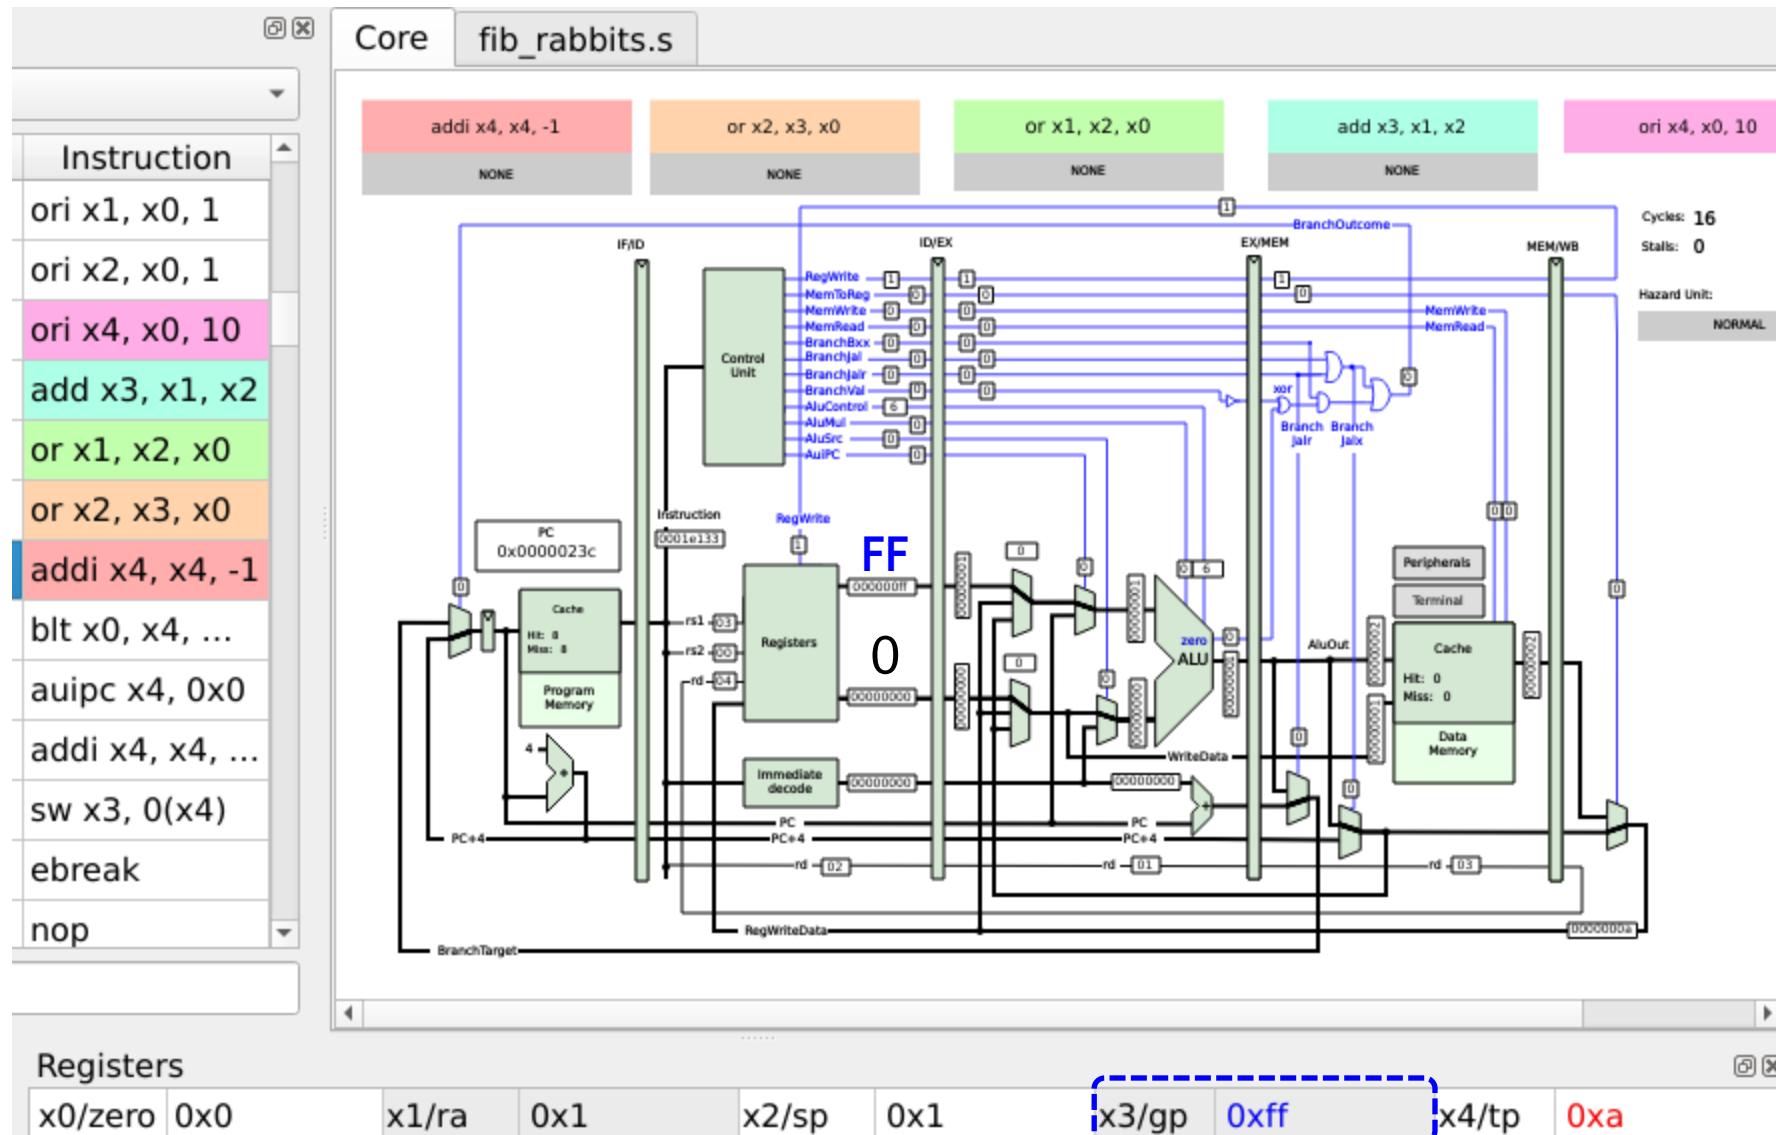


6th instruction + input correction add x3,x1,x2



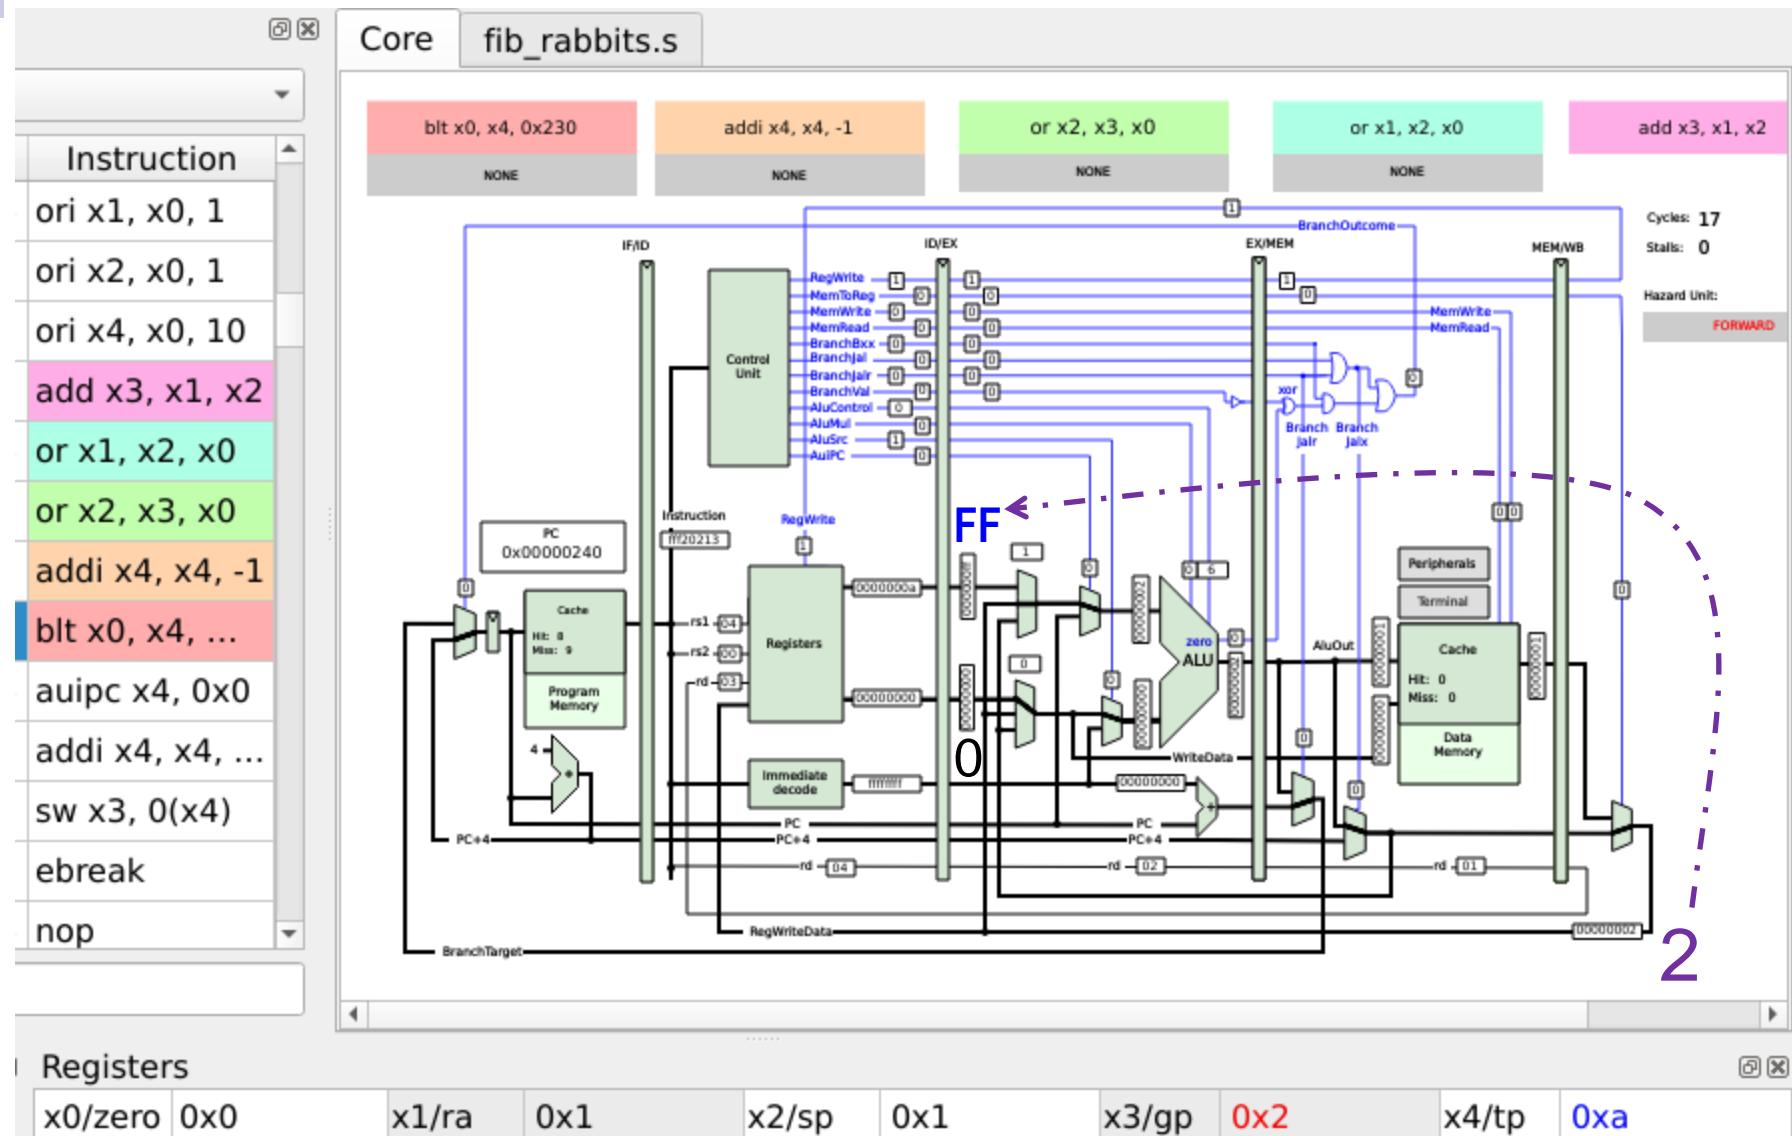
x1 value is sent from WR-BACK to the EXE stage and overwrites 0xFF

7th instruction



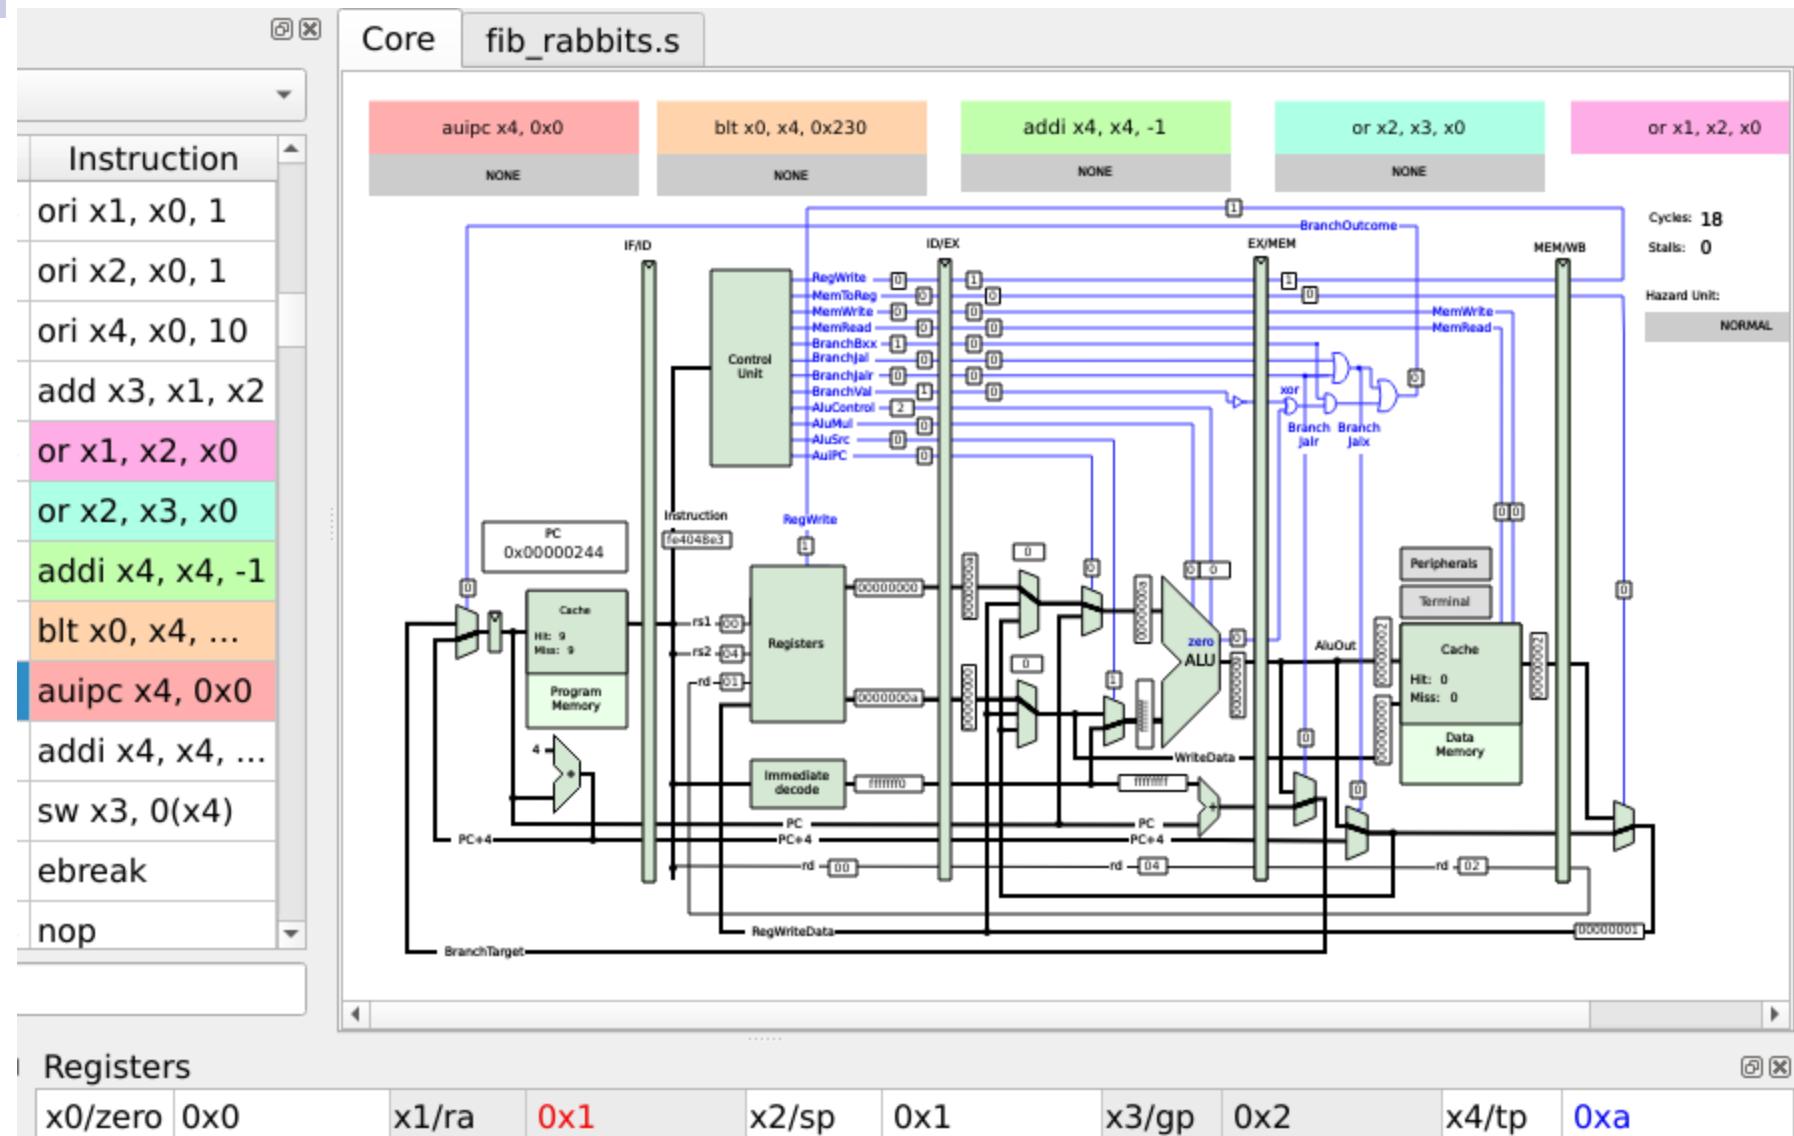
the operands were loaded again from the immediate values in the registers

8th instruction



x3 value is sent from WR-BACK to the EXE stage and overwrites 0xFF

9th instruction + correct result or x2,x3,x0



Not only from Write-Back is transferred to the EXE stage, but also from Memory, which occurs with a pair of instructions:

la x4, fib13 // x4=&fib13

sw x3, 0(x4) // *fib13=x3

translated to

auipc x4, 0x0 // $x4 \leftarrow PC + 0$

addi x4, x4, 480 // $x4 \leftarrow x4 + 0x1E0$

sw x3, 0(x4)

Here the value x4 is transferred twice from MEM to EXE.

When do we care about data hazards?

- We don't have to if we are translating for a large processor equipped with Hazard Unit.
- However, we will treat them on smaller processors without a Hazard Unit, where the code is often corrected by the compiler:

_start: ori x1, zero, 1 // $x_1 \leftarrow 0$ or 1

 ori x2, zero, 1 // $x_2 \leftarrow 0$ or 1

 ori x4, zero, 10 // $x_4 \leftarrow 0$ or 10

nop // wait for x2

loop: add x3, x1, x2 // $x_3 \leftarrow x_1 + x_2$

 or x1, x2, zero // $x_1 \leftarrow x_2$ or 0 $\equiv x_1 \leftarrow x_2$

nop // wait for x3

 or x2, x3, zero // $x_2 \leftarrow x_3$ or 0 $\equiv x_2 \leftarrow x_3$

 addi x4, x4, -1 // $x_4 \leftarrow x_3 + -1 \equiv x_4 --$

 blt zero, x4, loop // if($0 < x_4$) goto loop;

 auipc x4, 0x0 // $x_4 \leftarrow PC + 0$; *after adding NOP, PC=0x204*

nop // wait for x4

nop // wait for x4

 addi x4, x4, 476 // $x_4 \leftarrow x_4 + 0x1DC$

nop // wait for x4

nop // wait for x4

 sw x3, 0(x4) // $*x_4 \leftarrow x_3$

When do we care about data hazards?

_start: ori x1, zero, 1 // $x_1 \leftarrow 0$ or 1

ori x2, zero, 1 // $x_2 \leftarrow 0$ or 1

ori x4, zero, 10 // $x_4 \leftarrow 0$ or 10

nop // wait for x2

loop: add x3, x1, x2 // $x_3 \leftarrow x_1 + x_2$ // loop optimized by our instruction reorder

addi x4, x4, -1 // $x_4 \leftarrow x_3 + -1 \equiv x_4 --$

or x1, x2, zero // $x_1 \leftarrow x_2$ or 0 $\equiv x_1 \leftarrow x_2$

or x2, x3, zero // $x_2 \leftarrow x_3$ or 0 $\equiv x_2 \leftarrow x_3$

blt zero, x4, loop // if($0 < x_4$) goto loop;

auipc x4, 0x0 // $x_4 \leftarrow \text{PC} + 0$; after adding NOP, PC=204

nop // wait for x4

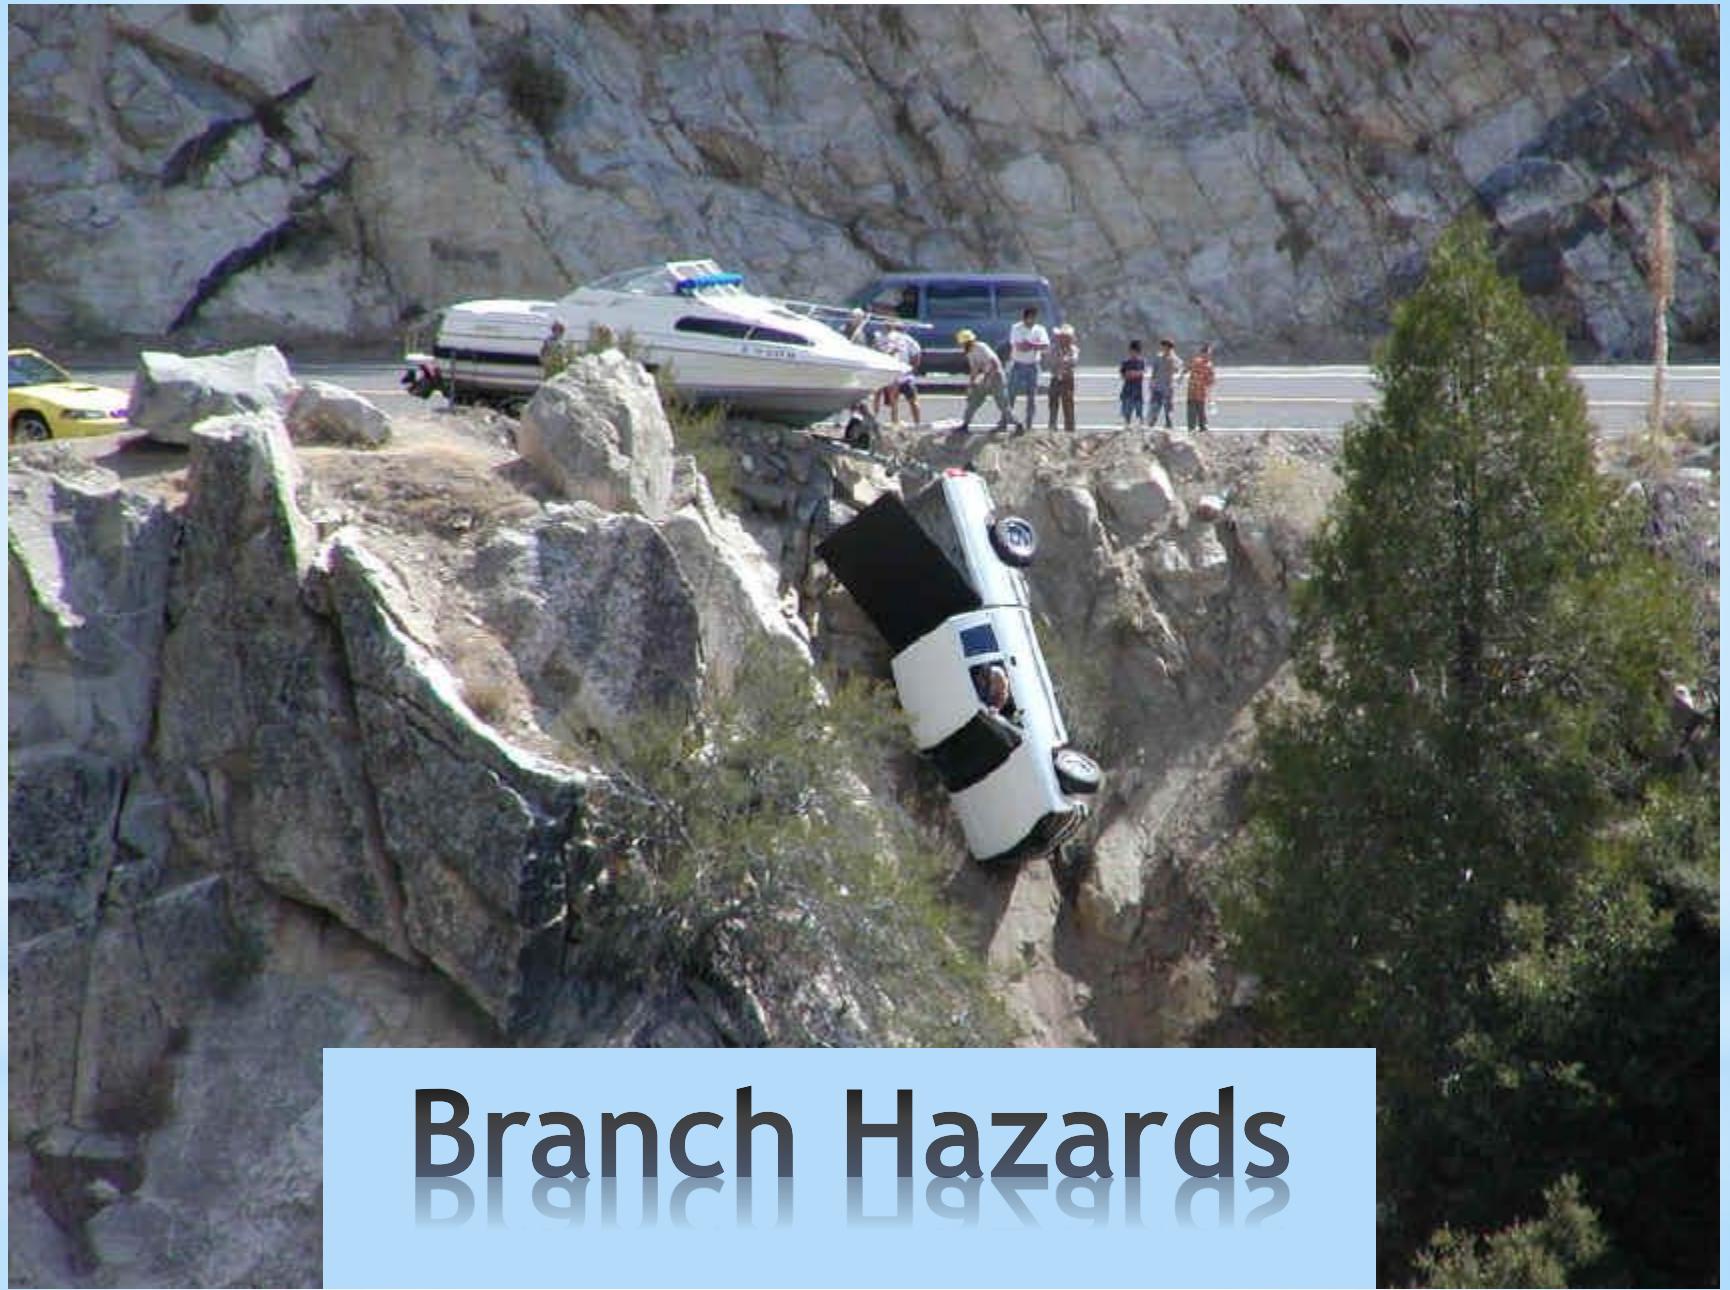
nop // wait for x4

addi x4, x4, 476 // $x_4 \leftarrow x_4 + 0x1DC$

nop // wait for x4

nop // wait for x4

sw x3, 0(x4) // Fib13=x4; Fib13 address = 0x400



Branch Hazards

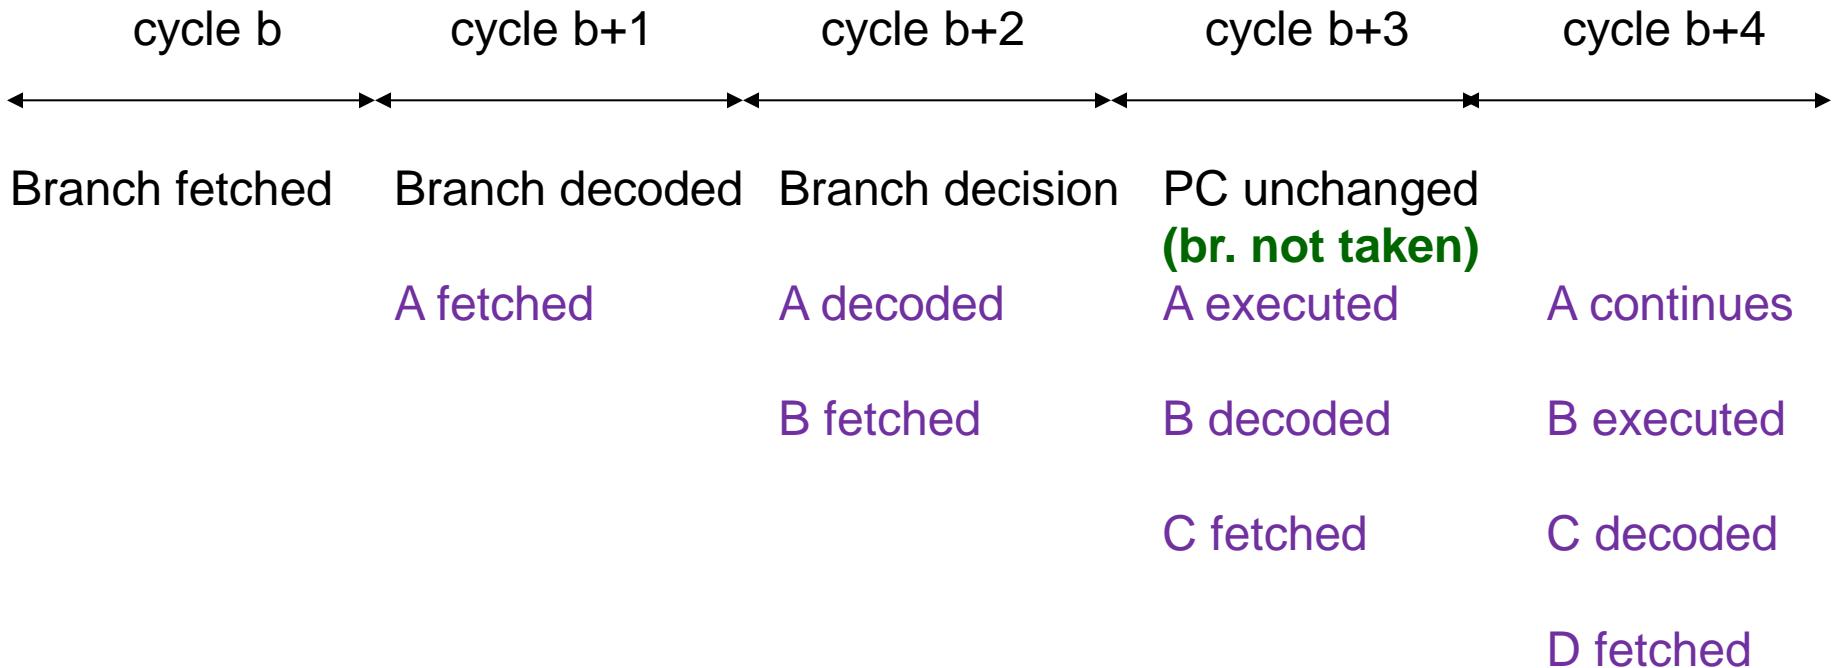
Source: <http://www.spaceg.com/multimedia/collection/automobiles/> (Aug 20, 2014)

Branch Not Taken

Branch to Z
A
B
C
D
Z

Assume the heuristic: no branching.

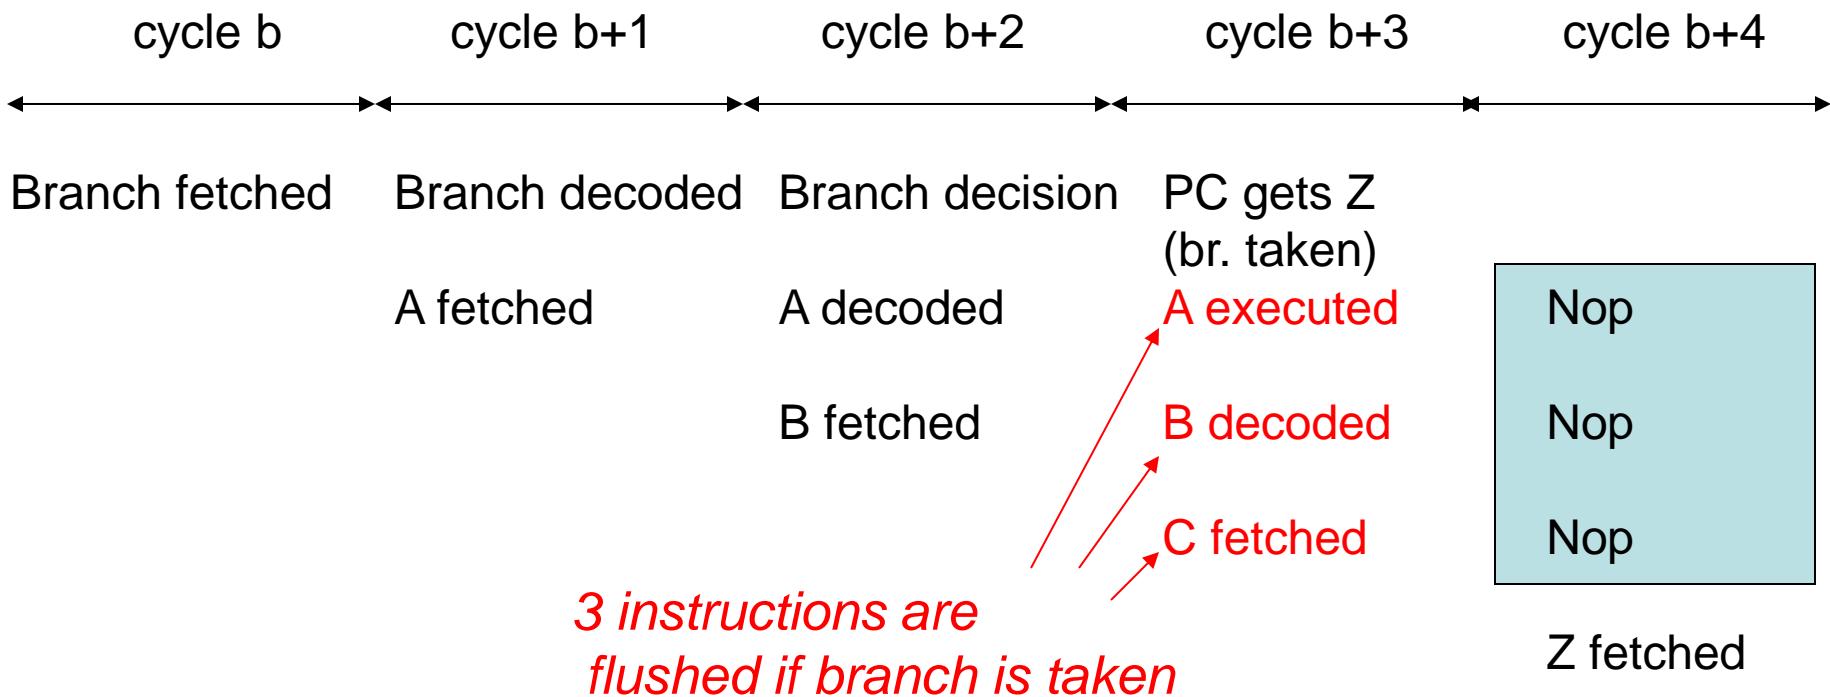
The instructions continue to be loaded in their order



Branch Taken

Branch to Z
A
B
C
D
Z

If branching, branching is generated in the MEM stage.
All work in progress is deleted and a new instruction is loaded.



Branch Hazard

Core fib_rabbits.s we already know we're gonna branch

Instruction

- ori x1, x0, 1
- ori x2, x0, 1
- ori x4, x0, 10
- add x3, x1, x2
- or x1, x2, x0
- or x2, x3, x0
- addi x4, x4, -1
- blt x0, x4, 0x230**
- auipc x4, 0x0
- addi x4, x4, ...
- sw x3, 0(x4)
- ebreak
- nop

Cycles: 19
Stalls: 0
Hazard Unit: FORWARD

0=condition was met.

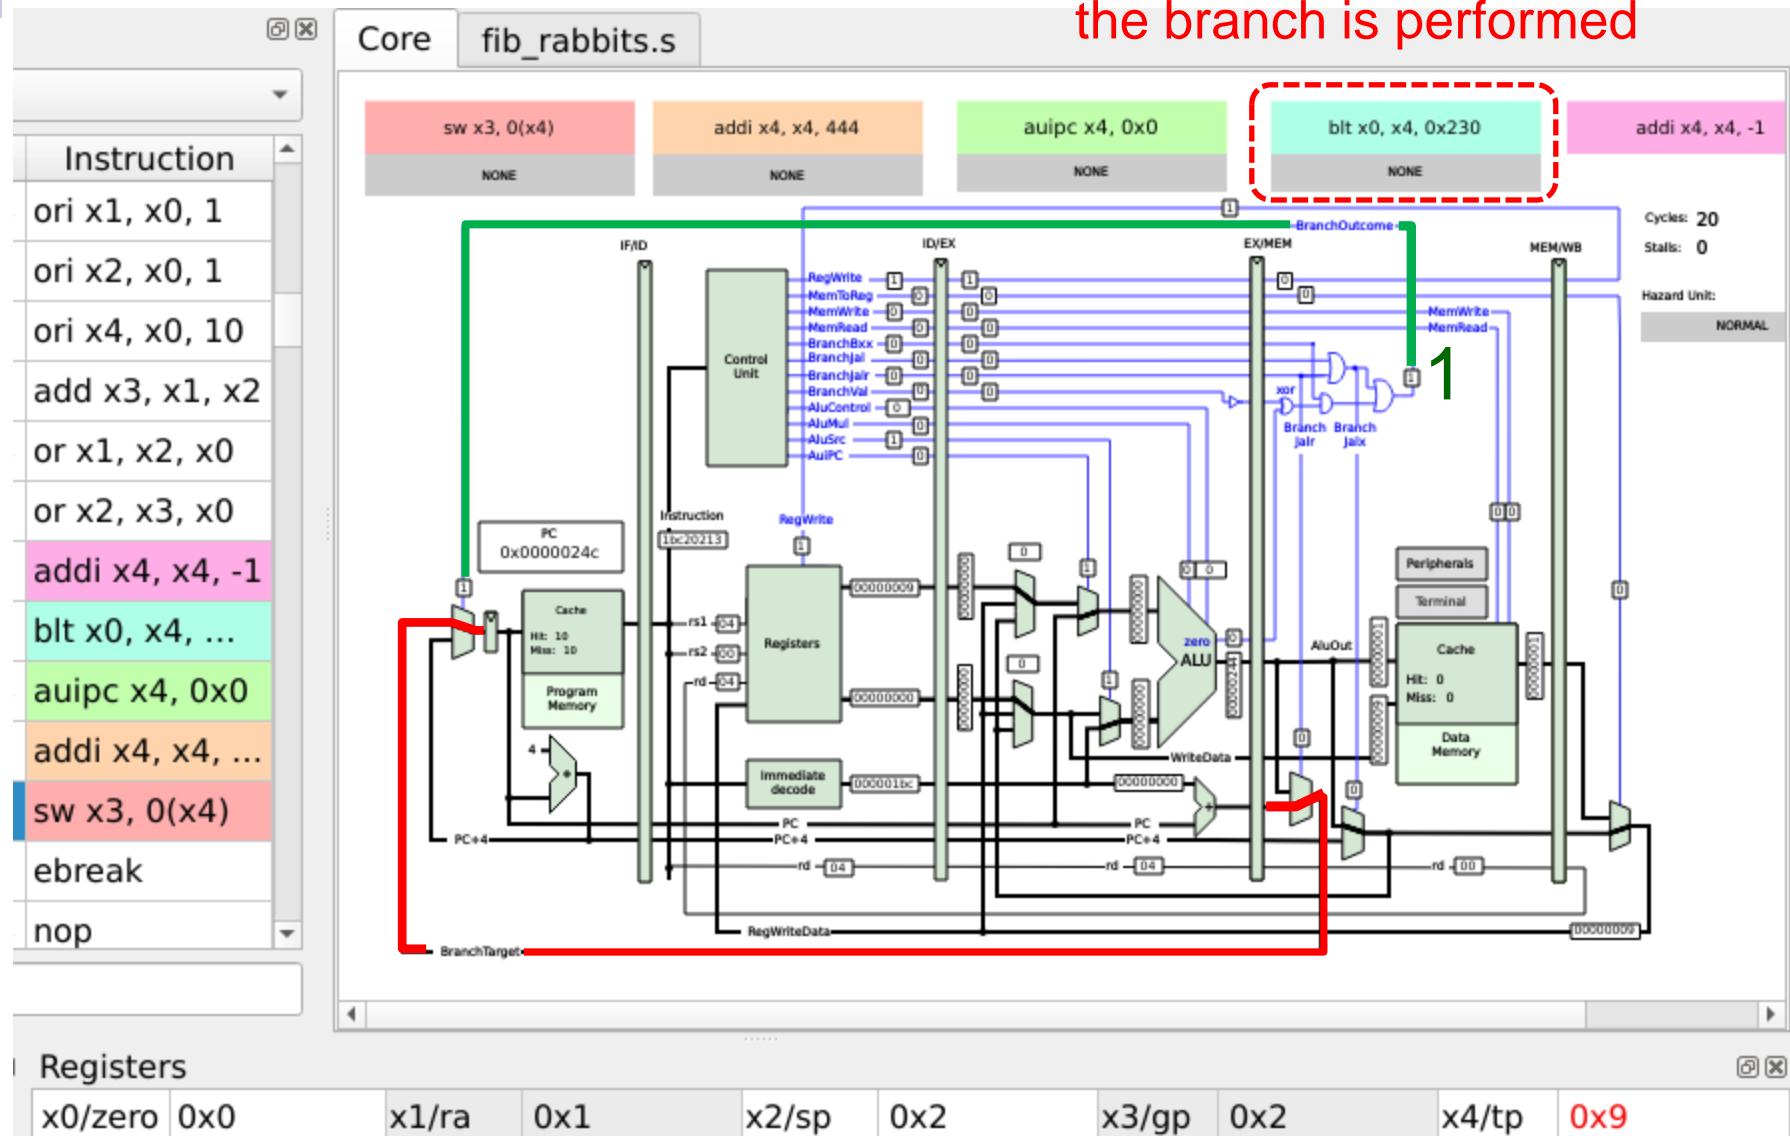
branch address

Registers

x0/zero	0x0	x1/ra	0x1	x2/sp	0x2	x3/gp	0x2	x4/tp	0xa
---------	-----	-------	-----	-------	-----	-------	-----	-------	-----

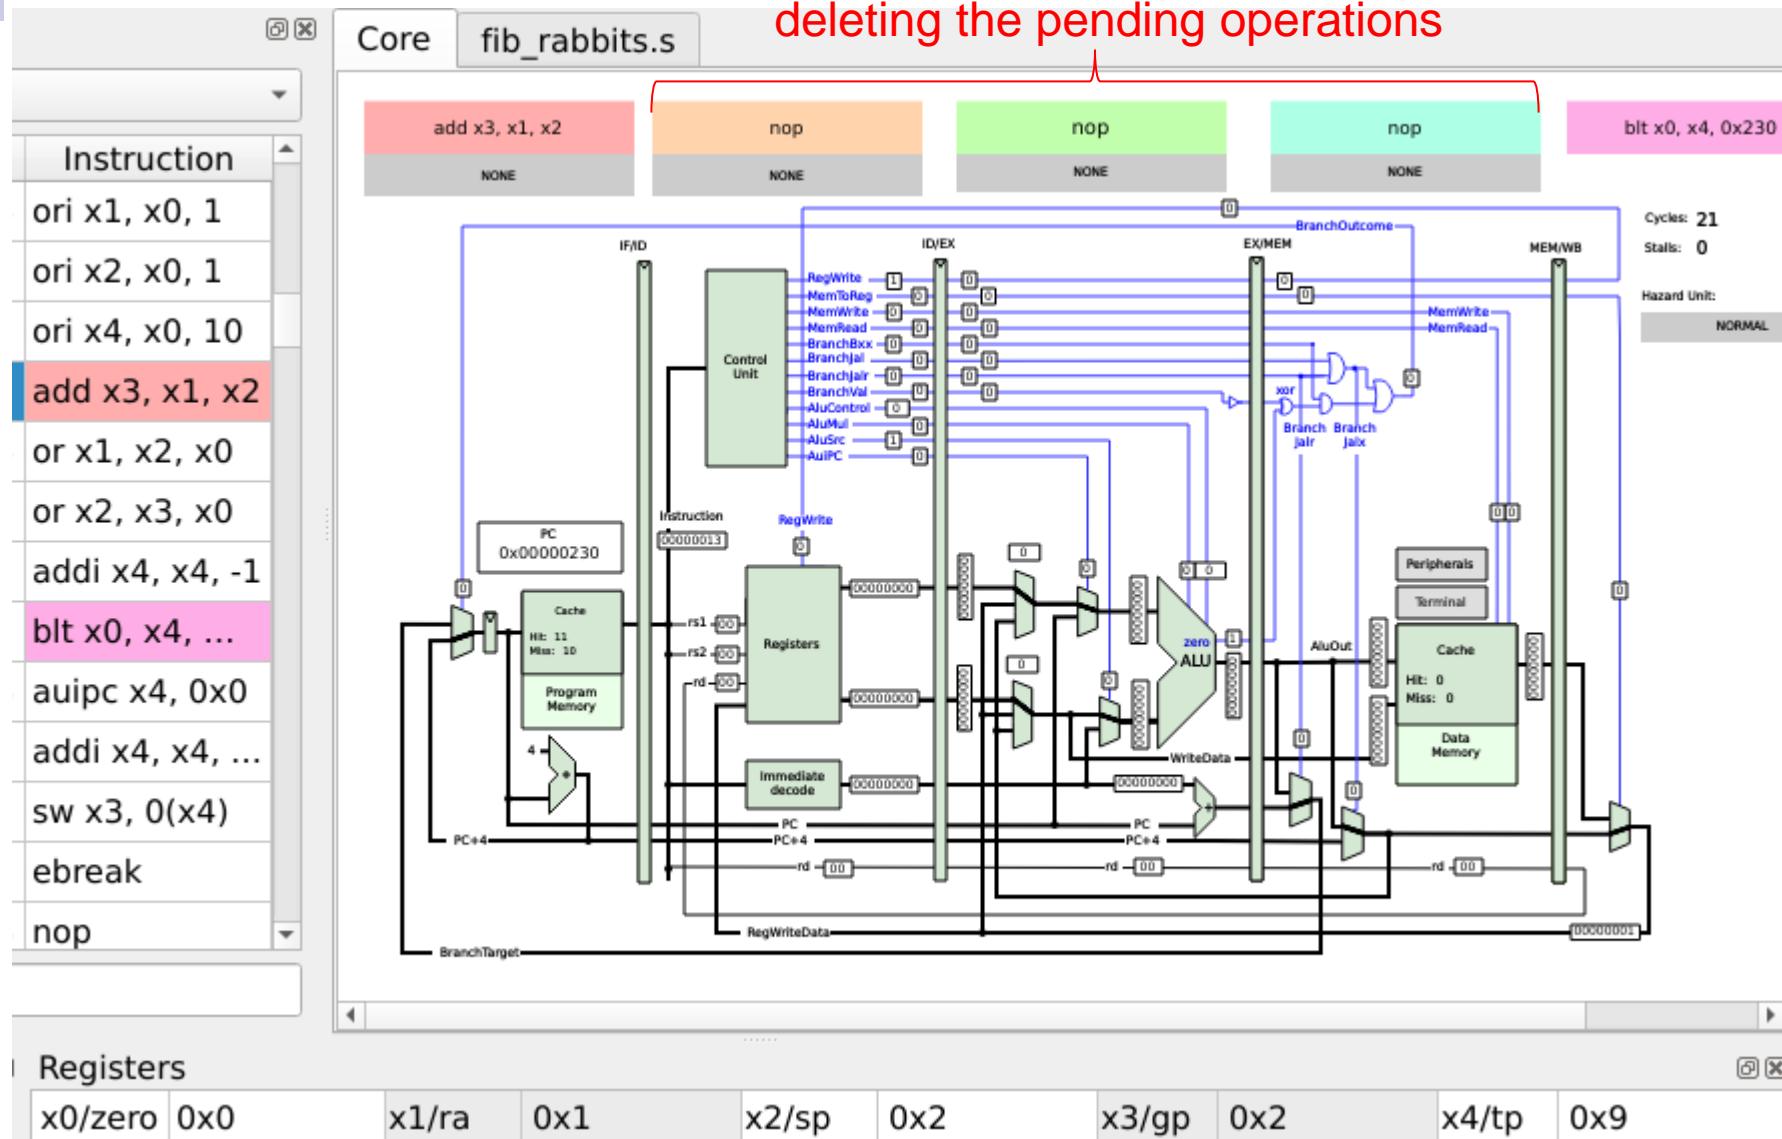
In the MEM stage, it will be decided whether to branch

the branch is performed

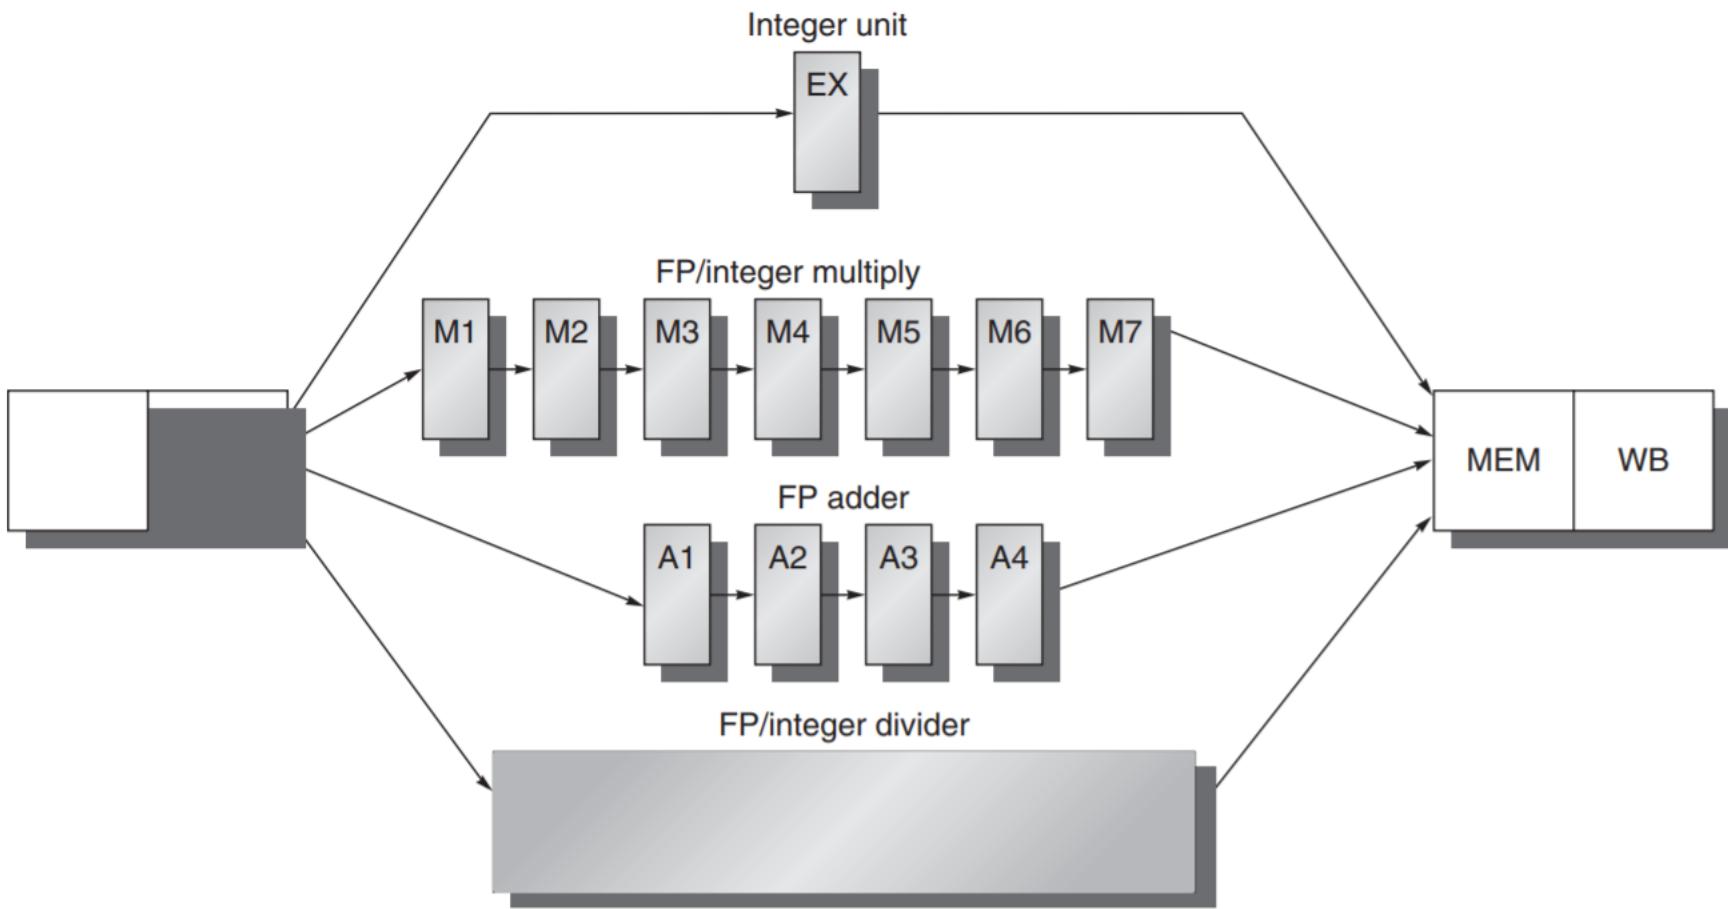


Read instructions are flushed

deleting the pending operations



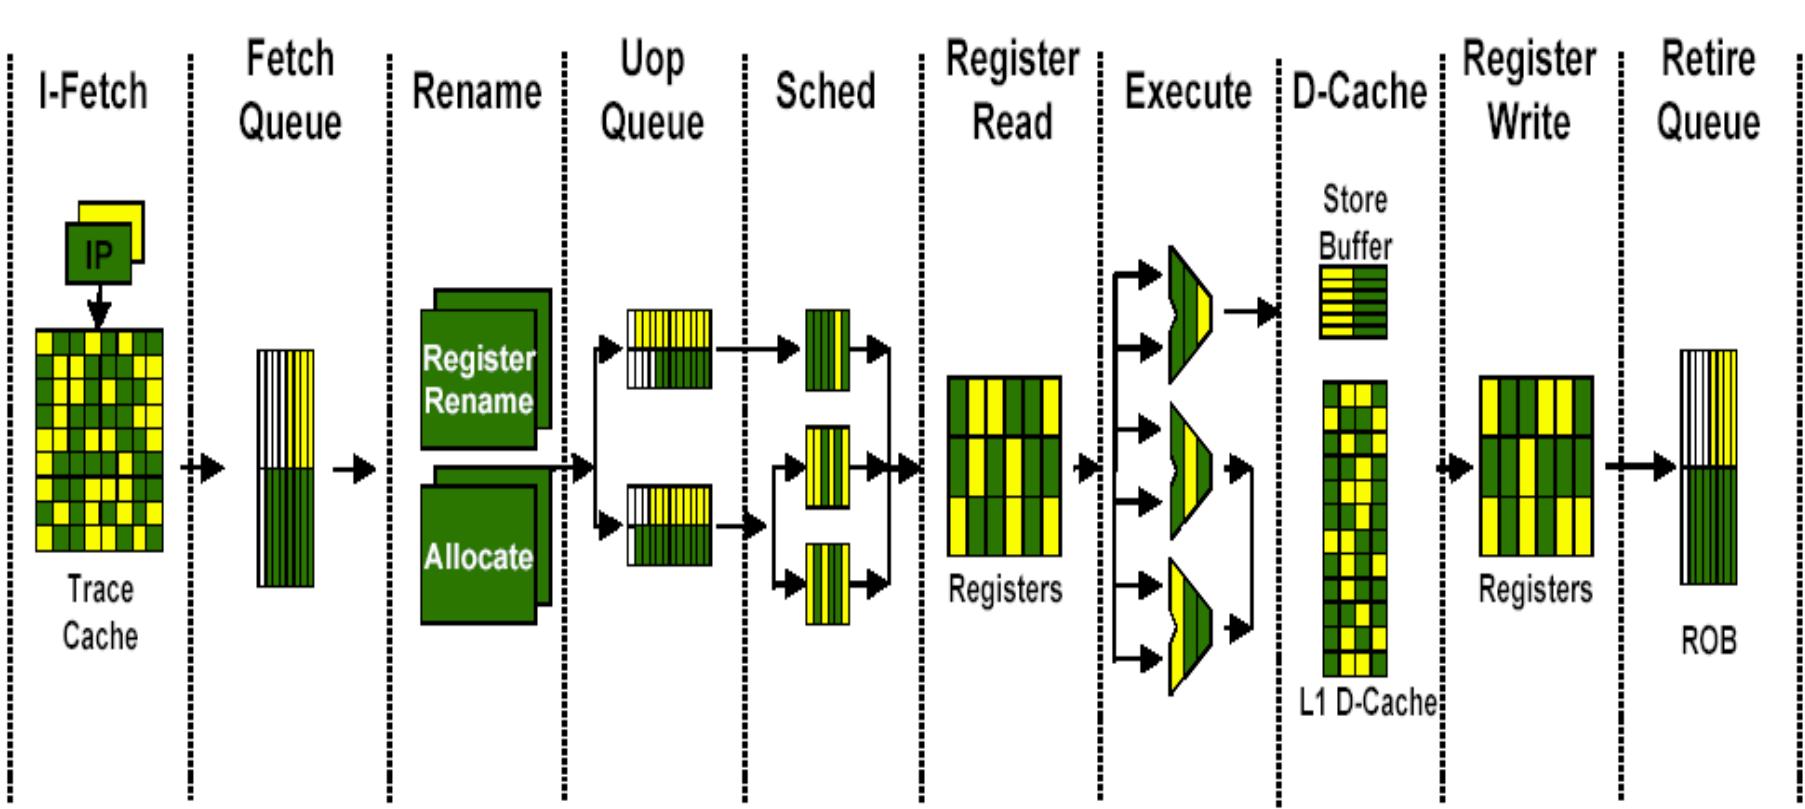
ALU can have multiple parallel data paths



Source of picture: J. L. Hennessy and D. A. Patterson,
Computer Architecture: A Quantitative Approach.

Bigger processors have longer pipelines

Pentium 4: Netburst Microarchitecture's



*Simplified picture, the pipeline actually has 20 steps.
The penalty for branching is high.*

What is the length of pipelines? – approx. ...

5	1993	P5 (Pentium)
10	1995	P6 (Pentium 3)
14	1995	P6 (Pentium Pro)
20	2000	NetBurst (Willamette, 180 nm) - Celeron, Pentium 4:
20	2002	NetBurst (Northwood, 130 nm) - Celeron, Pentium 4+ HT
31	2004	NetBurst (Prescott, 90 nm) - Celeron D, Pentium 4 ExEd
31	2005	NetBurst (Cedar Mill, 65 nm)
31	2006	NetBurst (Presler 65 nm) – Pentium D

14 2020- generation Intel I3, I5, I7 , 10 and following

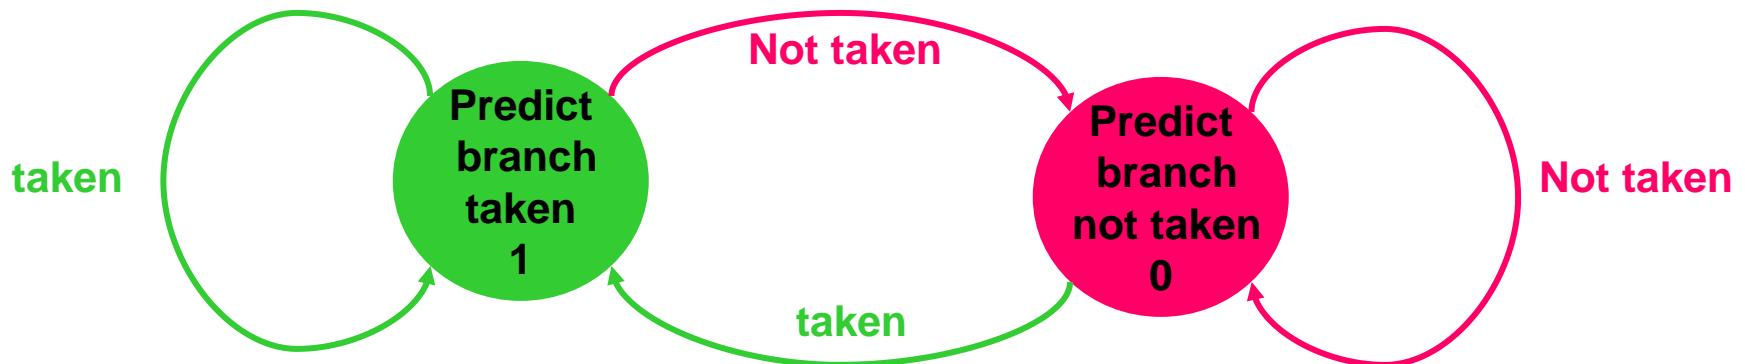
5	2005	ARM 8-9: 5
13	2005	Cortex A8: 13
5-25	2010	Cortex A15: 15-25
8-10	2011	Cortex A7: 8-10
10	2021-	Cortex X2, X3

- The Optimum Pipeline Depth for a Microprocessor:

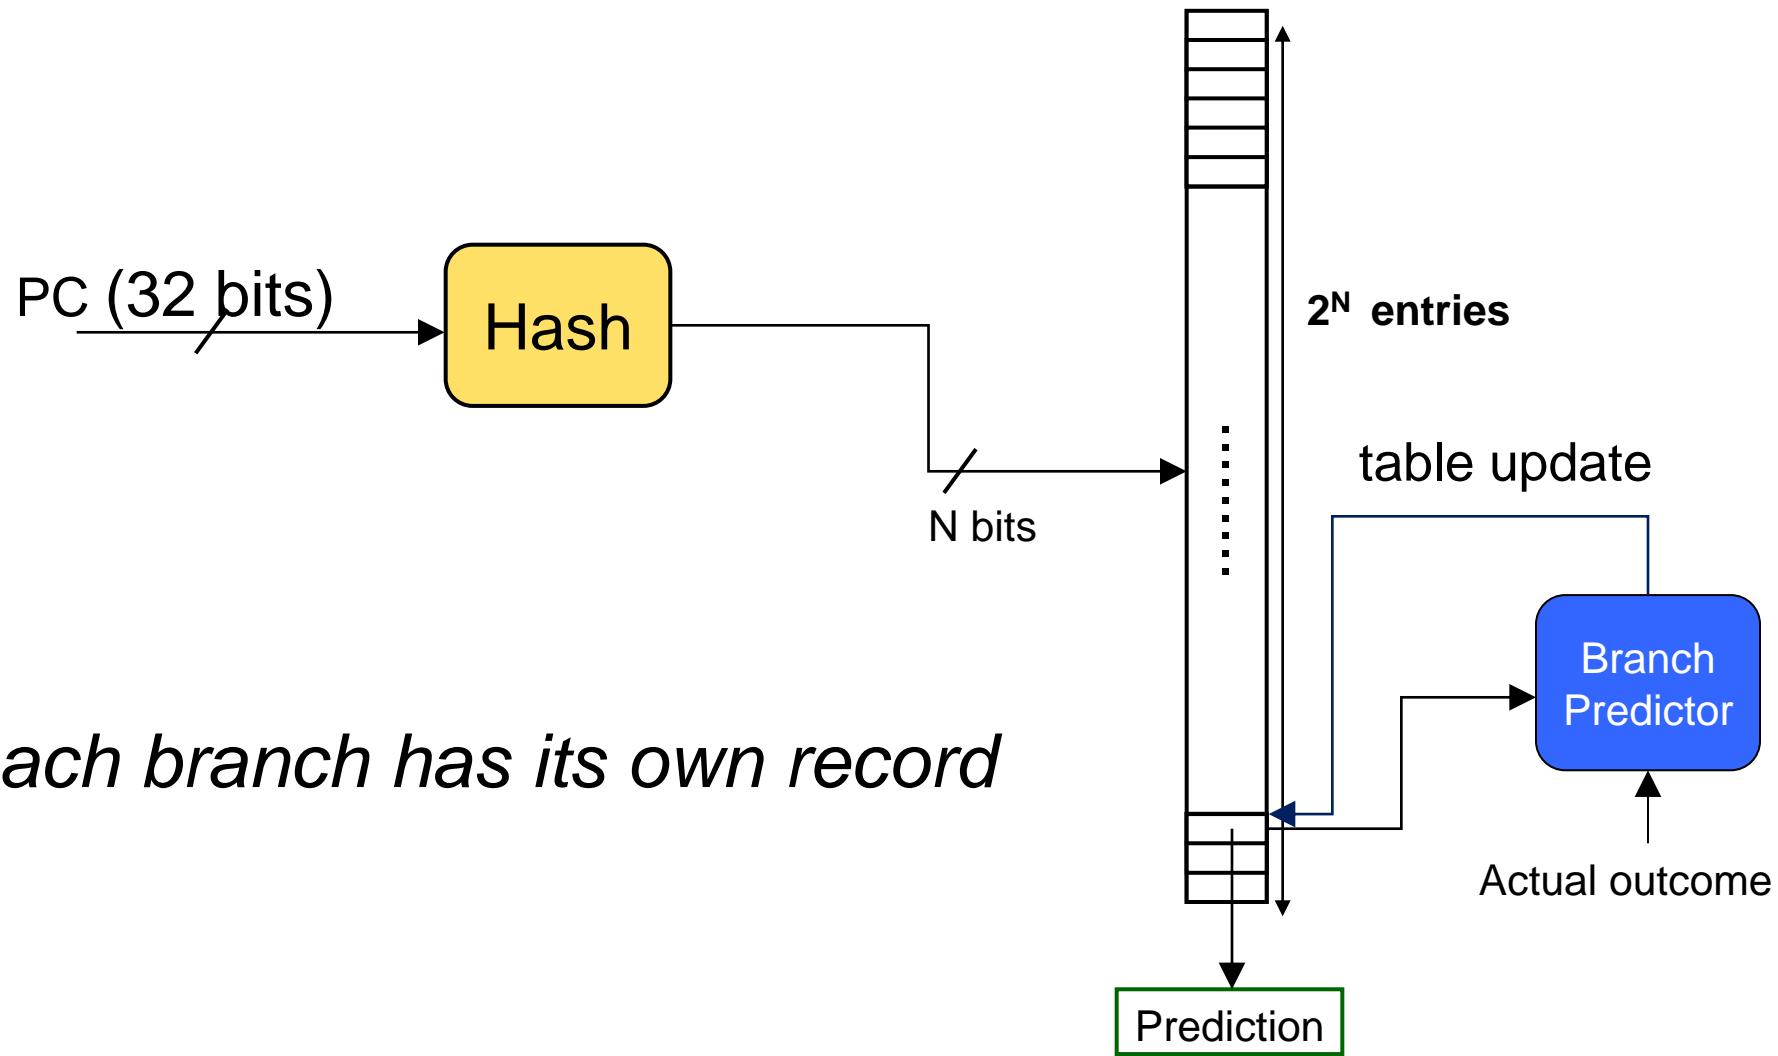
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.93.4333&rep=rep1&type=pdf>

* Prediction of Branches

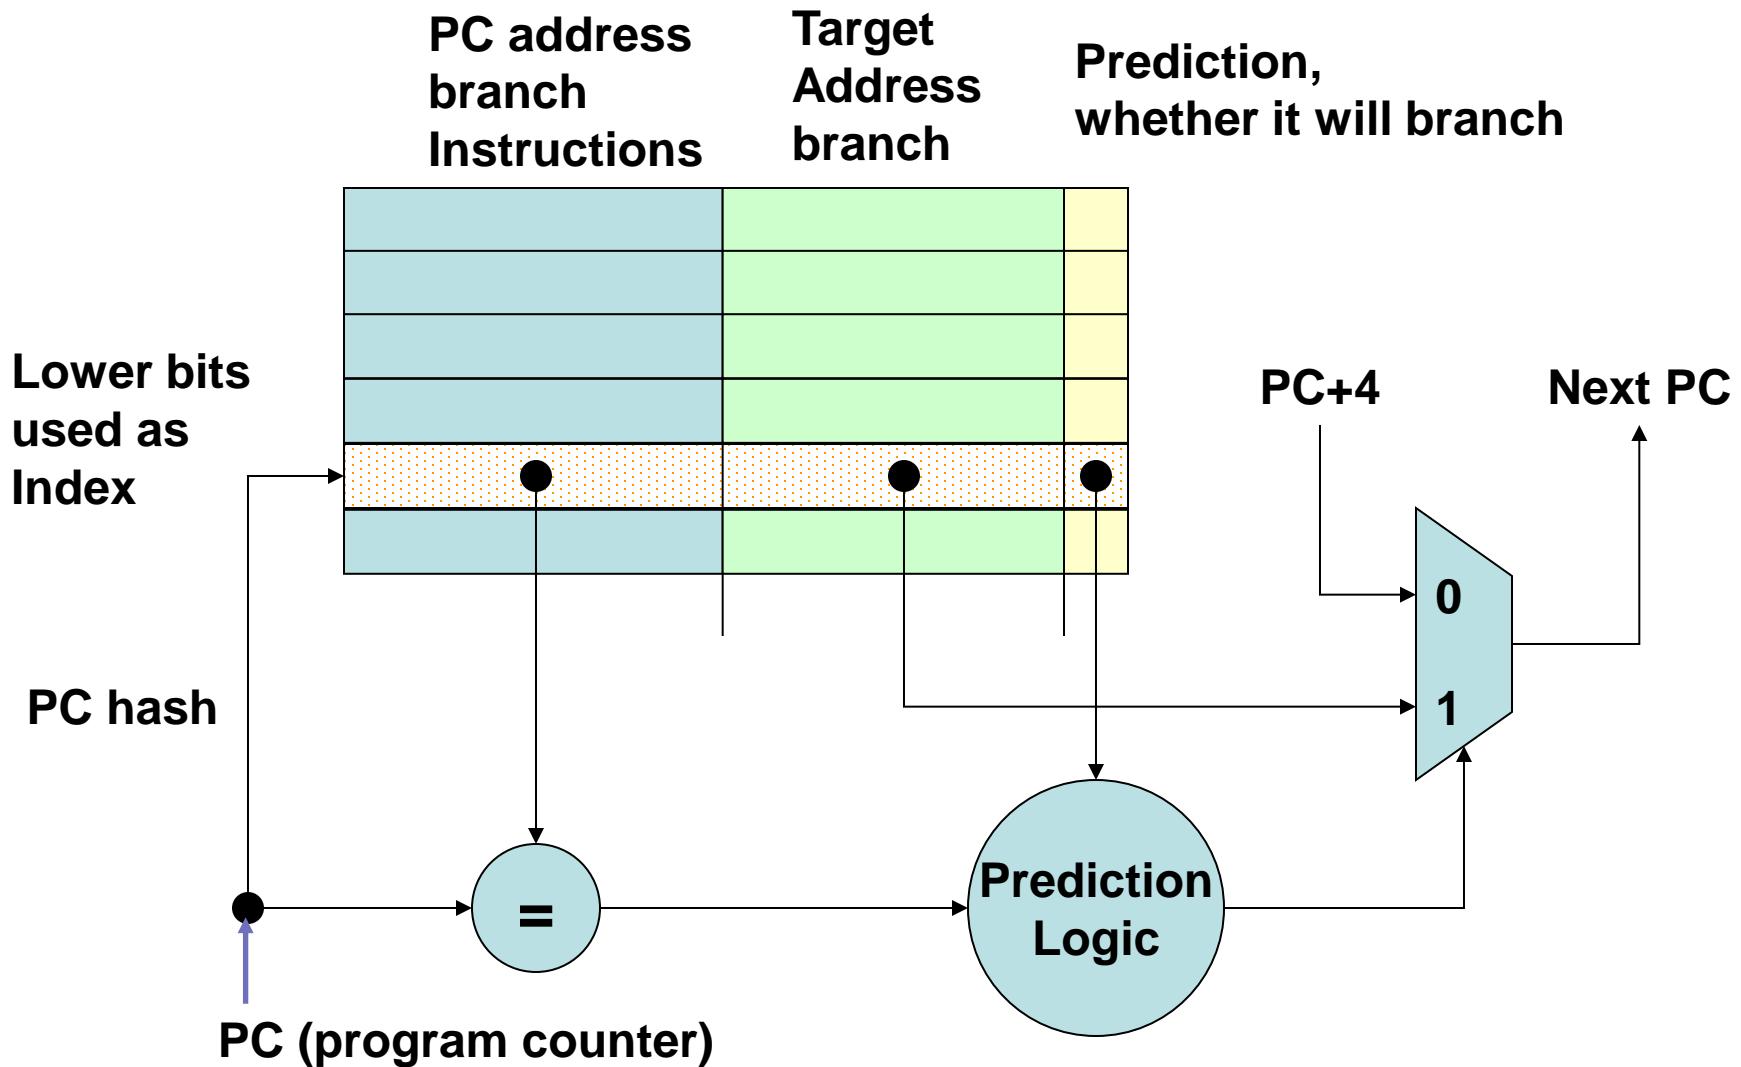
- 1-bit prediction:
has 1 "history bit" that remembers what was done,
when the instruction was last executed:
 - History bit = 1, branch was previously **Taken**
 - History bit = 0, branch was previously **Not taken**



Normal organization of the prediction table

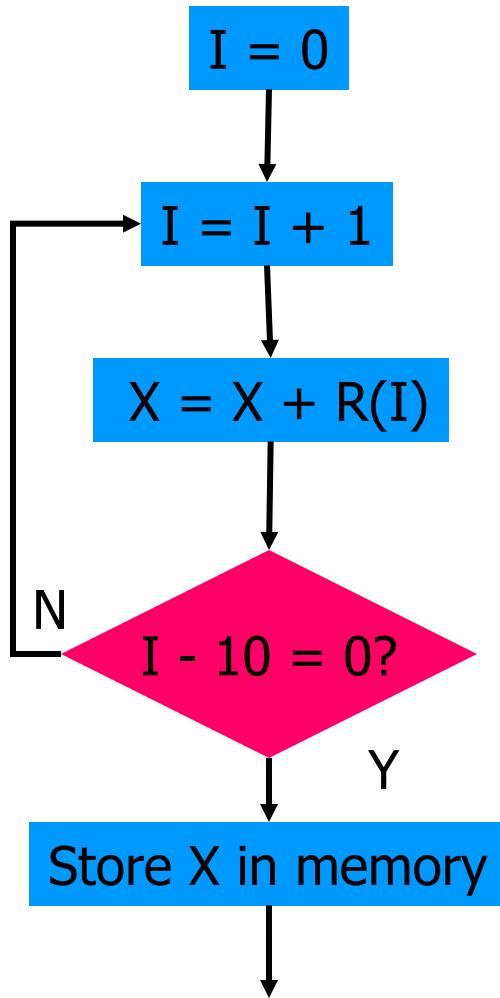


Branch Prediction



Branch Prediction for a Loop

1
2
3
4
5

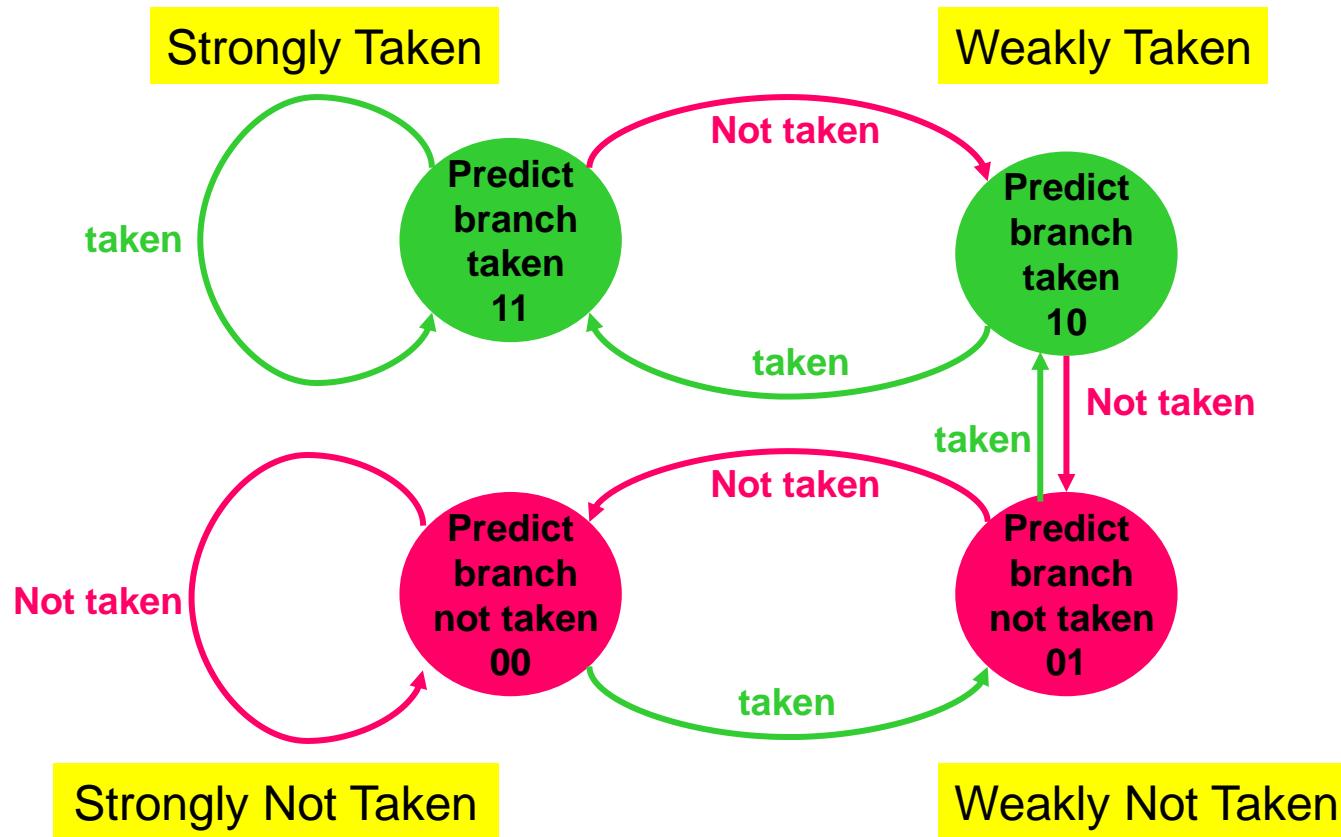


Execution of Instruction 4

Execution seq.	Old Hist. bit	Next instr.			New hist. bit	Prediction
		Before	I	Act.		
1	0	5	1	2	1	Bad
2	1	2	2	2	1	Good
3	1	2	3	2	1	Good
4	1	2	4	2	1	Good
5	1	2	5	2	1	Good
6	1	2	6	2	1	Good
7	1	2	7	2	1	Good
8	1	2	8	2	1	Good
9	1	2	9	2	1	Good
10	1	2	10	5	0	Bad

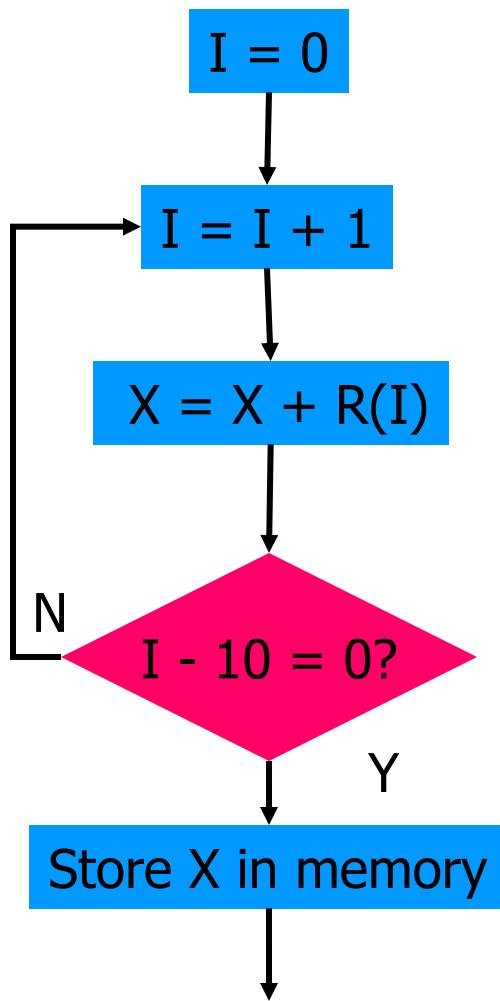
bit = 0 branch not taken, bit = 1 branch taken.

- It is also called "saturating counter!".



Branch Prediction for a Loop

1
2
3
4
5



Execution of Instruction 4

Execution seq.	Old Before e.Buf	Next instr.			New pred. Buf	Prediction
		Befor e.	I	Act.		
1	10	2	1	2	11	Good
2	11	2	2	2	11	Good
3	11	2	3	2	11	Good
4	11	2	4	2	11	Good
5	11	2	5	2	11	Good
6	11	2	6	2	11	Good
7	11	2	7	2	11	Good
8	11	2	8	2	11	Good
9	11	2	9	2	11	Good
10	11	2	10	5	10	Bad

```
for (int i = 0; i < 100; i++)
{
    for (int j = 0; j < 10; j++)
    {
        // some code
    }
}
```

Number of wrong predictions:

1-bit predictor: outer loop 2 + inner loop $100 \times 2 = 202$

2-bit predictor: outer loop 1 + inner loop $100 \times 1 = 101$

Example 2

```
int arr[] = { 1, 2, 3, -10, -20, 30, 100, -200, -300};  
int count = 0;  
for (int i = 0; i < sizeof(arr)/sizeof(int); i++)  
{ if (arr[i] < 0) count++; } // assembler if(arr[i]>=0) goto labelNext
```

Number of wrong predictions:

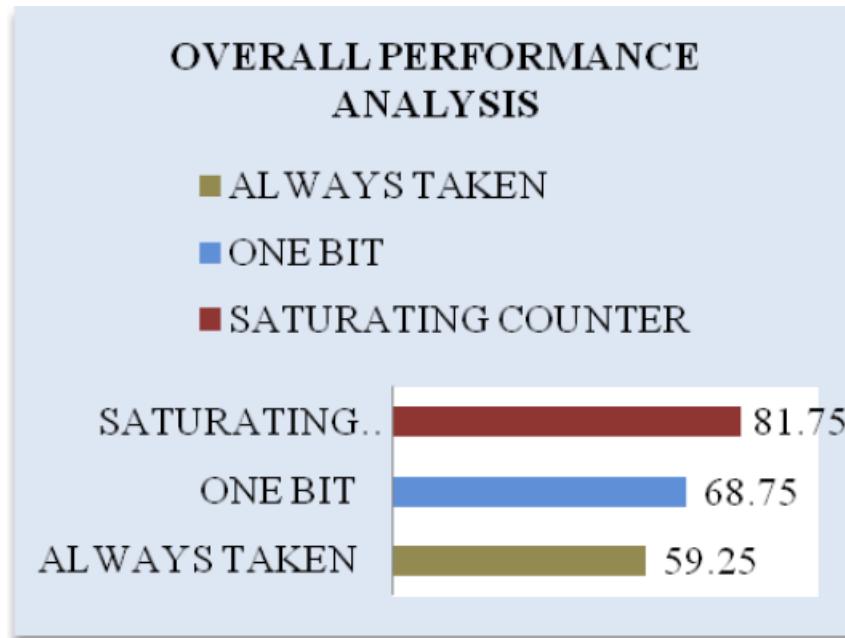
1-bit init NT predictor: outer loop 2 + inner loop 4

2-bit init WNT predictor: outer loop 2 + inner loop 6

2-bit init WT predictor: outer loop 1+ inner loop 5

Number	1	2	3	-10	-20	30	100	-200	-300
Assembler Needs	T	T	T	NT	NT	T	T	NT	NT
1bit Predictor	NT	T	T	T	NT	NT	T	T	NT
2bit Predictor WNT	WNT	WT	T	T	WT	WNT	WT	T	WT
2bit Predictor WT	WT	T	T	T	WT	WNT	WT	T	WT

Some Available Statistics



The larger numbers in the picture mean better prediction

Source : <https://ieeexplore.ieee.org/document/6918861>

H. Arora, S. Kotecha and R. Samyal, "Dynamic Branch Prediction Modeller for RISC Architecture," *2013 International Conference on Machine Intelligence and Research Advancement*, Katra, 2013, pp. 397-401.

They consider previous branches

```
if (x==2) // branch b1
```

...

```
if (y==2) // branch b2
```

...

```
if(x!=y) { ... } // branch b3 depends on the  
results of b1 and b2
```

(2,1) Correlated predictor

We switch between the results of 4 predictors: **P00 | P01 | P10 | P11**

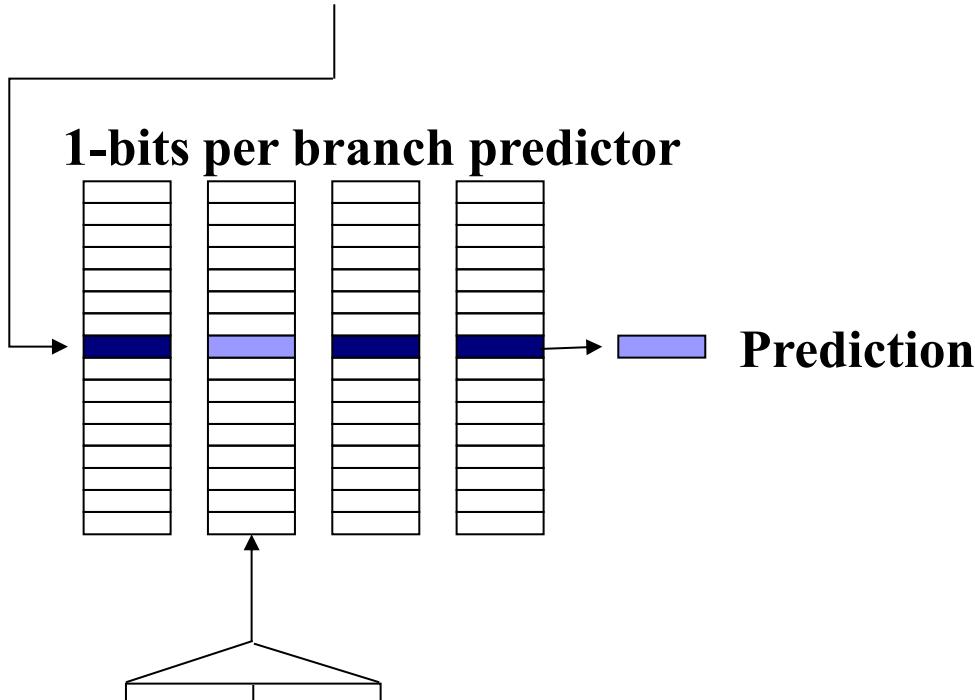
The predictor is used when	P00	P01	P10	P11
1. the previous branch was	Not taken	Not taken	Taken from	Taken from
2. the previous branch was	Not taken	Taken from	Not taken	Taken from

A (2,1) correlation predictor

- (2,1) means $2^2 = 4$ predictors, 1 bit each
- 2 indicates that the result of the previous 2 branches is selected.

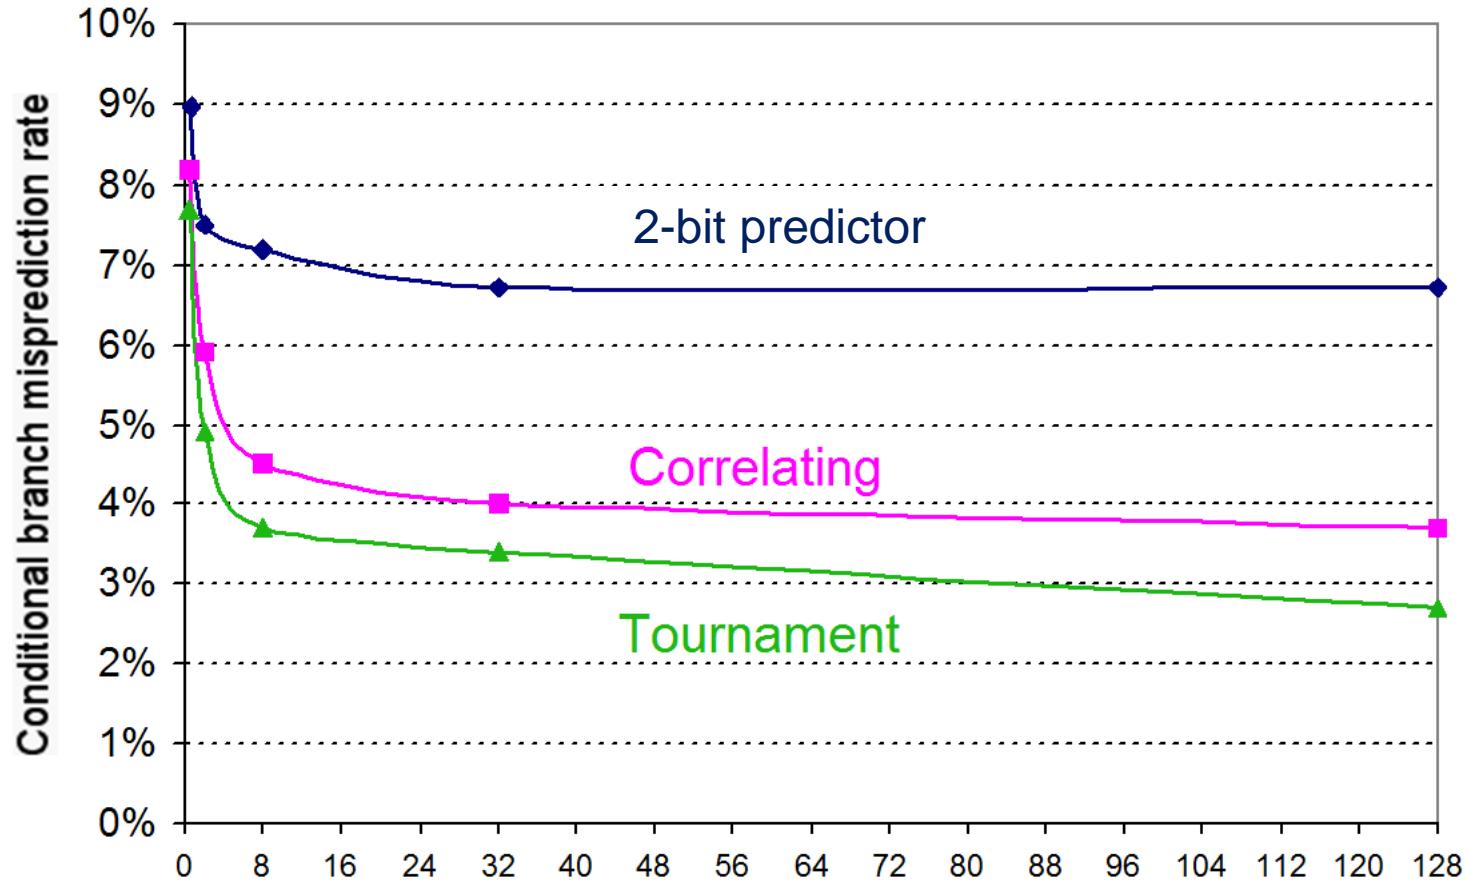
➤ Example (2,1) predictor

Hash of branch address



2-bit global branch history

A little bit of statistics again



Tournament - multiple predictors of different types. They all predict, but the result of the one with the most previous successes is taken. They are used in large processors, but the manufacturers hide their exact principles.

Branches can sometimes be avoided

There are many tricks on the web suitable for time-critical loops. For example, the absolute value of a 32-bit signed integer x can be calculated without jumps.

Code with *unpredictable branch* dependable on data

C code	RISC V if x in $t1$	Comment
if($x < 0$) $x = -x;$	blt x1, zero, Skip1	// if(tmp==0) goto Skip
	sub x1, zero, x1	// $x = -x;$
Skip1:	...	

Faster C code	RISC V if x in $x1$	Comment
int tmp = $x >> 31$;	sra x2, x1, 31	// signed right shift; tmp = $x < 0 ? -1 : 0$
$x \wedge= tmp$;	xor x1, x1, x2	// 1 st compliment of x, if tmp=-1 (0xFFFFFFFF)
$x -= tmp$;	sub x1, x1, x2	// add +1 if tmp = -1

Note: For RISCV, we save only 1 instruction out of 3 that are deleted from the pipeline.
If we compile C code for a processor with a longer pipeline, the misprediction will be more expensive.

```
int CountMinus(int* arr, int len)
{ int count = 0;
  for (int i = 0; i < len; i++) { if (arr[i] < 0) count++; }
  return count;
}

int CountMinusNoBranch(int* arr, int len)
{
  int count = 0;
  for (int i = 0; i < len; i++) { count -= (arr[i]>>31); }
  return count;
}
```

Visual Studio 2022 C++ Compiler

In DEBUG mode (), CountMinusNoBranch () is translated to code, which will be 3 faster than CountMinus().*

In RELEASE mode, both variants run equally fast because the compiler has replaced

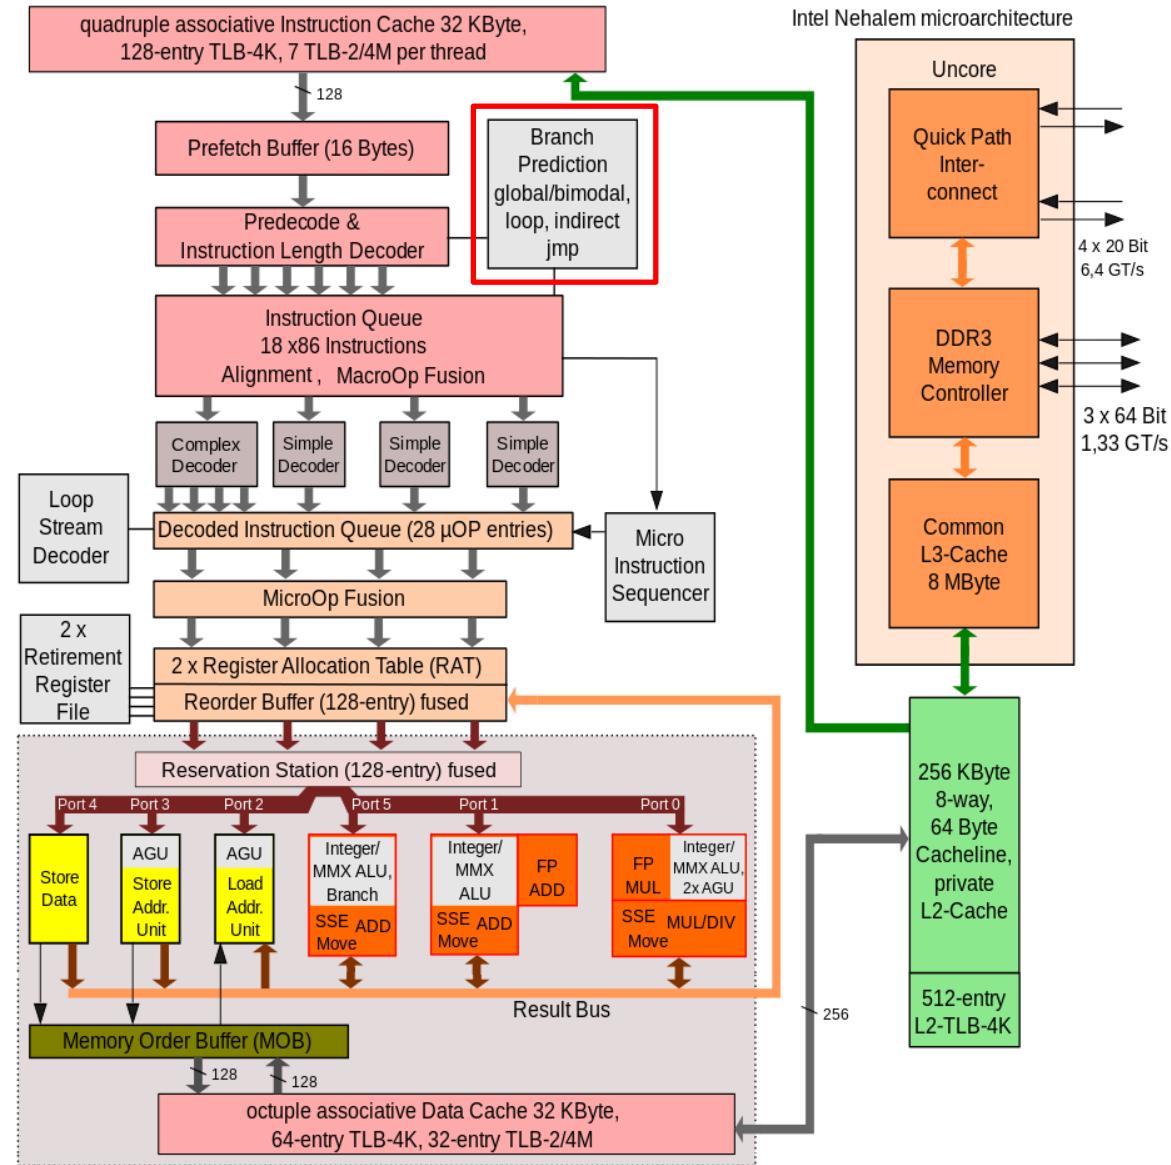
```
if (arr[i] < 0) count++;
```

during jump-free optimizations

```
count -= (arr[i]>>31);
```

However, it cannot handle more complex data-dependent jumps. It's always worth considering whether your algorithm could be modified so that there aren't so many of them.

Motivation for the lecture - Intel Nehalem (Core i7) - 2008



- http://en.wikipedia.org/wiki/File:Intel_Nehalem_arch.svg

ALU and Ripple Effect



„Ježku, vstávej! Pozor dávej:
dobře vlákno pročesávej,

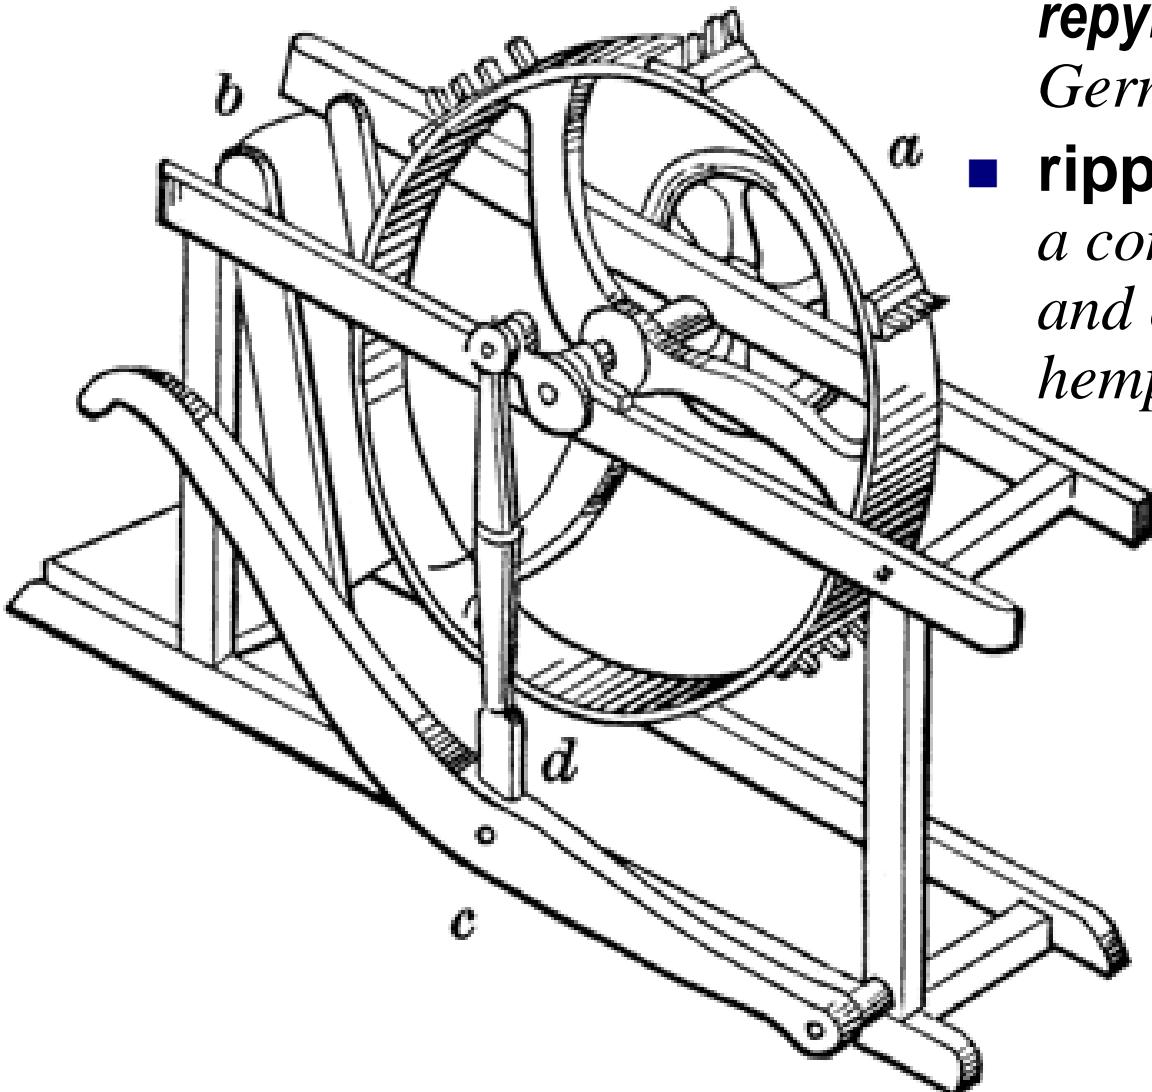
at' je jako vánek,
tenčí nežli sponky!“

pohadkyprodeti.cz



Ripple (1st meaning)

ripple



- **ripple** *from Middle English.*
repyll *akin to Old High German **riffila**= saw*
- **ripple** *a large instrument like a comb for removing seeds and other matter from flax or hemp*



saw



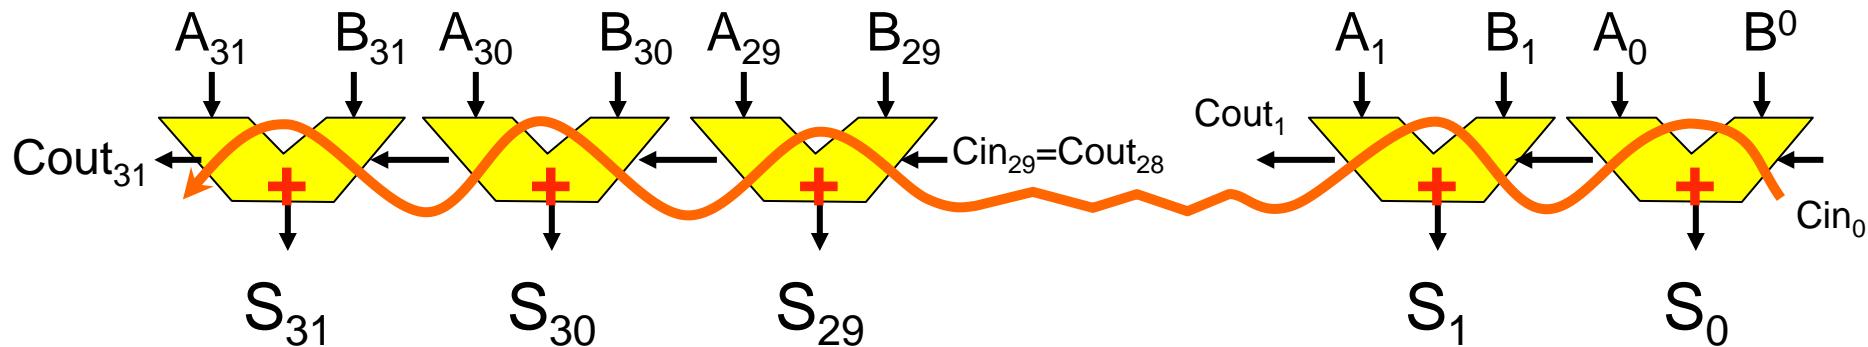
Ripple (2nd meaning)

- **ripple:** *slightly undulating on the surface or appearing in small waves*
- *Ripple Effect in economics: the Domino or Avalanche effect*
in physics: Chain Reaction



- The SoP (coverage 1) and PoS (coverage 0) methods express logic functions as expressions with three levels:
SoP: input inverters, AND gates of implicants, and a merging OR gate.
PoS: inverters of inputs, OR gates of implicants and merging AND gate.
- Both adders and multipliers are described by Karnaugh maps (KM), which contain many small groups of 0 and 1. Their coverage then has many implicants and their number increases exponentially with the bit length of the inputs.
- We need to describe them with multilevel logic, but there is currently no method that can efficiently decompose a complicated KM into smaller minimizable blocks in acceptable time. We have to invent the decomposition ourselves.
- Division is converted to multiplication in circuits. A general divisor can be used, but its circuit is slow and takes too many resources.

For more details, see the Logic Circuits textbook, Chapter 6.



■ Simplest N-bit adder

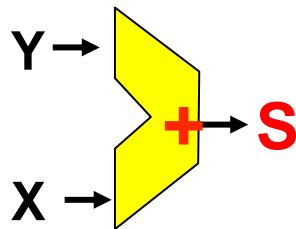
- we chain 1-bit full adders
- "Carry" ripple through their chain
- Minimal number of logical elements
- Delay is given by the last Cout - $2^*(N-1) + 3$ gates of the last adder
 $= (2 N+1)$ times propagation delay of 1 gate



Karnaugh maps for bits of adder

S_0	0	1	X
0	0	1	
1	1	0	

Y



S_1	00	01	11	10	X
00	0	0	1	1	
01	0	1	0	1	
11	1	0	1	0	
10	1	1	0	0	

Y

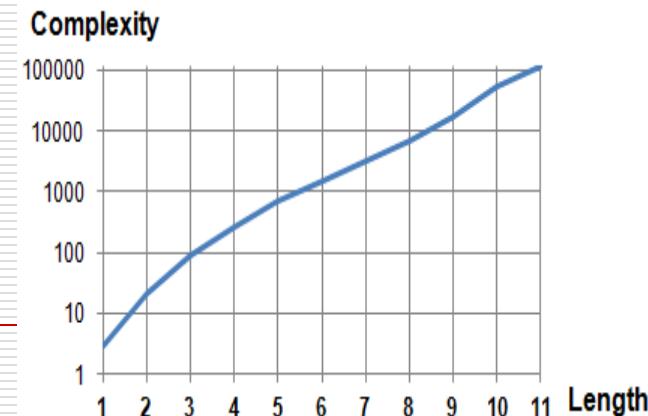
$S = A + B$
S bity [$S_{n-1} \dots S_2 S_1 S_0$]

Too small groups of neighbour bits!
The coverage of adder KMs has exponential complexity in gates and interconnections.

S_2	0	1	3	2	6	7	5	4	X
0 000	0	0	0	0	1	1	1	1	
1 001	0	0	1	0	1	0	1	1	
3 011	0	1	1	1	0	0	0	1	
2 010	0	0	1	0	0	0	1	1	
6 110	1	0	0	0	1	1	0	0	
7 111	1	0	0	0	1	0	1	0	
5 101	1	0	0	1	0	1	0	0	
4 100	1	0	1	1	0	0	0	0	

Y

Design of adder directly from KM by Sum-Of-Products



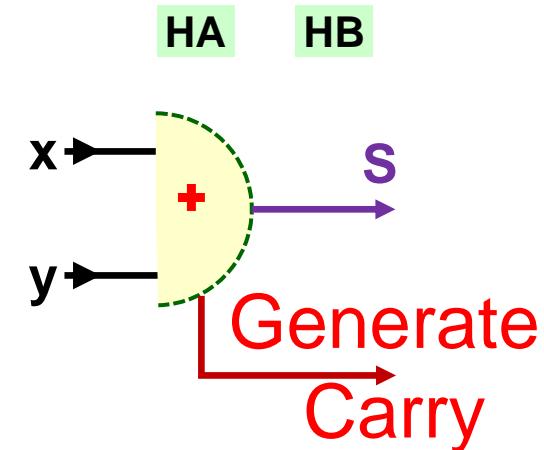
bit n of sum	Bit n of sum		Total operations bit 0 až n
	OR operations	AND operations	
0	1	2	3
1	5	14	22
2	15	52	89
3	35	148	272
4	75	380	727
5	152	688	1567
6	320	1400	3287
7	768	3072	7127
9	2304	8192	17623
10	8194	27648	53465
11	16388	46080	115933

*The calculation was performed by logic minimizer BOOM
developed by Dept. of Computers of CTU-FEE*

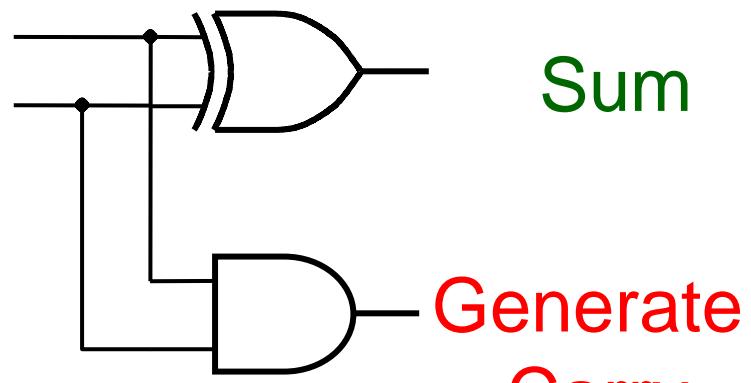
Half adder without Carry-In

$$\begin{array}{r}
 0 \\
 + 0 \\
 \hline
 0
 \end{array}
 \quad
 \begin{array}{r}
 0 \\
 + 1 \\
 \hline
 1
 \end{array}
 \quad
 \begin{array}{r}
 1 \\
 + 0 \\
 \hline
 1
 \end{array}
 \quad
 \begin{array}{r}
 1 \\
 + 1 \\
 \hline
 10
 \end{array}$$

Sum Generate Carry



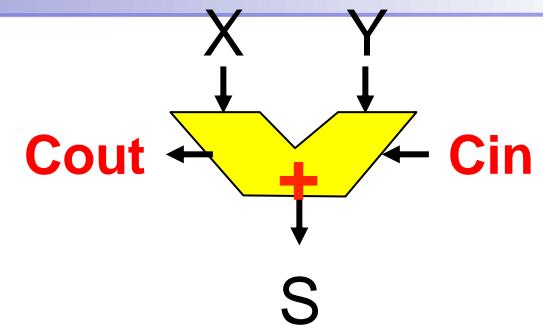
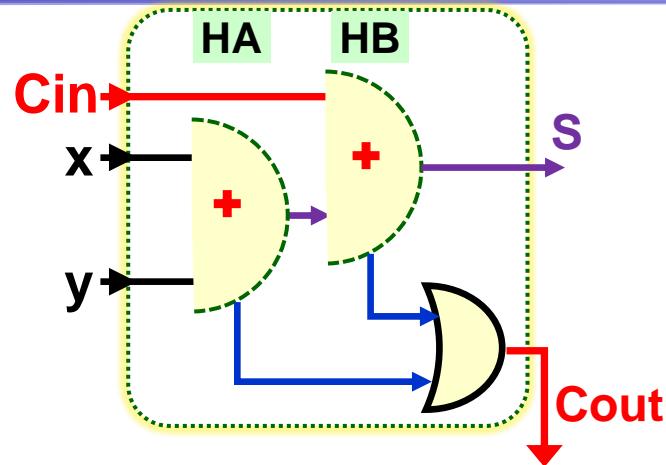
X	Y	Sum	G
0	0	0	0
0	1	1	0
1	0	1	0
1	1	0	1



Sum $\leq X \text{ xor } Y;$
 G $\leq X \text{ and } Y$

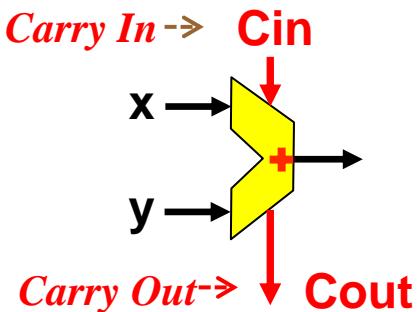
Full 1bit adder

Half Adder A + Half Adder B

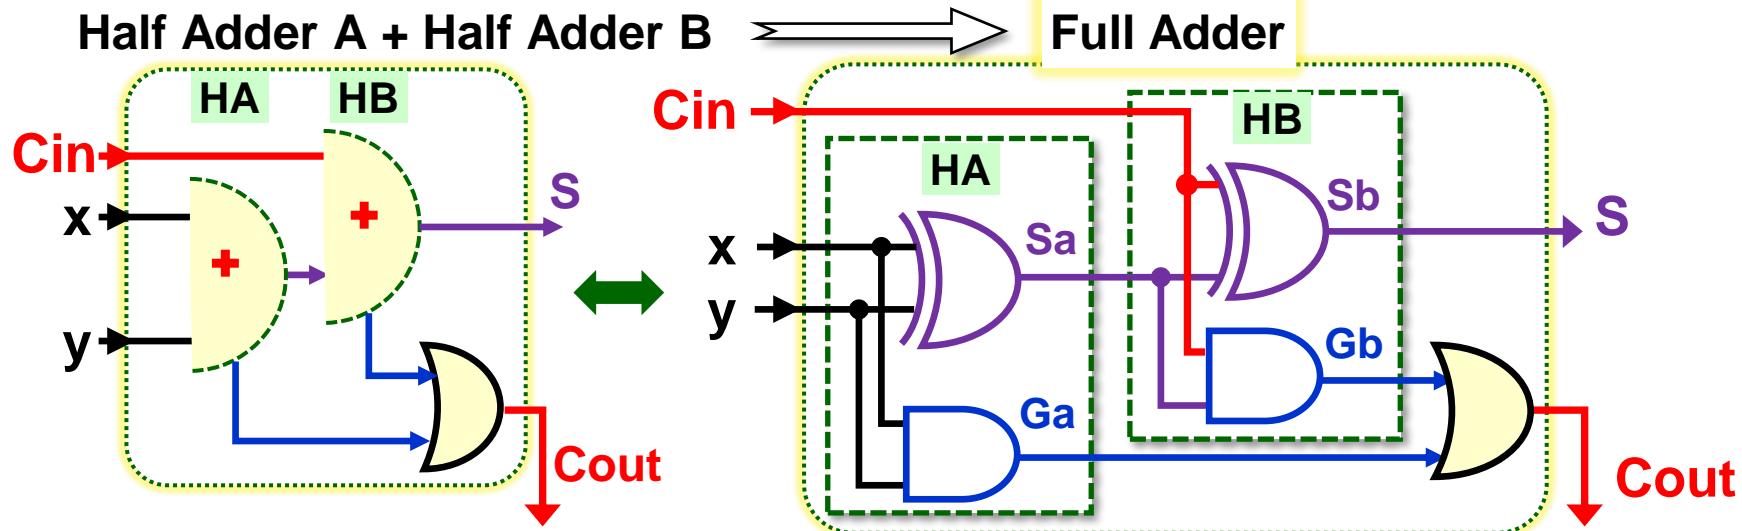
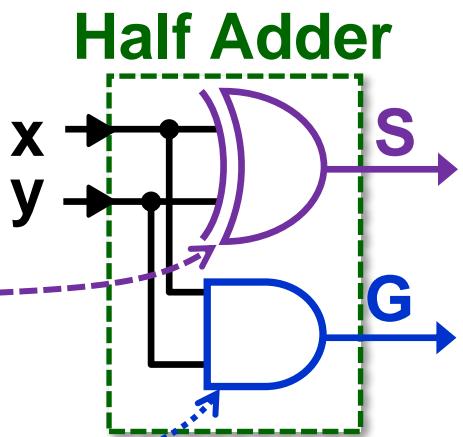


X	0	0	1	1	0	0	1	1
$+Y$	0	1	0	1	0	1	0	1
Sum	00	01	01	10	00	01	01	10
+ Carry-In	0	0	0	0	1	1	1	1
CarryOut Sum	00	01	01	10	01	10	10	11

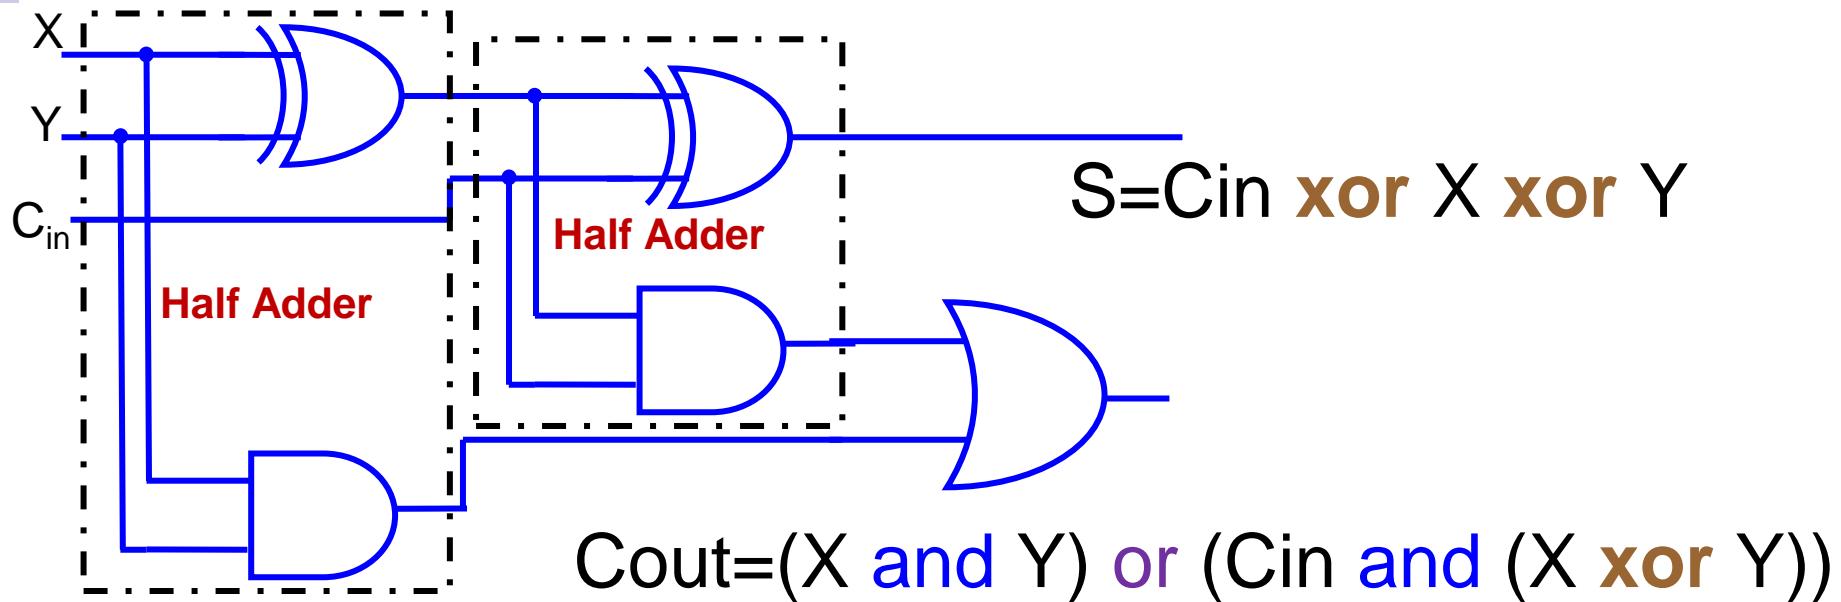
1 bit adder



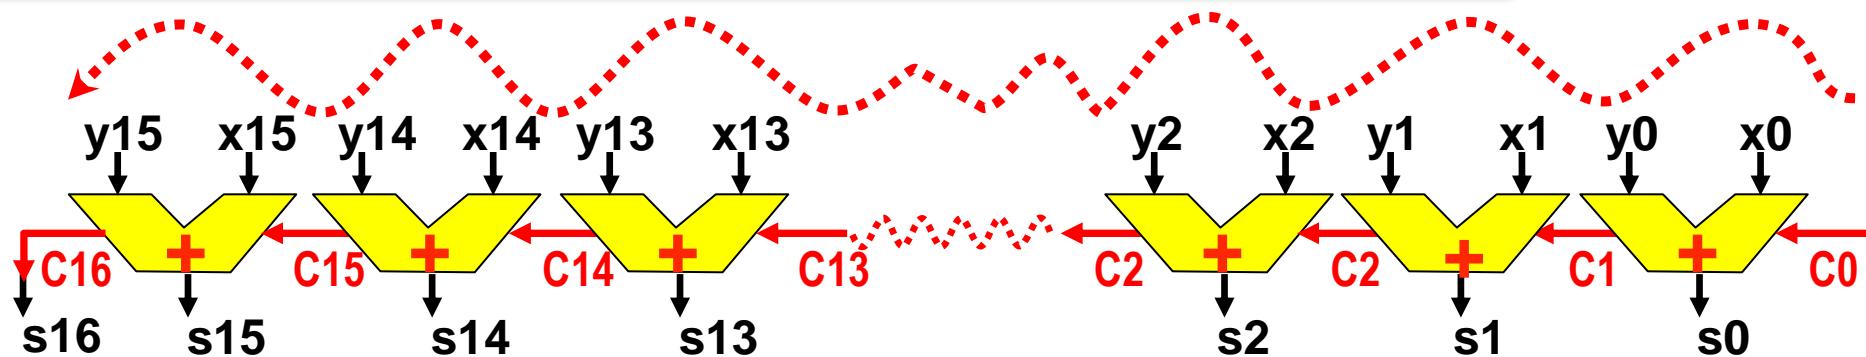
$(x+y)_{dec}$	x	y	$(x+y)_{bin}$
			G S
0	0	0	0 0
1	0	1	0 1
1	1	0	0 1
2	1	1	1 0



Full Adder Equations



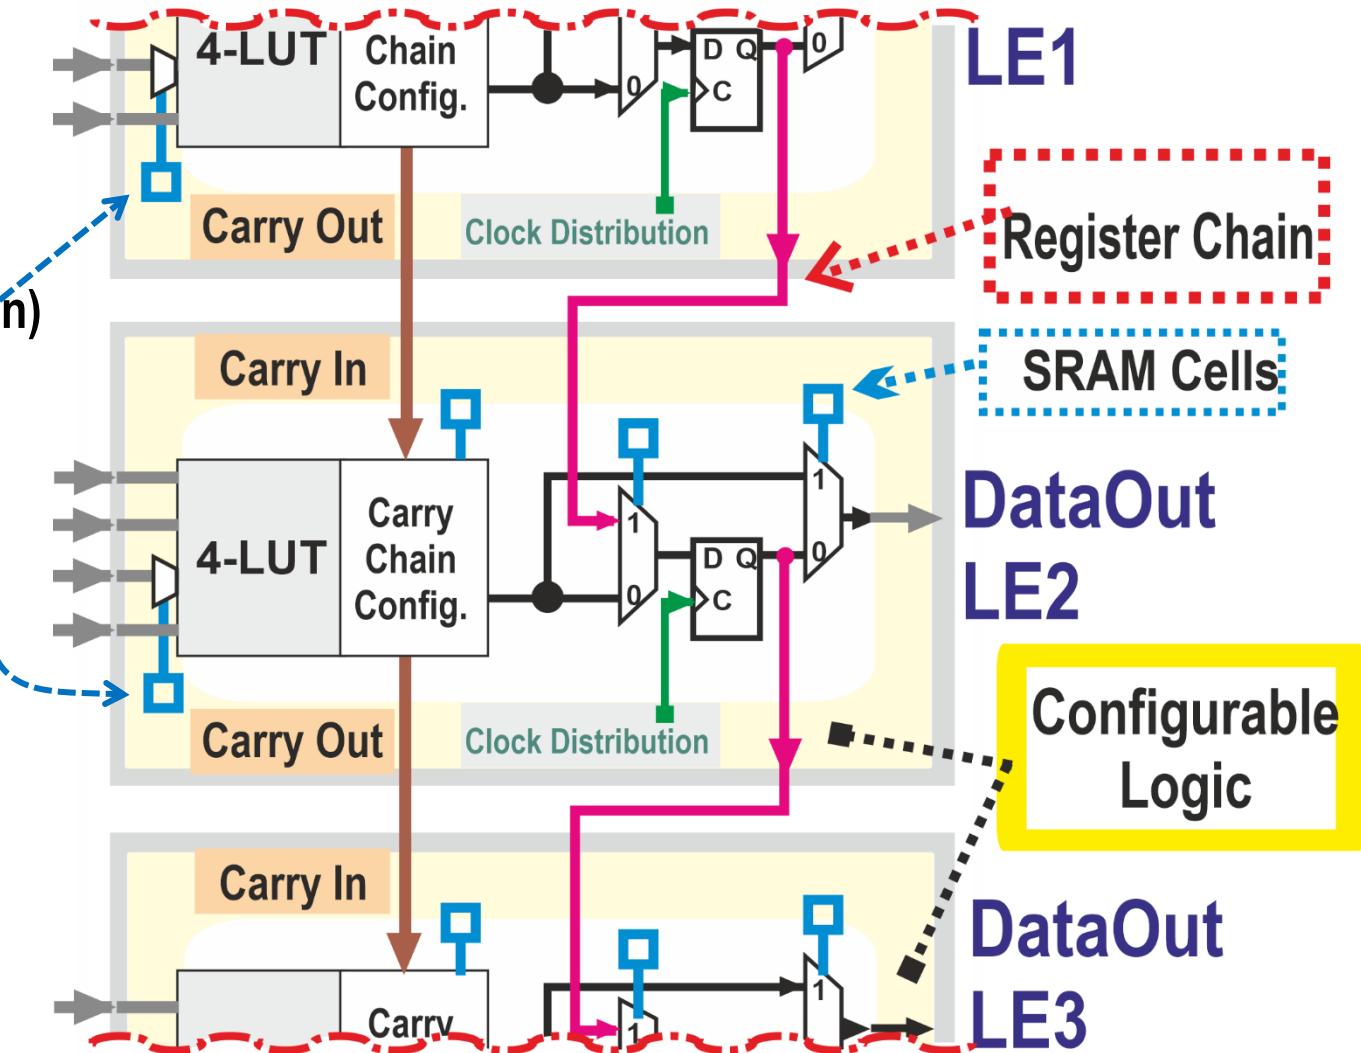
Decomposition into a chain of 1-bit adders



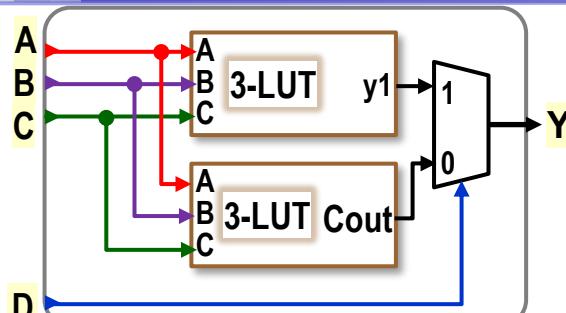
For **unsigned** addition,
 C_{16} announces that the result overflowed,
i.e., it is out of range.
 C_{16} has no such meaning for the addition
of **signed** numbers.
See [Binary Prerequisite](#), page 11 and 19.

Ripple support in Cyclone II and IV Logical Elements (LEs)

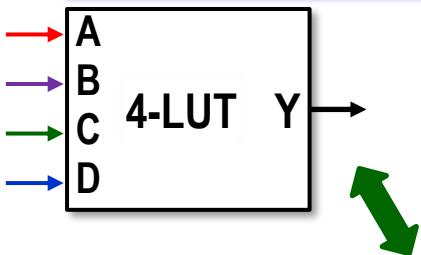
For configurations
Carry Chain
(transmission propagation)
will disconnect
C 4-LUT input
and it's being run
Carry In.



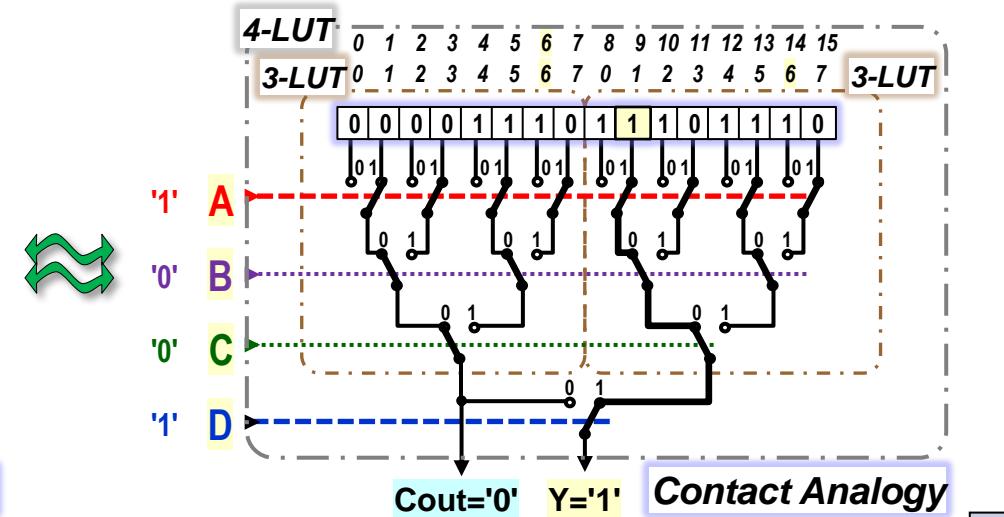
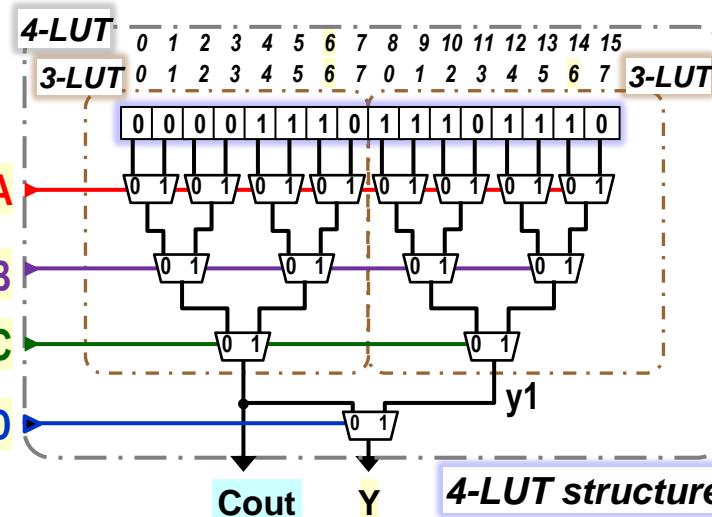
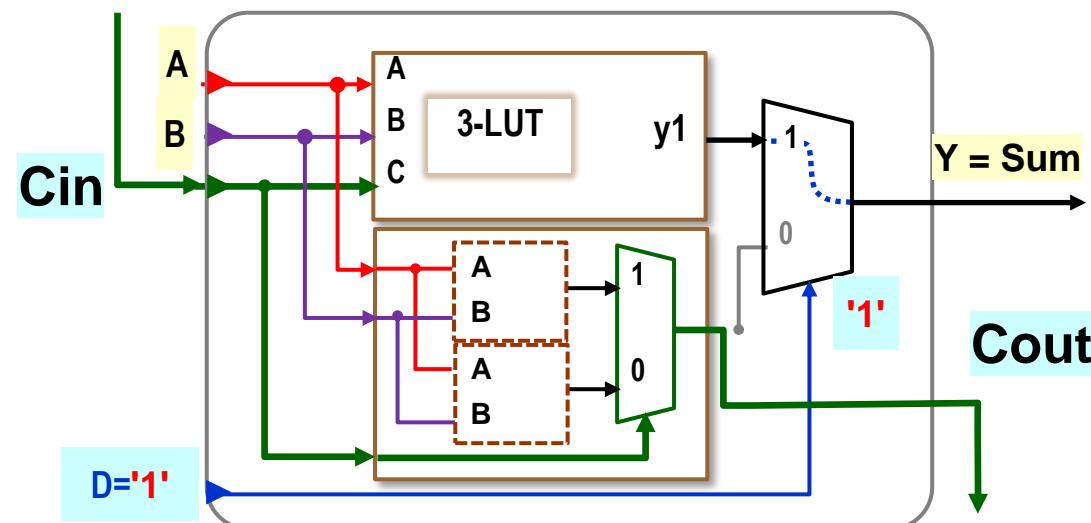
FPGA Carry Chain for "Ripple"



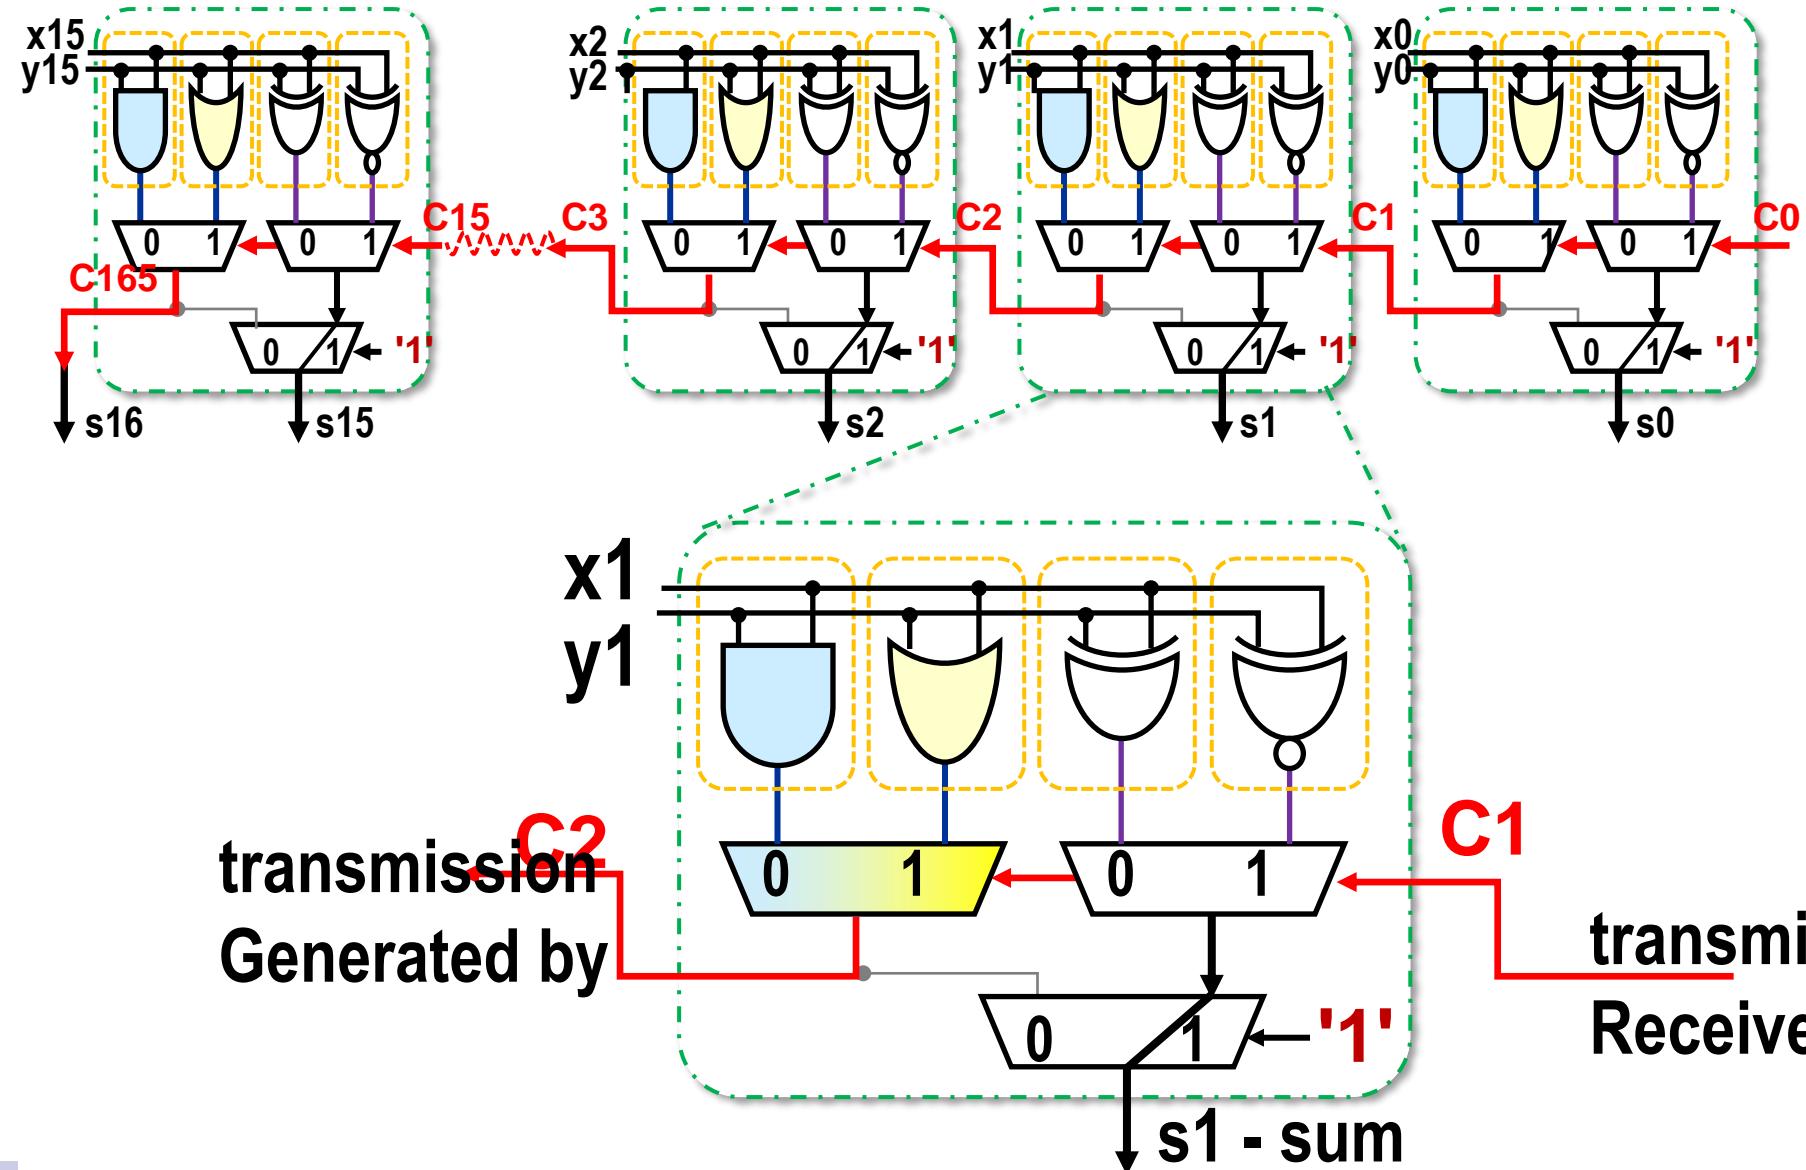
Primary Configuration $1 \times 4\text{-LUT}$



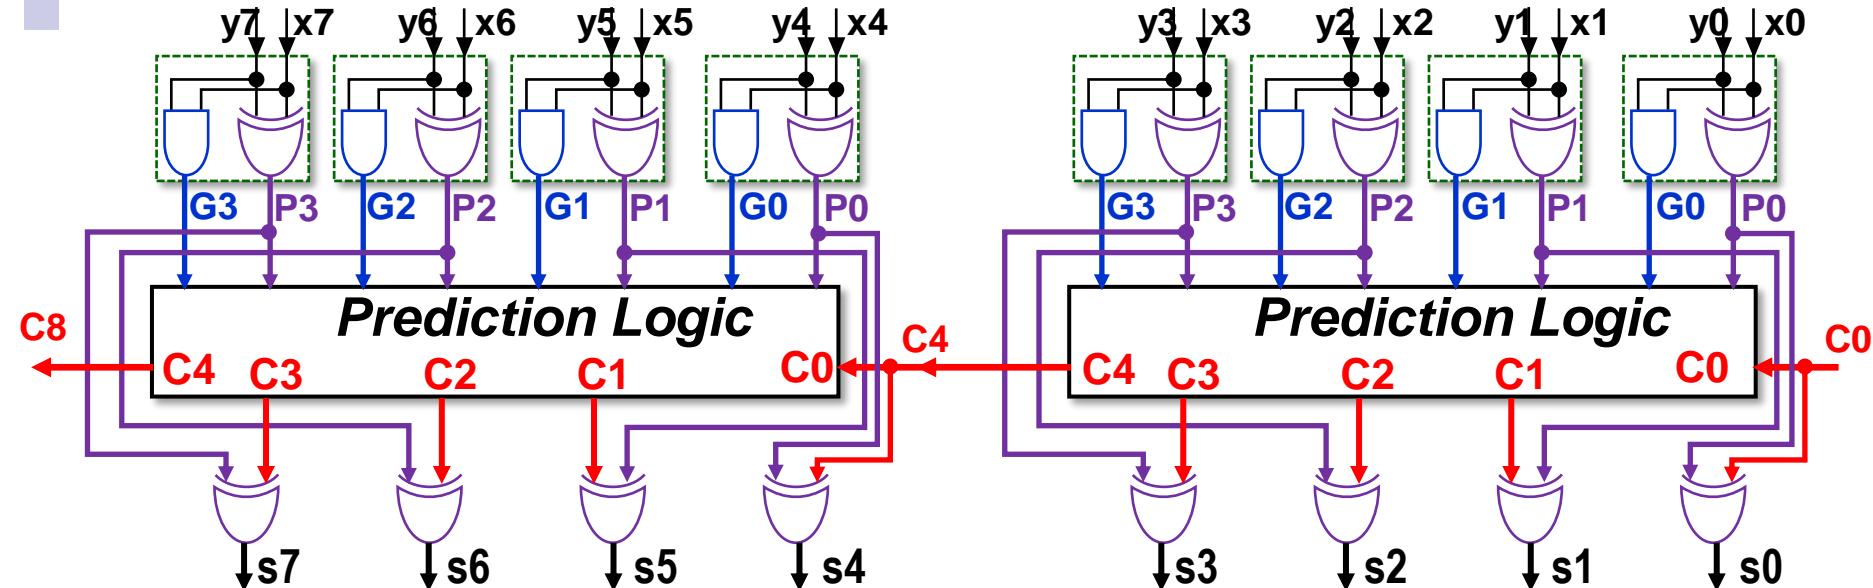
Carry Chain Configuration $2 \times 3\text{-LUT}$



Ripple Adder in FPGA



CLA (Carry-lookahead Adder)



Transfer equations:

$$C_1 = G_0 + P_0 \cdot C_0;$$

$$C_2 = G_1 + P_1 \cdot (G_0 + P_0 \cdot C_0) = G_1 + P_1 \cdot G_0 + P_1 \cdot P_0 \cdot C_0;$$

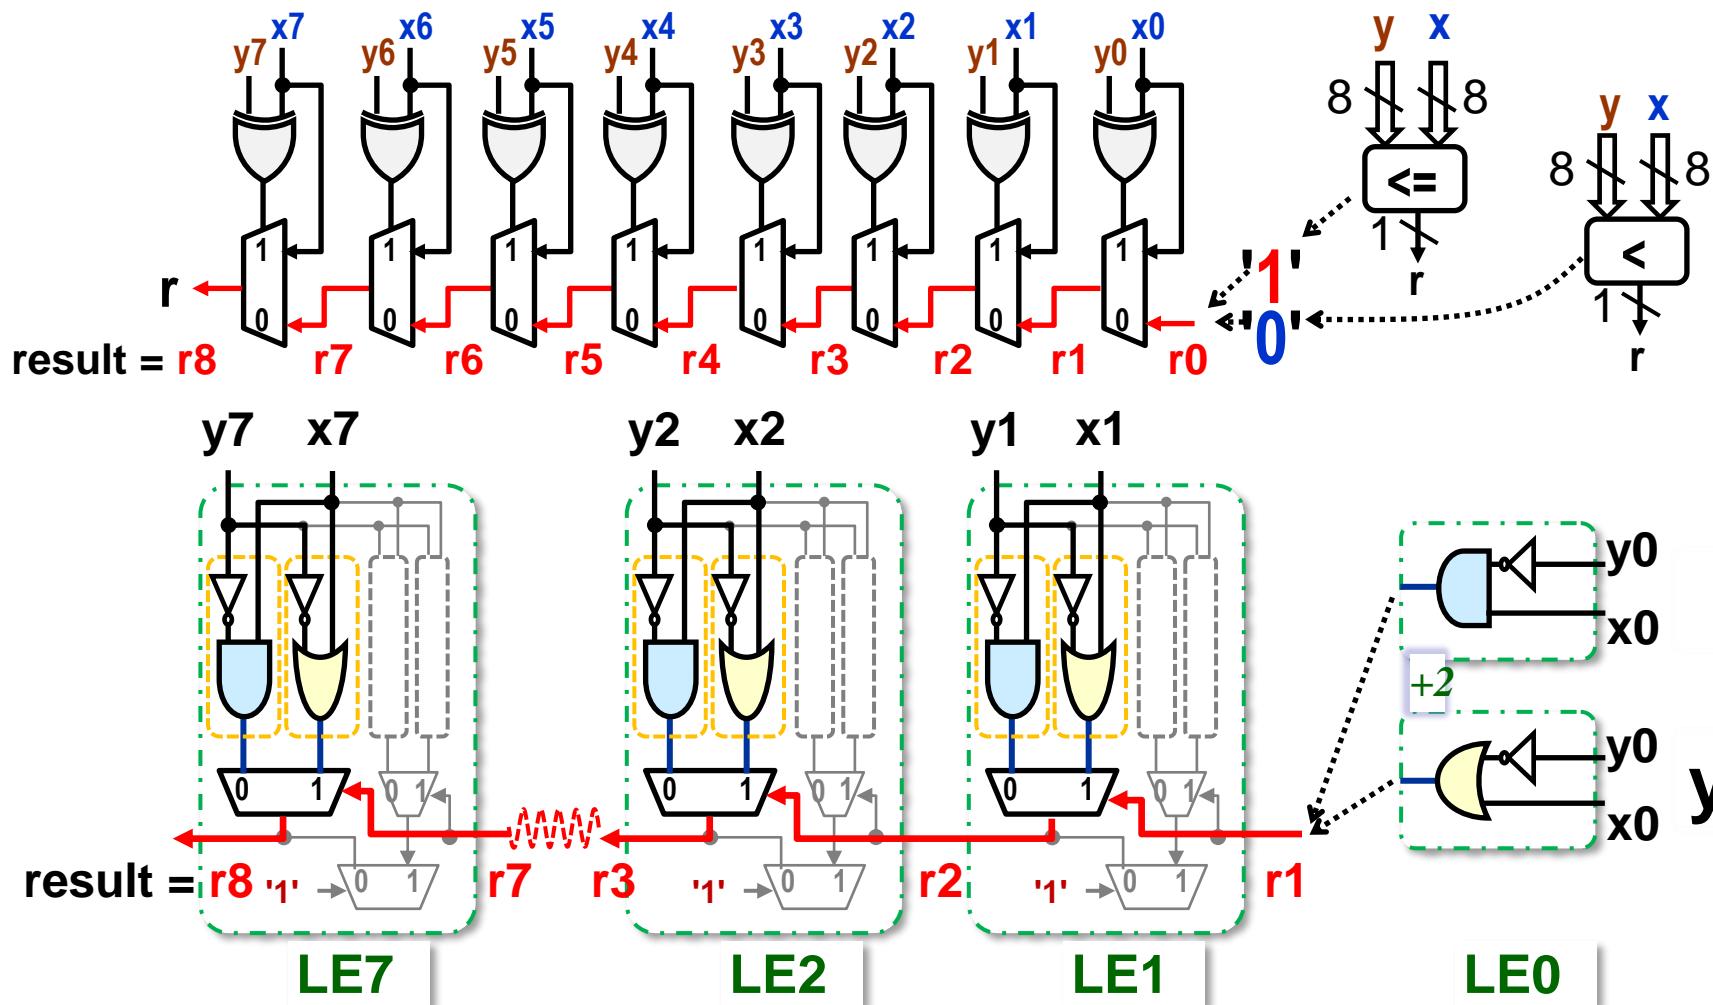
$$C_3 = G_2 + P_2 \cdot (G_1 + P_1 \cdot (G_0 + P_0 \cdot C_0)) = G_2 + P_2 \cdot G_1 + P_2 \cdot P_1 \cdot G_0 + P_2 \cdot P_1 \cdot P_0 \cdot C_0;$$

$$\begin{aligned} C_4 &= G_3 + P_3 \cdot (G_2 + P_2 \cdot (G_1 + P_1 \cdot (G_0 + P_0 \cdot C_0))) \\ &= G_3 + P_3 \cdot G_2 + P_3 \cdot P_2 \cdot G_1 + P_3 \cdot P_2 \cdot P_1 \cdot G_0 + P_3 \cdot P_2 \cdot P_1 \cdot P_0 \cdot C_0; \end{aligned}$$

etc.

They can be very quickly computed even over 128 bits on a tree structure, as is done by Kogge-Stone Adder. It requires about 50% more area than RCA.

Ripple structure has also comparator $x >= y$

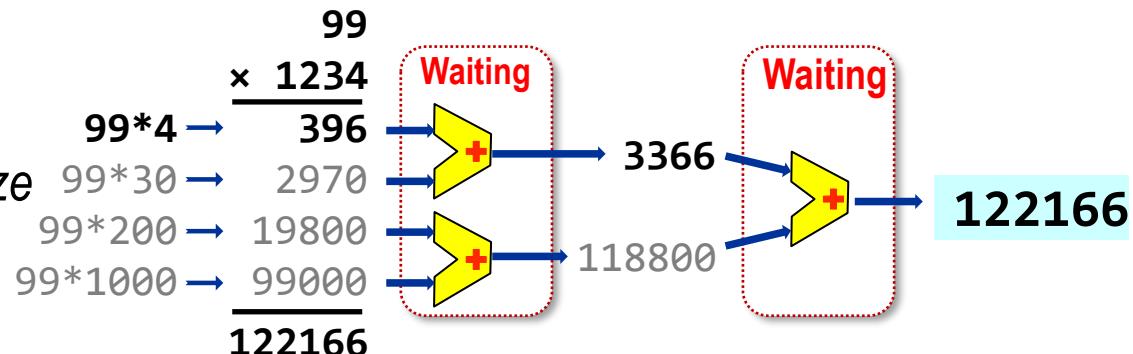


*RCA adders are preferred in FPGAs,
as their transfer chain can be modified to other operations.*

Wallace Tree Multiplier

The count is slow

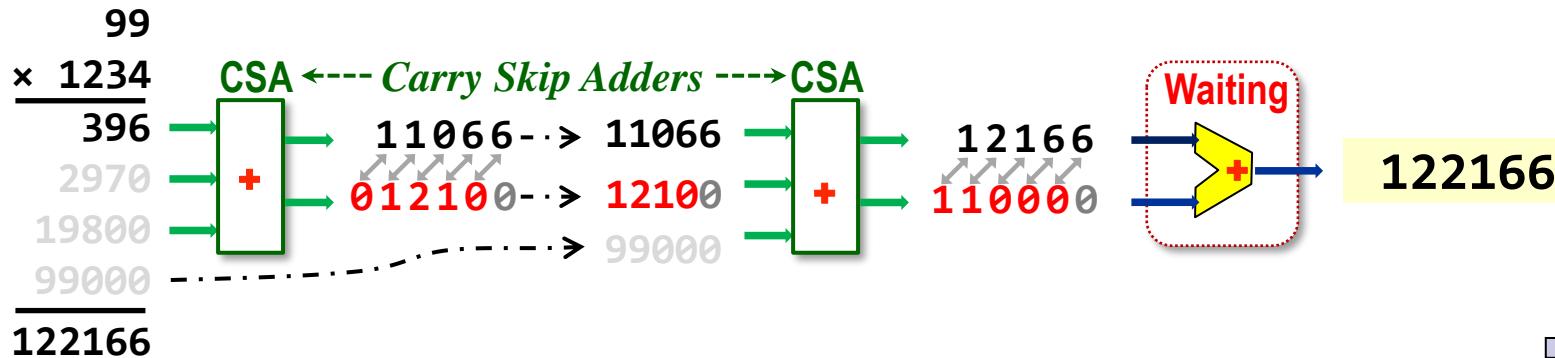
waiting for transmissions to stabilize
all intermediate results

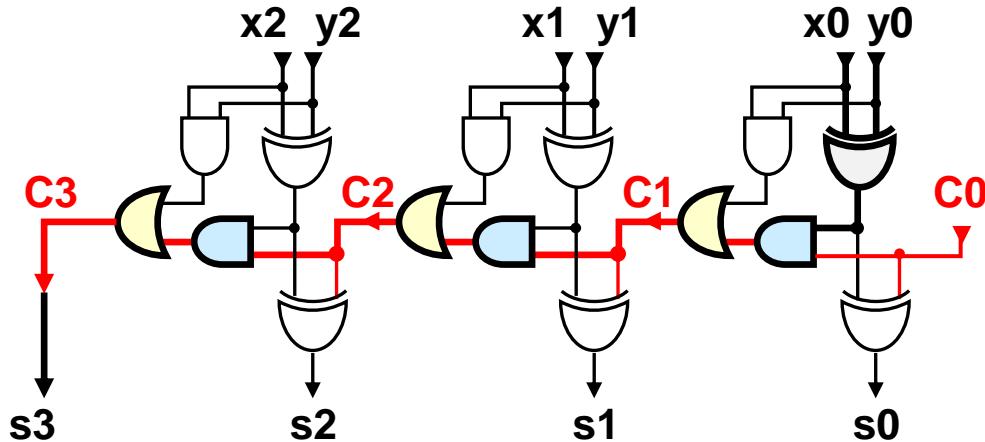


CSA (Carry Skip Adder): adds in orders of magnitude - reduces 3 input addends to 2 output values

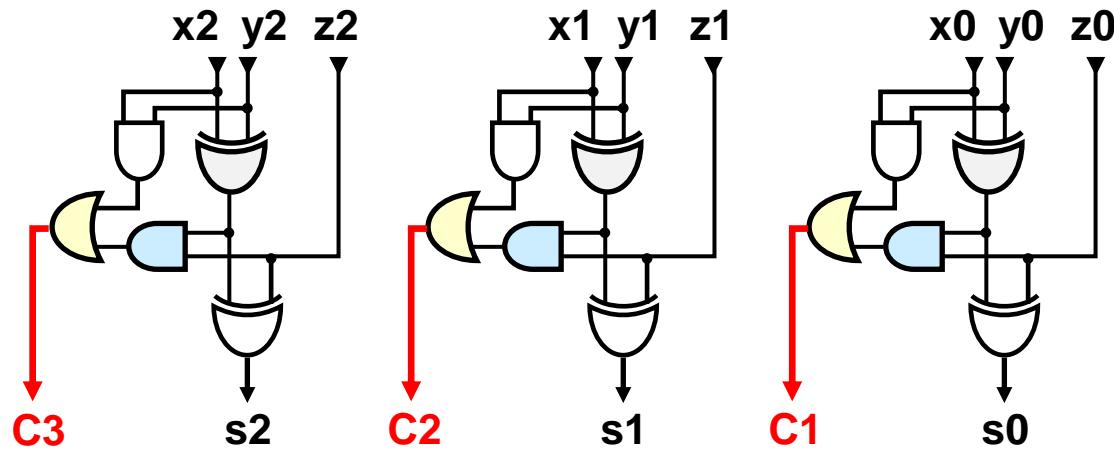
$$\begin{array}{rcl}
 396 & = & 300 + 90 + 6 \\
 2970 & = & 2000 + 900 + 70 + 0 \\
 +19800 & = & 10000 + 9000 + 800 \\
 \hline
 23166 & & 10000 + 11000 + 2000 + 160 + 06 \rightarrow 23166
 \end{array}$$

Fast multiplier reduces subtotals. One CSA only adds them at the end.





RCA - Ripple Carry Adder



CSA - Carry-skip Adder